

Homework 1

IDS 576

Name: Isaac Salvador
Email: isalva2@uic.edu
UIN: 6669845132

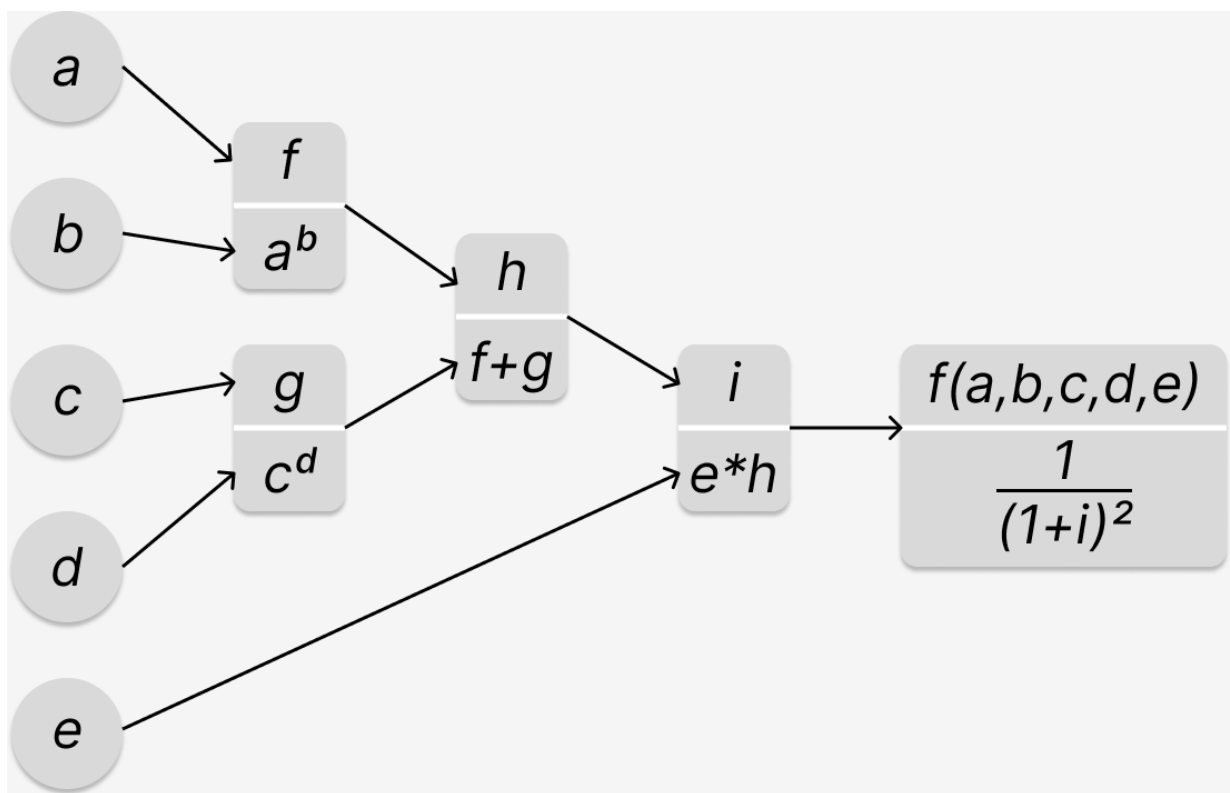
Name: Ahreum Kim
Email: akim239@uic.edu
UIN: 653241895

Name: Sadjad Bazarnovi
Email: sbazar3@uic.edu
UIN: 679314994

1. Backpropogation

Computational Graph

The multivariable equation $f(a, b, c, d, e) = \frac{1}{(1+(a^b+c^d)*e)^2}$ can be expressed with the following computational graph:



This graph begins with inputs a, b, c, d, e and terminates with the function in question.

This graph was created in [figma](#).

Gradient Computation

We wish to compute the gradient of f :

$$\nabla f = \left(\frac{\partial f}{\partial a}, \frac{\partial f}{\partial b}, \frac{\partial f}{\partial c}, \frac{\partial f}{\partial d}, \frac{\partial f}{\partial e} \right)$$

via the **chain rule**. The first step is to compute the partial derivatives of the **compute nodes** with respect to their inputs.

For the last compute node (let us define it as *node F*), we compute the partial derivative $\frac{\partial f}{\partial i}$ w.r.t. i as:

$$\frac{\partial f}{\partial i} = -\frac{2}{(1+i)^3}$$

Implemented in python, we get:

```
In [ ]: # import math and numpy
import math
import numpy as np
```

```

# assign input values
a,b,c,d,e = 1,1,1,1,1

# calculate the value of i
i = (a**b + c**d)*e

# compute the partial derivative df/di
def dF_di(i):
    return -2/(1+i)**3

```

We can then obtain the partial derivatives for the rest of the **compute nodes** w.r.t. their inputs:

```

In [ ]: # di/dh
h = a**b + c**d

def di_dh(e,h):
    return e

# di/de
def di_de(e,h):
    return h

# dh/df
def dh_df():
    return 1

# dh/dg
def dh_dg():
    return 1

# note that partial derivatives w.r.t f & g are linear and equate to 1

# df/da
def df_da(a,b):
    return b*a**(b-1)

# df/db
def df_db(a,b):
    return np.log(a)*a**b

# dg/dc
def dg_dc(c,d):
    return d*c**(d-1)

# dg/dd
def dg_dd(c,d):
    return np.log(c)*c**d

```

The first partial derivative $\frac{\partial f}{\partial a}$ can now be calculated:

$$\frac{\partial f}{\partial a} = \frac{\partial F}{\partial i} \frac{\partial i}{\partial h} \frac{\partial h}{\partial f} \frac{\partial f}{\partial a}$$

In python:

```
In [ ]: partial_a = dF_di(i)*di_dh(e,h)*dh_df()*df_da(a,b)

print(partial_a)
```

-0.07407407407407407

and summarily:

```
In [ ]: # compute partial derivatives w.r.t. b, c, d & e
partial_b = dF_di(i)*di_dh(e,h)*dh_df()*df_db(a,b)

partial_c = dF_di(i)*di_dh(e,h)*dh_dg()*dg_dc(c,b)

partial_d = dF_di(i)*di_dh(e,h)*dh_dg()*dg_dd(c,d)

partial_e = dF_di(i)*di_de(e,h)

# summary
gradient = [partial_a, partial_b, partial_c, partial_d, partial_e]

# chr index for summary
char = 97

print("Summary of Gradient:\n")

for partial in gradient:
    print(f"Partial derivative {chr(char)}: {format(partial, '.6f')}")
    char += 1

print(f"\nGradient of f(1,1,1,1,1) = {sum(gradient)}")
```

Summary of Gradient:

Partial derivative a: -0.074074
 Partial derivative b: -0.000000
 Partial derivative c: -0.074074
 Partial derivative d: -0.000000
 Partial derivative e: -0.148148

Gradient of f(1,1,1,1,1) = -0.2962962962962963

2. Gradient Descent

MSE Function

We wish to define a function in python that calculates *Mean Square Error*, defined by the

following function:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

In python:

```
In [ ]: def MSE(true, pred):  
  
    # user input error  
    if len(true) != len(pred):  
        print("Arrays are not the same size!")  
  
    # obtain sum of squares  
    sum_squares = np.square(true-pred)  
  
    # calculate MSE  
    mse = np.sum(sum_squares)/len(true)  
  
    return mse
```

Functionality test:

```
In [ ]: # initialize sample data obtained from statology.org  
observed = np.array([34, 37, 44, 47, 48, 48, 46, 43, 32, 27, 26, 24])  
predicted = np.array([37, 40, 46, 44, 46, 50, 45, 44, 34, 30, 22, 23])  
  
# compute MSE  
print(MSE(observed, predicted))
```

5.916666666666667

Note: the MSE Function computes the same answer as the generated answer by the [statology.org MSE Calculator](https://statology.org/mse-calculator/).

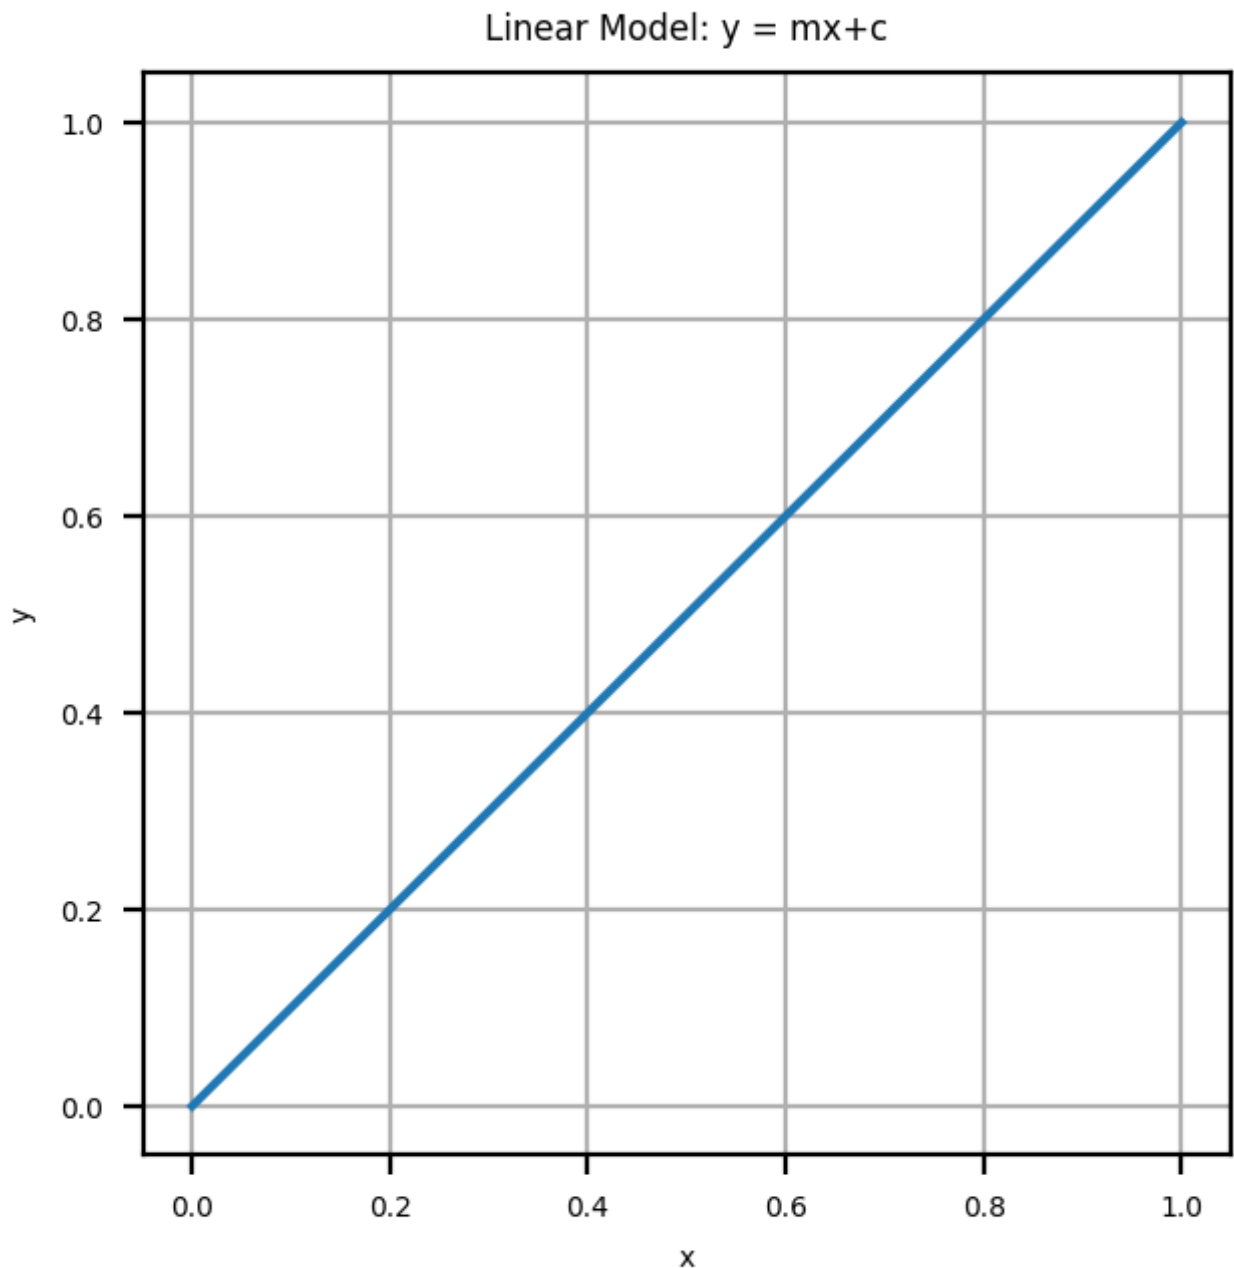
Linear Model

For the linear model $y = mx + c$, where the model parameters $m = 1$, $c = 0$, and $x \in (0, 1)$, the plot looks like this:

```
In [ ]: # import matplotlib  
import matplotlib.pyplot as plt  
plt.rcParams.update({'font.size': 5})  
  
# initialize model parameters  
m = 1  
c = 0  
  
# obtain x and y values for the plot  
x = np.linspace(0,1,2)
```

```
y = m*x+c

# plot linear model
fig, ax = plt.subplots(figsize=(3.5,3.5), dpi=200)
ax.plot(x, y)
ax.set_title("Linear Model: y = mx+c")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.grid(True)
```




Generate Example Data

We can now generate example data using `numpy`'s `random` module.

```
In [ ]: # initialize seed for reproducibility
np.random.seed(567)

# create random example data
x_example = np.random.rand(100)
y_example = m*x_example+c

# plot linear model
fig, ax = plt.subplots(figsize=(3.5, 3.5), dpi=200)
ax.scatter(x_example, y_example, zorder=1, c="r", s = 5)
ax.set_title("Random Example Data")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.grid(True)
ax.set_axisbelow(1)
```



Linear Gradient Descent

For the linear model with unknown parameters m and c , let us outline the **Gradient Descent** algorithm we shall employ. I implemented this [tutorial](#) as a guide.

1. Assign a random value to m and c such that $m \in [0, 1]$ and $c \in [0, 1]$.
2. Choose the learning rate α .
3. Compute the partial derivatives of the MSE loss function $\frac{\partial L}{\partial m}$ and $\frac{\partial L}{\partial c}$. More formally:

$$\frac{\partial L}{\partial m} = -\frac{2}{N} \sum_{i=1}^N (y_i - \hat{y}_i) x_i$$

$$\frac{\partial L}{\partial c} = -\frac{2}{N} \sum_{i=1}^N (y_i - \hat{y}_i)$$

4. Update the values of m and c in the direction of their respective partial derivative such that

$$m_{n+1} = m_n - \alpha_n \frac{\partial L}{\partial m}$$

$$c_{n+1} = c_n - \alpha_n \frac{\partial L}{\partial c}$$

5. Iterate this process for a set number of epochs or until a desirable loss value is achieved (we will implement the former).

To perform this task let's create a python `class` with the following inputs:

- `np.array` : input data x and y
- `float` : random initial starting parameters m and c between $[0,1]$
- `float` : learning rate lr
- `float` : learning rate $alpha$
- `int` : number of *epochs*
- `object` : partial derivatives of m and c

```
In [ ]: # create partial derivative function w.r.t. m
def LinearPartialM(x: np.array, y : np.array, yHat: np.array):

    return (-2/len(y))*np.sum((y-yHat)*x)

# create partial derivative function w.r.t. c
def LinearPartialC(y : np.array, yHat: np.array):
```



```
return (-2/len(y))*np.sum(y-yHat)
```

```
# create Gradient Descent Class
```

```
class GradientDescent:
```

```
    def __init__(self,  
                  x_example: np.array,  
                  y_example: np.array,  
                  lr: float,  
                  epochs: int,  
                  dLdm: object,  
                  dLdc: object,  
                  seed: int):
```

```
    # input variables
```

```
    self.xTrue = x_example # input x  
    self.yTrue = y_example # input y  
    self.learningRate = lr # learning rate  
    self.numEpochs = epochs # number of epochs  
    self.partialM = dLdm # partial w.r.t. m  
    self.partialC = dLdc # partial w.r.t. c  
    self.seed = seed # random seed for reproducibility
```

```
    # variables to be populated after runtime
```

```
    self.m = 0  
    self.c = 0  
    self.loss = 0
```

```
    # tracking metrics
```

```
    self.lossHistory = []
```

```
# run gradient descent
```

```
def start(self):
```

```
    # initialize index
```

```
    i = 0
```

```
    # generate random values for m and c
```

```
    np.random.seed(self.seed)  
    self.m = np.random.rand()  
    self.c = np.random.rand()
```

```
    # begin Gradient Descent algorithm
```

```
    while i < self.numEpochs:
```

```
        # obtain predictions based on current m and c
```

```
        yHat = self.m*self.xTrue + self.c
```

```
        # update parameters
```

```
        self.m += -1*self.learningRate*self.partialM(self.xTrue, self.y  
        self.c += -1*self.learningRate*self.partialC(self.yTrue, yHat)
```

```
        # update loss value
```

```

        self.loss = MSE(self.yTrue, yHat)
        self.lossHistory.append(self.loss)

        if i % 10 == 0:
            print(f"epoch: {i}, loss: {self.loss.round(10)}, m: {self.m}

        i += 1

    print(f"\nTraining Summary ({i} epochs): loss: {self.loss.round(10)}

    def history(self):
        return self.lossHistory, self.learningRate, self.m, self.c

```

Now we initialize an instance of the `GradientDescent` class and run the `start()` method.

```

In [ ]: # instantiate a GradientDescent instance
LinearModel1 = GradientDescent(x_example, y_example, 0.5, 200, LinearPartial

# run the algorithm
LinearModel1.start()

```

```

epoch: 0, loss: 0.5503705653, m: 0.30405, c: 0.17979
epoch: 10, loss: 0.0094255094, m: 0.68943, c: 0.16335
epoch: 20, loss: 0.0023711224, m: 0.84423, c: 0.08193
epoch: 30, loss: 0.0005964899, m: 0.92187, c: 0.04109
epoch: 40, loss: 0.0001500556, m: 0.96081, c: 0.02061
epoch: 50, loss: 3.77486e-05, m: 0.98035, c: 0.01034
epoch: 60, loss: 9.4962e-06, m: 0.99014, c: 0.00518
epoch: 70, loss: 2.3889e-06, m: 0.99506, c: 0.0026
epoch: 80, loss: 6.01e-07, m: 0.99752, c: 0.0013
epoch: 90, loss: 1.512e-07, m: 0.99876, c: 0.00065
epoch: 100, loss: 3.8e-08, m: 0.99938, c: 0.00033
epoch: 110, loss: 9.6e-09, m: 0.99969, c: 0.00016
epoch: 120, loss: 2.4e-09, m: 0.99984, c: 8e-05
epoch: 130, loss: 6e-10, m: 0.99992, c: 4e-05
epoch: 140, loss: 2e-10, m: 0.99996, c: 2e-05
epoch: 150, loss: 0.0, m: 0.99998, c: 1e-05
epoch: 160, loss: 0.0, m: 0.99999, c: 1e-05
epoch: 170, loss: 0.0, m: 1.0, c: 0.0
epoch: 180, loss: 0.0, m: 1.0, c: 0.0
epoch: 190, loss: 0.0, m: 1.0, c: 0.0

```

Training Summary (200 epochs): loss: 0.0, m: 1.0, c: 0.0

Our completed linear model is $y = 1.0x + 0.0$.

Plotting the error

We will now instantiate a second `GradientDescent` object with a different `lr` parameter and plot the error for the two objects.

```
In [ ]: # create second linear model
LinearModel2 = GradientDescent(x_example, y_example, 0.01, 200, LinearPartia

# run the algorithm
LinearModel2.start()
```

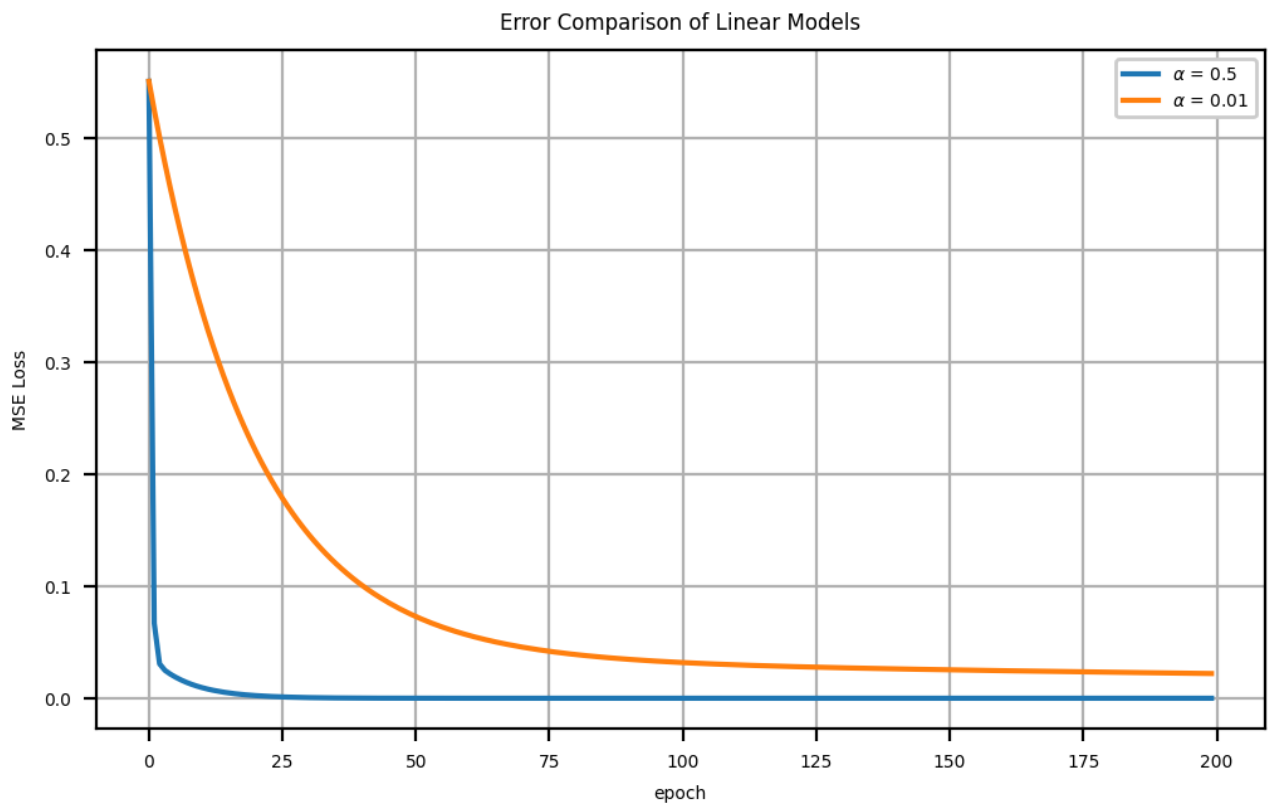
```
epoch: 0, loss: 0.5503705653, m: 0.62717, c: 0.89935
epoch: 10, loss: 0.3445637126, m: 0.57077, c: 0.7708
epoch: 20, loss: 0.2205747043, m: 0.52891, c: 0.67019
epoch: 30, loss: 0.1457405002, m: 0.49831, c: 0.59125
epoch: 40, loss: 0.1004413013, m: 0.47641, c: 0.52912
epoch: 50, loss: 0.0728919338, m: 0.46121, c: 0.48003
epoch: 60, loss: 0.0560132147, m: 0.4512, c: 0.44105
epoch: 70, loss: 0.0455526911, m: 0.44518, c: 0.40993
epoch: 80, loss: 0.0389558344, m: 0.44224, c: 0.38491
epoch: 90, loss: 0.0346880055, m: 0.44164, c: 0.36463
epoch: 100, loss: 0.0318271826, m: 0.44286, c: 0.34803
epoch: 110, loss: 0.0298193282, m: 0.44545, c: 0.33429
epoch: 120, loss: 0.028331445, m: 0.44908, c: 0.32279
epoch: 130, loss: 0.0271633124, m: 0.4535, c: 0.31304
epoch: 140, loss: 0.0261944696, m: 0.45851, c: 0.30464
epoch: 150, loss: 0.0253523695, m: 0.46396, c: 0.29732
epoch: 160, loss: 0.0245932505, m: 0.46972, c: 0.29082
epoch: 170, loss: 0.023890646, m: 0.47571, c: 0.28499
epoch: 180, loss: 0.0232284828, m: 0.48186, c: 0.27968
epoch: 190, loss: 0.0225969351, m: 0.4881, c: 0.27479
```

Training Summary (200 epochs): loss: 0.0220496515, m: 0.49377, c: 0.27067

We can now leverage the use of the `history()` method to obtain the errors of each model. MSE error is plotted in Log scale for clarity.

```
In [ ]: # obtain training history for both models
history1 = LinearModel1.history()
history2 = LinearModel2.history()

# plotting the errors
# plot linear model
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)
ax.plot(history1[0], label = r'$\alpha$ = 0.5')
ax.plot(history2[0], label = r'$\alpha$ = 0.01')
ax.set_title("Error Comparison of Linear Models")
ax.set_xlabel('epoch')
ax.set_ylabel('MSE Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)
```



Quadratic Gradient Descent

For the quadratic model $y = m_1x + m_2x^2 + c$, we obtain the following MSE loss function:

$$L = \frac{1}{n} \sum_{i=1}^n (y_i - (m_1x + m_2x^2 + c))^2$$

With the following partial derivatives:

$$\frac{\partial L}{\partial m_1} = -\frac{2}{n} \sum_{i=1}^n (y_i - (m_1x + m_2x^2 + c))x = -\frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i)x$$

$$\frac{\partial L}{\partial m_2} = -\frac{2}{n} \sum_{i=1}^n (y_i - (m_1x + m_2x^2 + c))x^2 = -\frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i)x^2$$

$$\frac{\partial L}{\partial c} = -\frac{2}{n} \sum_{i=1}^n (y_i - (m_1x + m_2x^2 + c)) = -\frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i)$$

In python:

```
In [ ]: # partial derivative w.r.t. m_1
def QuadraticPartialM1(x: np.array, y : np.array, yHat: np.array):
    return (-2/len(x))*np.dot(y-yHat,x)
```

```

def QuadraticPartialM2(x: np.array, y : np.array, yHat: np.array):
    return (-2/len(x))*np.dot(y-yHat,np.square(x))

def QuadraticPartialC(y : np.array, yHat: np.array):
    return (-2/len(y))*np.sum(y-yHat)

```

We can now modify the `GradientDescent` class to account for the third partial derivative.

```

In [ ]: # create Quadratic Gradient Descent Class
class QuadraticGradientDescent:
    def __init__(self,
                  x_example: np.array,
                  y_example: np.array,
                  lr: float,
                  epochs: int,
                  partialM1: object,
                  partialM2: object,
                  partialC: object,
                  seed: int):

        # input variables
        self.xTrue = x_example # input x
        self.yTrue = y_example # input y
        self.learningRate = lr # learning rate
        self.numEpochs = epochs # number of epochs
        self.partialM1 = partialM1 # partial w.r.t. m1
        self.partialM2 = partialM2 # partial w.r.t. m2
        self.partialC = partialC # partial w.r.t. c
        self.seed = seed # random seed for reproducibility

        # tracking metrics
        self.lossHistory = []

        # run gradient descent
        def start(self):

            # initialize index
            i = 0

            # generate random values for m and c
            np.random.seed(self.seed)
            self.m1 = np.random.rand()
            self.m2 = np.random.rand()
            self.c = np.random.rand()

            # begin Gradient Descent algorithm
            print(f"Initial Parameters: m1: {self.m1}, m2: {self.m2}, c: {self.c}")

            while i < self.numEpochs:

                # obtain predictions based on current m1, m2, and c

```

```

        yHat = self.m1*self.xTrue + self.m2*np.square(self.xTrue)+self.c

        # update parameters
        self.m1 += -1*self.learningRate*self.partialM1(self.xTrue, self.yTrue)
        self.m2 += -1*self.learningRate*self.partialM2(self.xTrue, self.yTrue)
        self.c += -1*self.learningRate*self.partialC(self.yTrue, yHat)

        # update loss value
        self.loss = MSE(self.yTrue, yHat)
        self.lossHistory.append(self.loss)

        if i % 10 == 0:
            print(f"epoch: {i}, loss: {self.loss.round(10)}, m1: {self.m1.round(10)}, m2: {self.m2.round(10)}, c: {self.c.round(10)}")

        i += 1

    print(f"\nTraining Summary ({i} epochs): loss: {self.loss.round(10)}, m1: {self.m1.round(10)}, m2: {self.m2.round(10)}, c: {self.c.round(10)}")

    def history(self):
        return self.lossHistory, self.learningRate, self.m1, self.m2, self.c

```

Now we create sample data and instantiate instances of the `QuadraticGradientDescent` object to create two models.

```

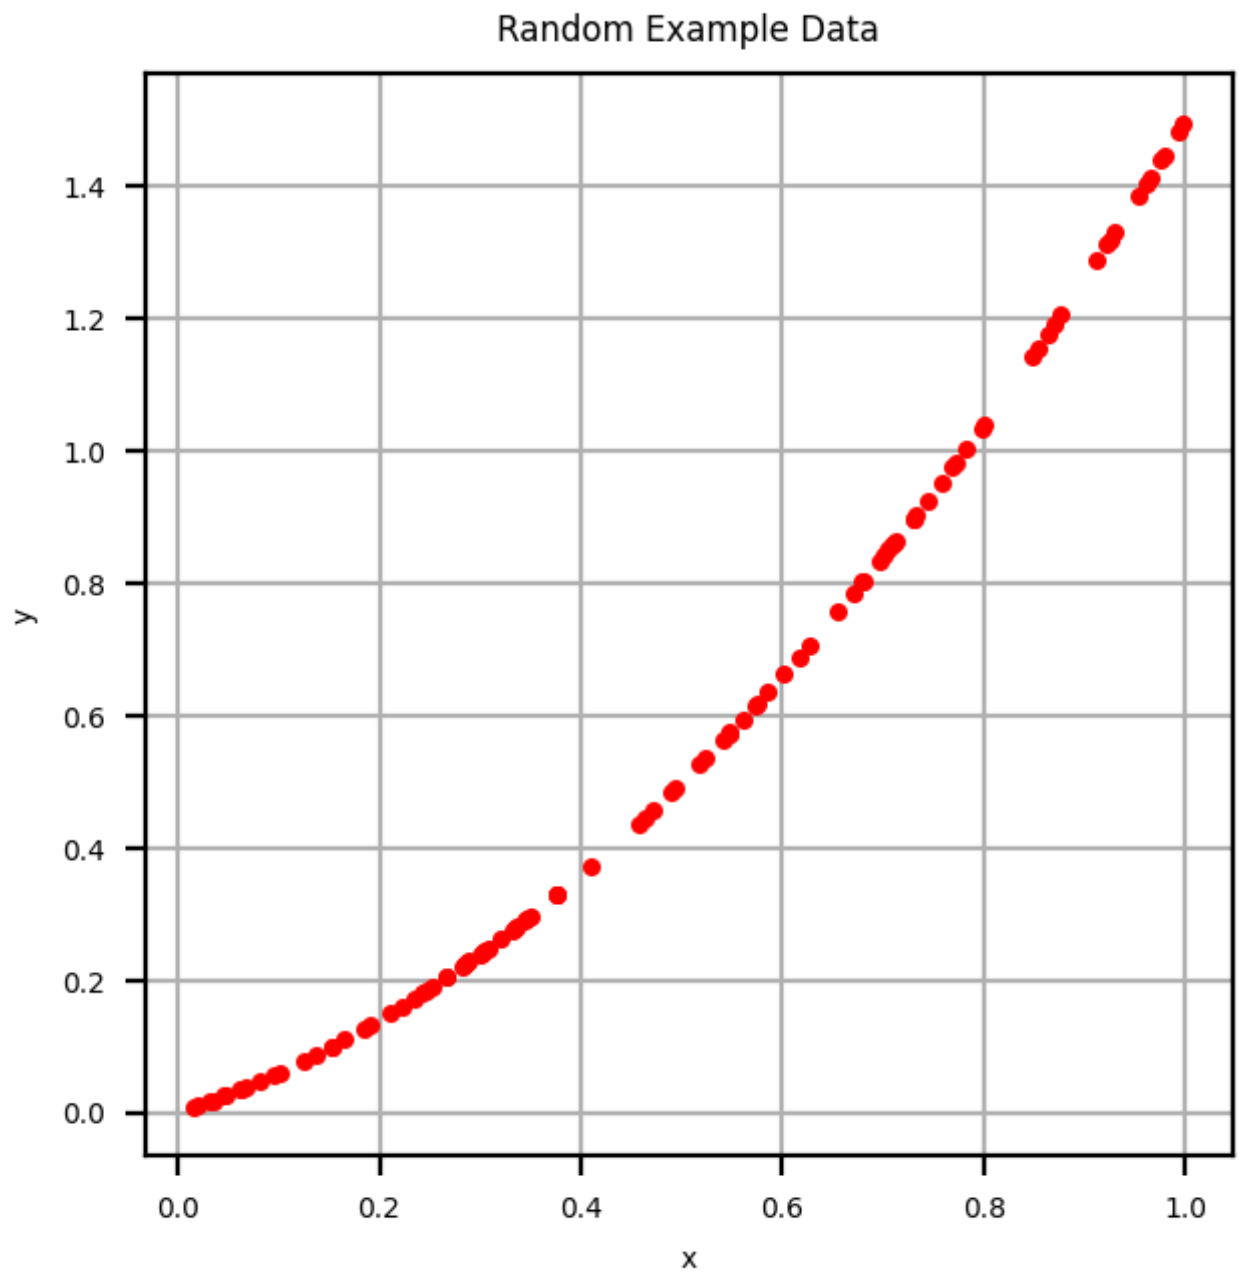
In [ ]: # initialize seed for reproducibility
np.random.seed(567)

m1 = 0.5
m2 = 1.0
c = 0

# create random example data
x_example_quadratic = np.random.rand(100)
y_example_quadratic = m1*x_example_quadratic + m2*x_example_quadratic**2 + c

# plot linear model
fig, ax = plt.subplots(figsize=(3.5, 3.5), dpi=200)
ax.scatter(x_example_quadratic, y_example_quadratic, zorder=1, c="r", s = 5)
ax.set_title("Random Example Data")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.grid(True)
ax.set_axisbelow(1)

```



```
In [ ]: # create object
QuadraticModel1 = QuadraticGradientDescent(x_example_quadratic, y_example_qu
0.1,
300,
QuadraticPartialM1,
QuadraticPartialM2,
QuadraticPartialC,
576)

# run algorithm
QuadraticModel1.start()
```

Initial Parameters: m1: 0.6337605448331272, m2: 0.9140319527336822, c: 0.9137978208648712

```
epoch: 0, loss: 0.9056425945, m1: 0.53955, m2: 0.85151, c: 0.72349
epoch: 10, loss: 0.0185745316, m1: 0.34439, m2: 0.7447, c: 0.20914
epoch: 20, loss: 0.010418527, m1: 0.37771, m2: 0.79118, c: 0.14784
epoch: 30, loss: 0.0062740964, m1: 0.4106, m2: 0.83222, c: 0.11318
epoch: 40, loss: 0.0037795234, m1: 0.43635, m2: 0.86429, c: 0.08677
epoch: 50, loss: 0.0022773384, m1: 0.45632, m2: 0.88921, c: 0.06631
epoch: 60, loss: 0.0013727439, m1: 0.47178, m2: 0.90857, c: 0.05044
epoch: 70, loss: 0.000828004, m1: 0.48375, m2: 0.92363, c: 0.03812
epoch: 80, loss: 0.0004999601, m1: 0.49301, m2: 0.93535, c: 0.02858
epoch: 90, loss: 0.0003024053, m1: 0.50016, m2: 0.94447, c: 0.02117
epoch: 100, loss: 0.0001834281, m1: 0.50568, m2: 0.95157, c: 0.01543
epoch: 110, loss: 0.0001117686, m1: 0.50994, m2: 0.95712, c: 0.01098
epoch: 120, loss: 6.86029e-05, m1: 0.51321, m2: 0.96145, c: 0.00754
epoch: 130, loss: 4.25955e-05, m1: 0.51572, m2: 0.96483, c: 0.00487
epoch: 140, loss: 2.69206e-05, m1: 0.51763, m2: 0.96749, c: 0.0028
epoch: 150, loss: 1.74678e-05, m1: 0.51909, m2: 0.96958, c: 0.0012
epoch: 160, loss: 1.1762e-05, m1: 0.52019, m2: 0.97123, c: -3e-05
epoch: 170, loss: 8.3127e-06, m1: 0.52102, m2: 0.97254, c: -0.00099
epoch: 180, loss: 6.2223e-06, m1: 0.52163, m2: 0.97359, c: -0.00172
epoch: 190, loss: 4.9504e-06, m1: 0.52208, m2: 0.97442, c: -0.00229
epoch: 200, loss: 4.1715e-06, m1: 0.5224, m2: 0.9751, c: -0.00272
epoch: 210, loss: 3.6896e-06, m1: 0.52261, m2: 0.97565, c: -0.00305
epoch: 220, loss: 3.3867e-06, m1: 0.52276, m2: 0.97611, c: -0.0033
epoch: 230, loss: 3.1918e-06, m1: 0.52284, m2: 0.97649, c: -0.00349
epoch: 240, loss: 3.0619e-06, m1: 0.52287, m2: 0.97681, c: -0.00363
epoch: 250, loss: 2.9714e-06, m1: 0.52287, m2: 0.97709, c: -0.00374
epoch: 260, loss: 2.9047e-06, m1: 0.52284, m2: 0.97733, c: -0.00382
epoch: 270, loss: 2.8526e-06, m1: 0.52279, m2: 0.97755, c: -0.00387
epoch: 280, loss: 2.8092e-06, m1: 0.52273, m2: 0.97774, c: -0.00391
epoch: 290, loss: 2.7713e-06, m1: 0.52265, m2: 0.97791, c: -0.00393
```

Training Summary (300 epochs): loss: 2.7402e-06, m1: 0.52257, m2: 0.97806, c : -0.00395

```
In [ ]: # create object
        QuadraticModel2 = QuadraticGradientDescent(x_example_quadratic, y_example_quadratic,
                                                    0.01,
                                                    300,
                                                    QuadraticPartialM1,
                                                    QuadraticPartialM2,
                                                    QuadraticPartialC,
                                                    576)

        # run algorithm
        QuadraticModel2.start()
```


Initial Parameters: m1: 0.6337605448331272, m2: 0.9140319527336822, c: 0.9137978208648712

```
epoch: 0, loss: 0.9056425945, m1: 0.62434, m2: 0.90778, c: 0.89477
epoch: 10, loss: 0.5256892006, m1: 0.54423, m2: 0.85508, c: 0.73049
epoch: 20, loss: 0.3092160459, m1: 0.48522, m2: 0.81706, c: 0.60523
epoch: 30, loss: 0.1856784656, m1: 0.44208, m2: 0.79007, c: 0.50942
epoch: 40, loss: 0.114983147, m1: 0.41088, m2: 0.77136, c: 0.43585
epoch: 50, loss: 0.0743429888, m1: 0.38864, m2: 0.75884, c: 0.37907
epoch: 60, loss: 0.0508064356, m1: 0.37313, m2: 0.75095, c: 0.33499
epoch: 70, loss: 0.0370117377, m1: 0.36266, m2: 0.74652, c: 0.30052
epoch: 80, loss: 0.0287740048, m1: 0.35595, m2: 0.74466, c: 0.27331
epoch: 90, loss: 0.0237138744, m1: 0.35206, m2: 0.74469, c: 0.25161
epoch: 100, loss: 0.0204781523, m1: 0.35026, m2: 0.74612, c: 0.2341
epoch: 110, loss: 0.0182968846, m1: 0.35, m2: 0.74856, c: 0.21977
epoch: 120, loss: 0.0167316289, m1: 0.35089, m2: 0.75174, c: 0.20786
epoch: 130, loss: 0.0155323905, m1: 0.3526, m2: 0.75544, c: 0.19781
epoch: 140, loss: 0.014556313, m1: 0.35492, m2: 0.7595, c: 0.18919
epoch: 150, loss: 0.0137214675, m1: 0.35767, m2: 0.7638, c: 0.18167
epoch: 160, loss: 0.0129805875, m1: 0.36071, m2: 0.76825, c: 0.17501
epoch: 170, loss: 0.0123061403, m1: 0.36396, m2: 0.77278, c: 0.16902
epoch: 180, loss: 0.0116818412, m1: 0.36733, m2: 0.77734, c: 0.16355
epoch: 190, loss: 0.0110978292, m1: 0.37078, m2: 0.7819, c: 0.15851
epoch: 200, loss: 0.0105479236, m1: 0.37425, m2: 0.78644, c: 0.15381
epoch: 210, loss: 0.0100280635, m1: 0.37773, m2: 0.79092, c: 0.14938
epoch: 220, loss: 0.0095354193, m1: 0.38119, m2: 0.79534, c: 0.14519
epoch: 230, loss: 0.0090678873, m1: 0.38462, m2: 0.79968, c: 0.1412
epoch: 240, loss: 0.0086238, m1: 0.388, m2: 0.80395, c: 0.13737
epoch: 250, loss: 0.008201761, m1: 0.39132, m2: 0.80813, c: 0.1337
epoch: 260, loss: 0.0078005501, m1: 0.39459, m2: 0.81222, c: 0.13015
epoch: 270, loss: 0.0074190681, m1: 0.39778, m2: 0.81622, c: 0.12672
epoch: 280, loss: 0.0070563041, m1: 0.40091, m2: 0.82014, c: 0.1234
epoch: 290, loss: 0.0067113167, m1: 0.40398, m2: 0.82396, c: 0.12018
```

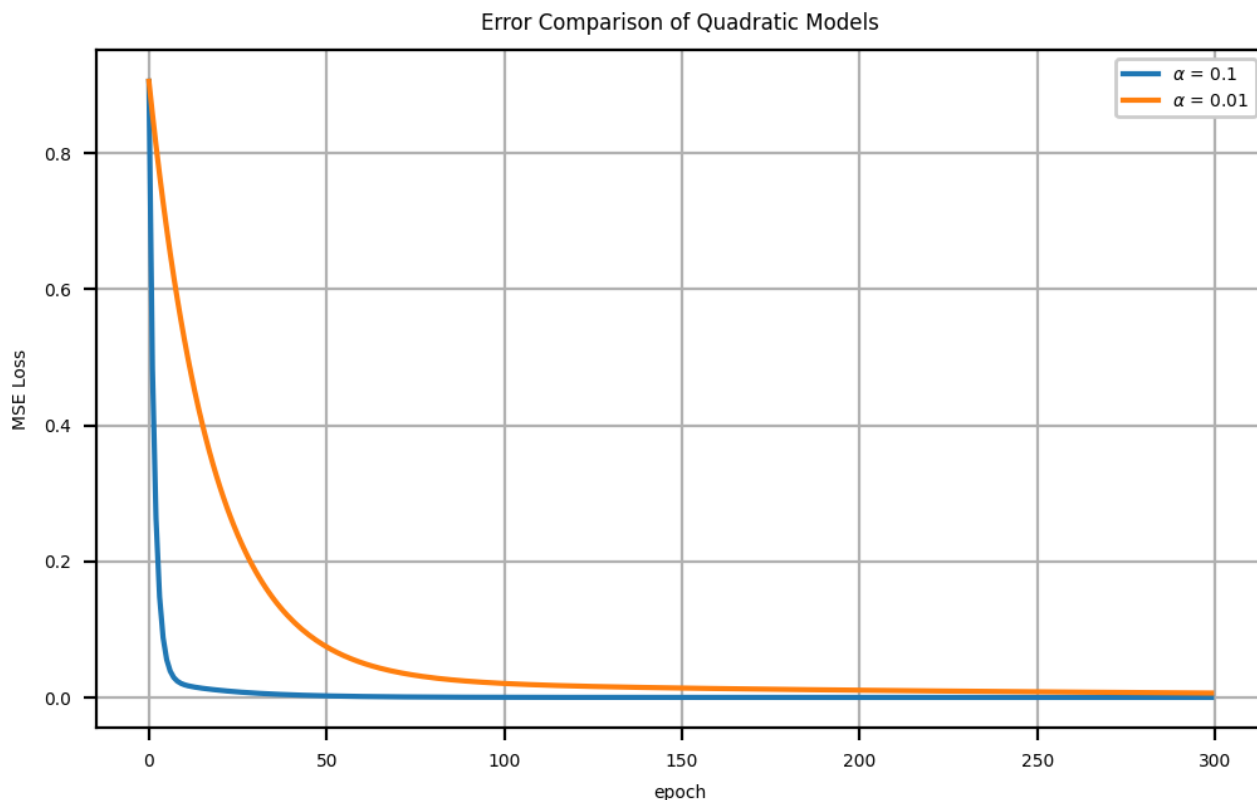
Training Summary (300 epochs): loss: 0.0064152952, m1: 0.40667, m2: 0.82732, c: 0.11736

We plot the loss of each quadratic model with the `history` method:

```
In [ ]: # obtain training history for both models
history1 = QuadraticModel1.history()
history2 = QuadraticModel2.history()

# plotting the errors
# plot linear model
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)
ax.plot(history1[0], label = r'$\alpha$ = 0.1')
ax.plot(history2[0], label = r'$\alpha$ = 0.01')
ax.set_title("Error Comparison of Quadratic Models")
ax.set_xlabel('epoch')
ax.set_ylabel('MSE Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
```

```
ax.grid(True)
```



Similarly, the parameters m_1 , m_2 , and c are stored in the `history` object and can be used to plot the models against the ground truth model.

```
In [ ]: # obtain model parameters
model1 = history1[2:]
model2 = history2[2:]

# ground truth model
truth_model = (0.5, 1, 0)

# define function for plotting
def model_plotting(model: tuple):

    # create x values
    x = np.linspace(0, 1)

    return x, model[0]*x + model[1]*x**2 + model[2]

# plot linear model
fig, ax = plt.subplots(figsize=(3.5, 3.5), dpi=200)

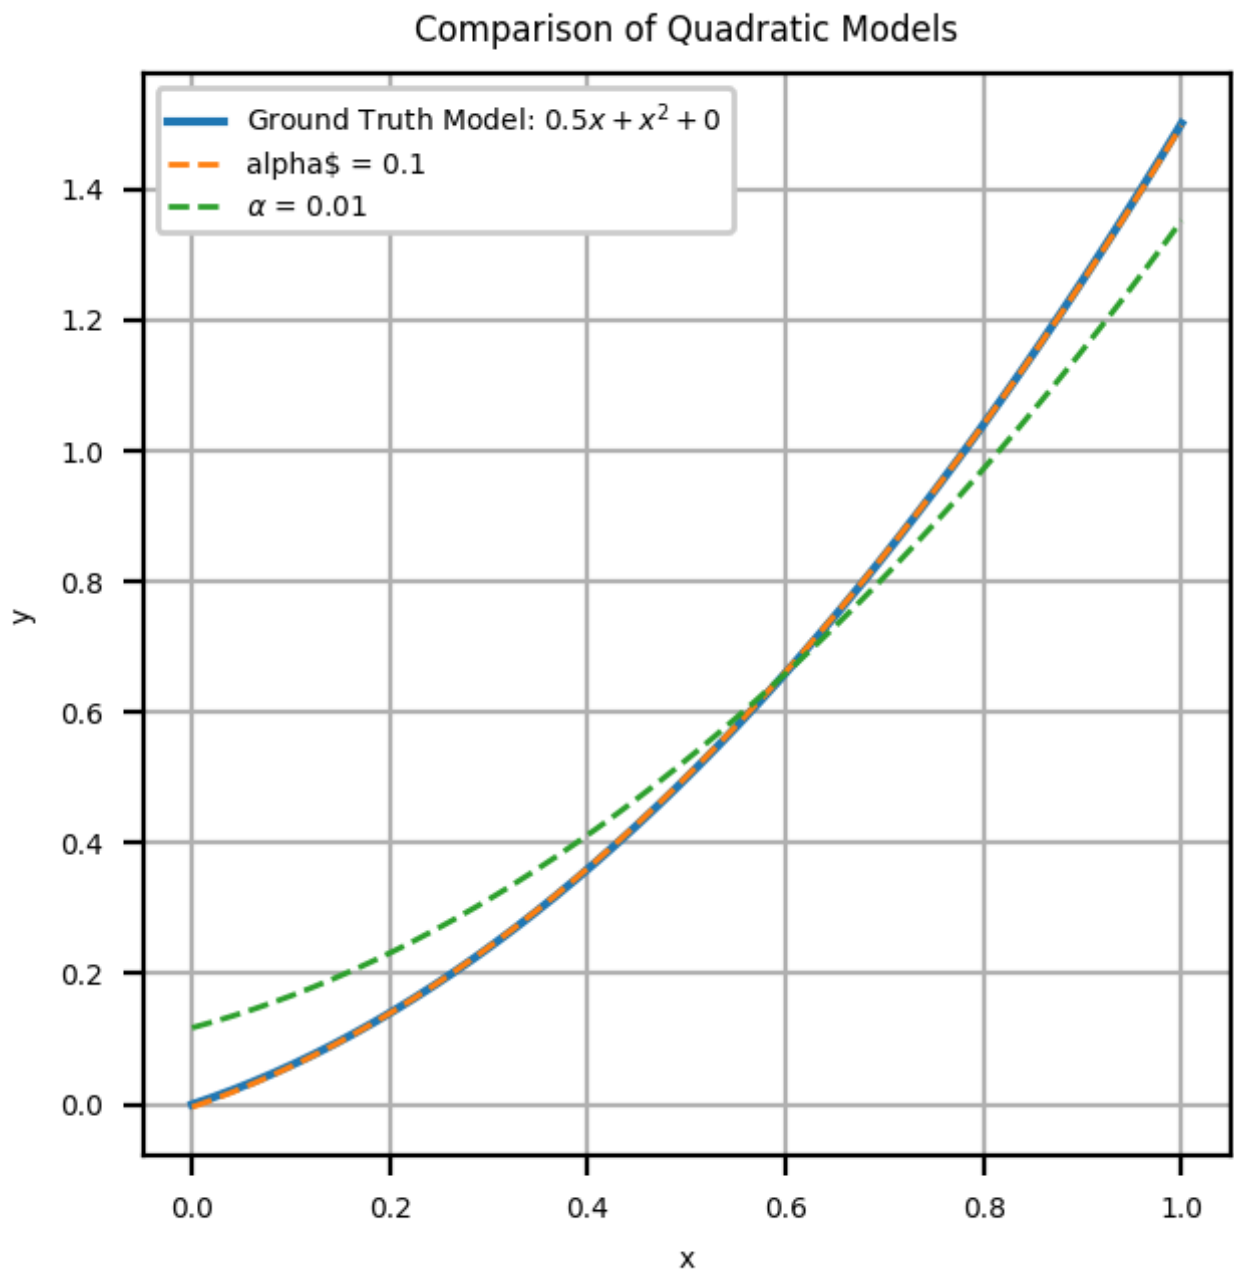
ax.plot(model_plotting(truth_model)[0],
        model_plotting(truth_model)[1],
        label = r'Ground Truth Model: $0.5x+x^2+0$'
        )

ax.plot(model_plotting(model1)[0],
```

```
    model_plotting(model1)[1],
    linestyle='dashed',
    label = r'alpha$ = 0.1',
    linewidth = 1)

ax.plot(model_plotting(model2)[0],
        model_plotting(model2)[1],
        linestyle='dashed',
        label = r'$\alpha$ = 0.01',
        linewidth=1)

ax.set_title("Comparison of Quadratic Models")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)
```



Gradient Descent techniques for both the Linear and Quadratic models exhibit better accuracy at lower learning rates. Large learning rates and a small number of epochs allows the algorithm to converge on the local minima of the MSE loss function. It is important to note that for both implementations of the gradient descent algorithm learning rates above $\alpha > 1, 0$ resulted in exploding gradients, making the models unable to learn. Error logs indicate that this is due to out of bound exceptions at run time.

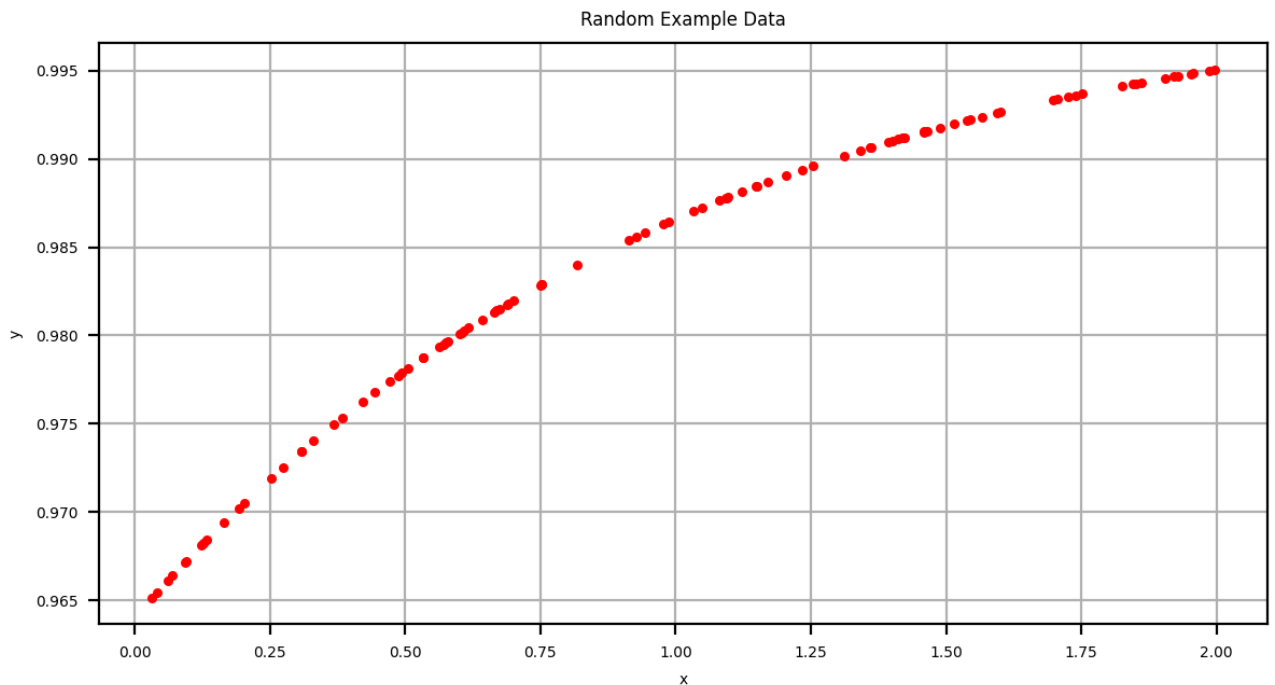
Hyperbolic Tangent Gradient Descent

For the hyperbolic tangent model $y = \tanh(m * x + c)$ we obtain example data for $x \in [0.2]$ in the following manner:

```
In [ ]: # initialize seed for reproducibility
np.random.seed(567)

# create random example data
x_hyperbolic = 2*np.random.rand(100)
y_hyperbolic = np.tanh(x_example+2)

# plot linear model
fig, ax = plt.subplots(figsize=(7, 3.5), dpi=200)
ax.scatter(x_hyperbolic, y_hyperbolic, zorder=1, c="r", s = 5)
ax.set_title("Random Example Data")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.grid(True)
ax.set_axisbelow(1)
```



similar to the quadratic mode, we obtain the loss function

$$L = \frac{1}{n} \sum_{i=1}^n (y_i - \tanh(m * x + c))^2$$

w.r.t. parameters c and m , the partial derivatives are as follows:

$$\frac{\partial L}{\partial m} = \frac{2}{n} \sum_{i=1}^n (y_i - \tanh(m * x + c))(1 - \tanh^2(m * x + c))x = \frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i)(1 - \hat{y}_i^2)x$$

$$\frac{\partial L}{\partial c} = \frac{2}{n} \sum_{i=1}^n (y_i - \tanh(m * x + c))(1 - \tanh^2(m * x + c)) = \frac{2}{n} \sum_{i=1}^n (y_i - \hat{y}_i)(1 - \hat{y}_i^2)$$

Implemented in python:

```
In [ ]: # partial w.r.t. m
def HyperbolicPartialM(x, y, yHat):
    return (2/len(x))*np.dot(y-yHat, (1 - np.square(yHat)*x))

# partial w.r.t. c
def HyperbolicPartialC(y, yHat):
    return (2/len(y))*np.dot(y-yHat, 1 - np.square(yHat))
```

We can now leverage a modified `GradientDescent` class to perform the algorithm.

```
In [ ]: # create Gradient Descent Class
class HyperbolicGradientDescent:
    def __init__(self,
                  x_example: np.array,
                  y_example: np.array,
                  lr: float,
                  epochs: int,
                  dLdm: object,
                  dLdc: object,
                  seed: int):

        # input variables
        self.xTrue = x_example # input x
        self.yTrue = y_example # input y
        self.learningRate = lr # learning rate
        self.numEpochs = epochs # number of epochs
        self.partialM = dLdm # partial w.r.t. m
        self.partialC = dLdc # partial w.r.t. c
        self.seed = seed # random seed for reproducibility

        # variables to be populated after runtime
        self.m = 0
        self.c = 0
        self.loss = 0

        # tracking metrics
        self.lossHistory = []

        # run gradient descent
        def start(self):

            # initialize index
            i = 0

            # generate random values for m and c
            np.random.seed(self.seed)
            self.m = np.random.rand()
            self.c = np.random.rand()

            # begin Gradient Descent algorithm
            print(f"Initial Parameters: m: {self.m} c: {self.c}\n")
```


Initial Parameters: m: 0.6337605448331272 c: 0.9140319527336822

```
epoch: 0, loss: 0.0126170849, m: 0.70904, c: 0.95649
epoch: 20, loss: 0.0013042466, m: 1.46273, c: 1.20034
epoch: 40, loss: 0.0006682789, m: 1.85625, c: 1.27604
epoch: 60, loss: 0.0004666872, m: 2.13117, c: 1.32135
epoch: 80, loss: 0.000372775, m: 2.34157, c: 1.35354
epoch: 100, loss: 0.0003202143, m: 2.5104, c: 1.37838
epoch: 120, loss: 0.0002874248, m: 2.64988, c: 1.39854
epoch: 140, loss: 0.00026544, m: 2.76737, c: 1.41544
epoch: 160, loss: 0.0002499179, m: 2.86769, c: 1.42997
epoch: 180, loss: 0.0002385219, m: 2.95421, c: 1.44267
epoch: 200, loss: 0.0002298932, m: 3.02939, c: 1.45394
epoch: 220, loss: 0.0002231933, m: 3.09505, c: 1.46405
epoch: 240, loss: 0.0002178798, m: 3.15264, c: 1.47321
epoch: 260, loss: 0.0002135887, m: 3.20327, c: 1.48158
epoch: 280, loss: 0.0002100671, m: 3.24786, c: 1.48927
epoch: 300, loss: 0.0002071355, m: 3.28716, c: 1.49637
epoch: 320, loss: 0.0002046628, m: 3.32178, c: 1.50298
epoch: 340, loss: 0.000202552, m: 3.35226, c: 1.50914
epoch: 360, loss: 0.0002007296, m: 3.37904, c: 1.51492
epoch: 380, loss: 0.0001991395, m: 3.40249, c: 1.52036
```

Training Summary (400 epochs): loss: 0.0001978043, m: 3.422, c: 1.52523

[illegible]

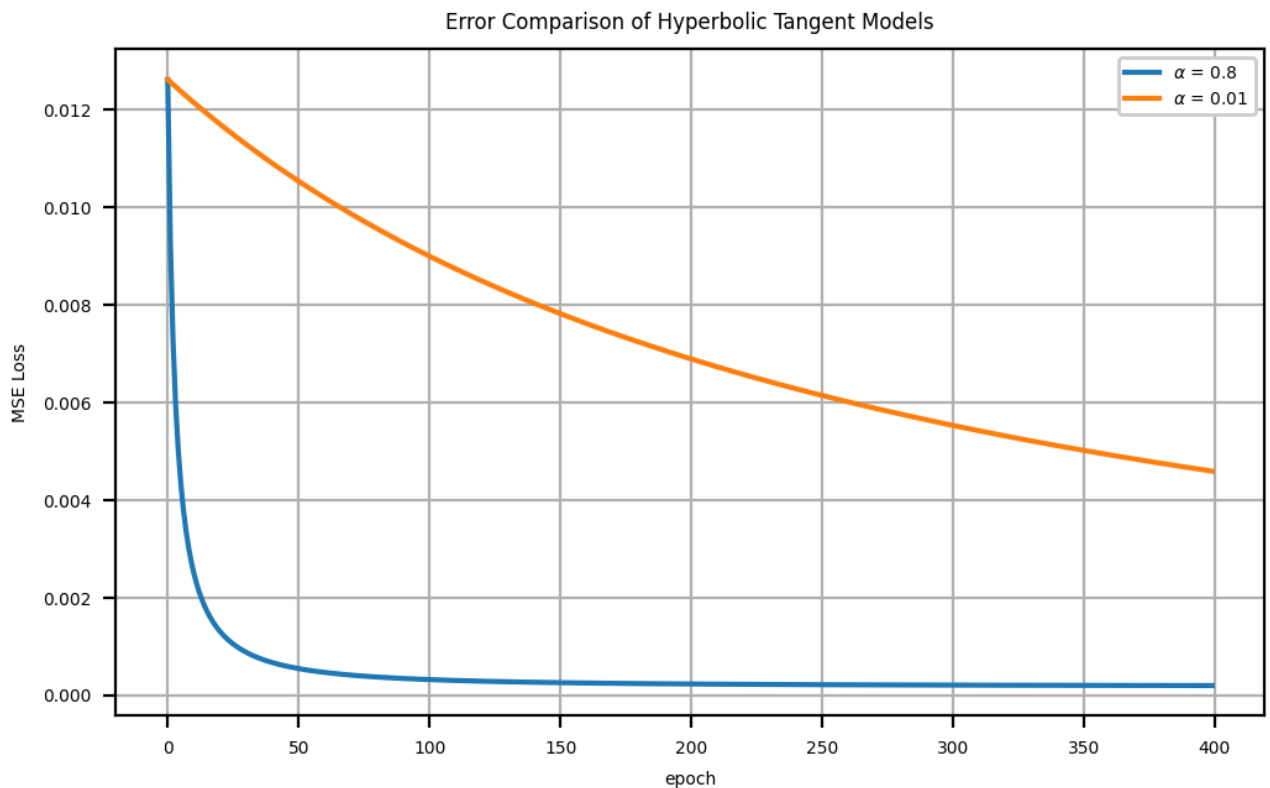
Initial Parameters: m: 0.6337605448331272 c: 0.9140319527336822

```
epoch: 0, loss: 0.0126170849, m: 0.6347, c: 0.91456
epoch: 20, loss: 0.0117031174, m: 0.65316, c: 0.9248
epoch: 40, loss: 0.0109012423, m: 0.67098, c: 0.93439
epoch: 60, loss: 0.0101926076, m: 0.68823, c: 0.94339
epoch: 80, loss: 0.0095623182, m: 0.70495, c: 0.95188
epoch: 100, loss: 0.0089984548, m: 0.72119, c: 0.9599
epoch: 120, loss: 0.0084913701, m: 0.73698, c: 0.9675
epoch: 140, loss: 0.008033175, m: 0.75237, c: 0.97472
epoch: 160, loss: 0.0076173582, m: 0.76738, c: 0.9816
epoch: 180, loss: 0.0072384996, m: 0.78203, c: 0.98816
epoch: 200, loss: 0.0068920527, m: 0.79635, c: 0.99444
epoch: 220, loss: 0.0065741761, m: 0.81035, c: 1.00044
epoch: 240, loss: 0.0062816037, m: 0.82407, c: 1.0062
epoch: 260, loss: 0.0060115409, m: 0.8375, c: 1.01174
epoch: 280, loss: 0.0057615836, m: 0.85067, c: 1.01706
epoch: 300, loss: 0.0055296526, m: 0.8636, c: 1.02218
epoch: 320, loss: 0.0053139411, m: 0.87628, c: 1.02713
epoch: 340, loss: 0.0051128719, m: 0.88874, c: 1.0319
epoch: 360, loss: 0.0049250623, m: 0.90099, c: 1.03651
epoch: 380, loss: 0.0047492951, m: 0.91303, c: 1.04097
```

Training Summary (400 epochs): loss: 0.0045924899, m: 0.92428, c: 1.04508

```
In [ ]: # obtain training history for both models
        history1 = HyperbolicModel1.history()
        history2 = HyperbolicModel2.history()

        # plotting the errors
        # plot linear model
        fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)
        ax.plot(history1[0], label = r'$\alpha$ = 0.8')
        ax.plot(history2[0], label = r'$\alpha$ = 0.01')
        ax.set_title("Error Comparison of Hyperbolic Tangent Models")
        ax.set_xlabel('epoch')
        ax.set_ylabel('MSE Loss')
        ax.legend()
        ax.legend().get_frame().set_alpha(1.0)
        ax.grid(True)
```



```
In [ ]: # obtain model parameters
model1 = history1[2:]
model2 = history2[2:]

# ground truth model
truth_model = (1.0, 2.0)

# define function for plotting
def model_plotting(model: tuple):

    # create x values
    x = np.linspace(0, 2)

    return x, np.tanh(model[0]*x+model[1])

# plot linear model
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(model_plotting(truth_model)[0],
        model_plotting(truth_model)[1],
        label = r'Ground Truth Model:  $\tanh(x+2)$ '
        )

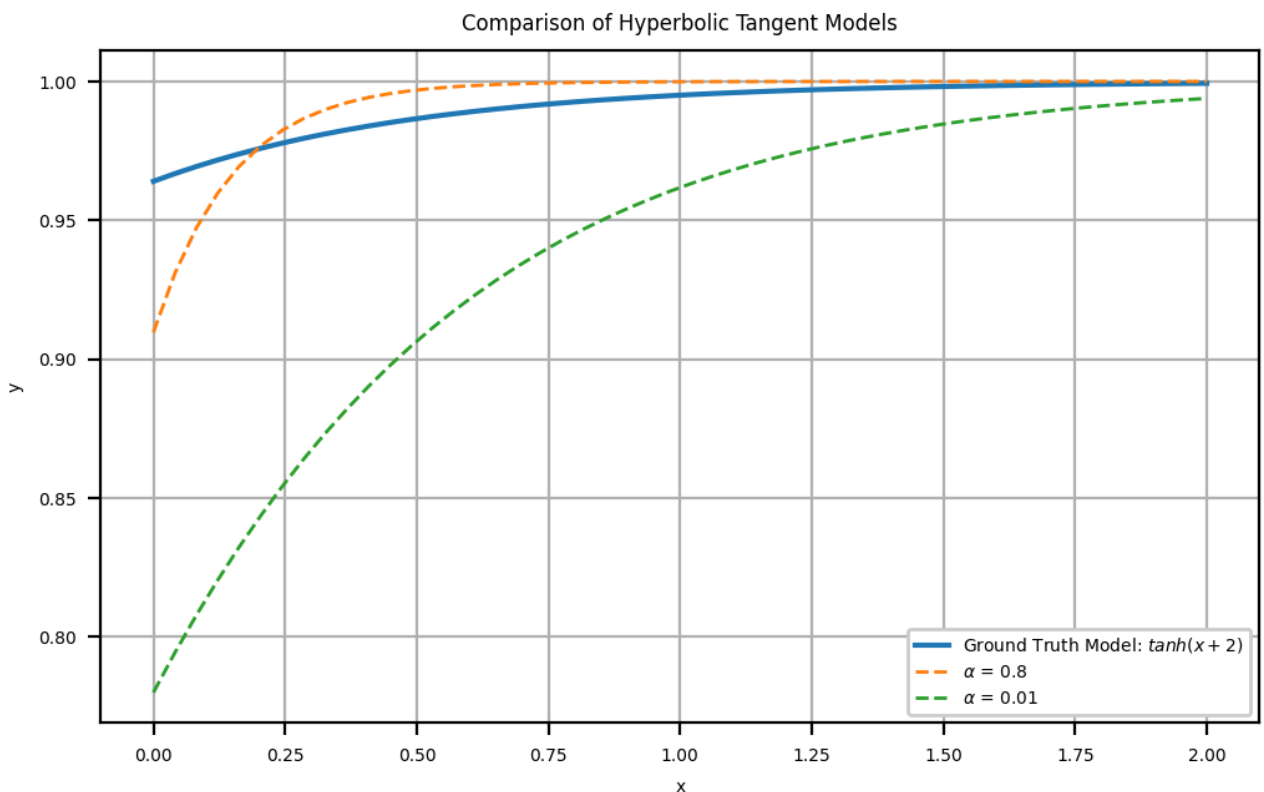
ax.plot(model_plotting(model1)[0],
        model_plotting(model1)[1],
        linestyle='dashed',
        label = r' $\alpha = 0.8$ ',
        linewidth = 1)
```

```

ax.plot(model_plotting(model2)[0],
        model_plotting(model2)[1],
        linestyle='dashed',
        label = r'$\alpha$ = 0.01',
        linewidth=1)

ax.set_title("Comparison of Hyperbolic Tangent Models")
ax.set_xlabel('x')
ax.set_ylabel('y')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)

```



3. ML Basics

Logistic (multiclass or cross-entropy) Loss

Begin by downloading the generated data from `Data_Linear_Classifier.ipynb`

```

In [ ]: import pickle


X = pickle.load(open('Misc_files/dataX.pickle', 'rb'))
y = pickle.load(open('Misc_files/dataY.pickle', 'rb'))

fig, ax = plt.subplots(figsize=(3.5, 3.5), dpi=200)

```

```
ax.scatter(X[:,0], X[:,1], c=y, s=5, cmap=plt.cm.Spectral, edgecolors="black")
ax.set_aspect('equal')
ax.set_xlim(-1, 1)
ax.set_ylim(-1, 1)
```

Out[]: (-1.0, 1.0)



In addition to the data we initialize random parameters W and b for testing

```
In [ ]: # random seed for reproducibility
np.random.seed(576)

# get Array W of shape (2, 3)
W = 0.01*np.random.randn(X.shape[1],max(y)+1)

# get Array K of shape (1, 3)
b = np.zeros((1, max(y)+1))
```

Next build the *cross-entropy* function with inputs (X, y, W, b)

```
In [ ]: def crossentropy_loss(X, y, W, b):

    # evaluate class scores
    scores = np.dot(X, W) + b

    # compute softmax of scores
    probs = np.exp(scores)/np.sum(np.exp(scores), axis=1, keepdims=True)

    # compute loss
    correct_logprobs = -np.log(probs[range(X.shape[0]),y])
    data_loss = np.sum(correct_logprobs)/X.shape[0]

    return data_loss

# test function
print("loss "+str(crossentropy_loss(X, y, W, b)))
```

loss 1.0988948999636798

l_1 and l_2 Regularization

We can revise the `crossentropy_loss` function by introducing L_1 and L_2 regularization. More formally:

$$L_1 = \sum_i \sum_j |W_{ij}|$$

$$L_2 = \sum_i \sum_j W_{ij}^2$$

By default, we will initialize function parameters `l1` and `l2` at `0.0`.

```
In [ ]: def crossentropy_loss_regularization(X, y, W, b, l1=0.0, l2=0.0):

    # evaluate class scores
    scores = np.dot(X, W) + b

    # compute softmax of scores
    probs = np.exp(scores)/np.sum(np.exp(scores), axis=1, keepdims=True)

    # compute l1 and l2 regularizers
    l1_loss = 0.5*l1*np.sum(np.abs(W))
    l2_loss = 0.5*l2*np.sum(np.square(W))

    # compute loss
    correct_logprobs = -np.log(probs[range(X.shape[0]),y])

    # return total loss
    data_loss = np.sum(correct_logprobs)/X.shape[0] - l1_loss - l2_loss
```

```

        return data_loss

# test function
print("loss "+str(crossentropy_loss_regularization(X, y, W, b, 1e-3, 1e-3)))

loss 1.0988840322643474

```

4. Classification Pipeline

Train Test Splits

We will use the [sci-kit learn](#) `train_test_split` utility to separate a new instance of the data and recreate a new set of initial (W, b) parameters.

```

In [ ]: # import utility
        from sklearn.model_selection import train_test_split

# reload X and y data from the pickle file
X = pickle.load(open('Misc_files/dataX.pickle','rb'))
y = pickle.load(open('Misc_files/dataY.pickle','rb'))

# random seed for reproducibility
np.random.seed(576)

# get Array W of shape (2, 3)
W = 0.01*np.random.randn(X.shape[1],max(y)+1)

# get Array K of shape (1, 3)
b = np.zeros((1, max(y)+1))

# create test train splits, note shuffling on creation
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, sh

```

Linear Classifier Construction

Let's build a class to instantiate Linear Classifier models. Per the instructions given in the assignment, `l1` regularization will not be taken as a parameter in the class.

```

In [ ]: class LinearClassifier:
        def __init__(self, X, y, W, b, lr, epochs, l2 = 1e-3):
            self.X = X.copy()
            self.y = y.copy()
            self.W = W.copy()
            self.b = b.copy()
            self.lr = lr
            self.epochs = epochs
            self.l2 = l2

            # metric tracking

```

```

self.loss = 0
self.loss_history = []
self.accuracy_history = []

# runtime method
def start(self, verbose = True):

    i = 0

    for i in range(self.epochs):

        # compute the gradient on scores
        num_examples = self.X.shape[0]
        dscores = np.dot(self.X, self.W) + self.b
        dscores[range(num_examples), self.y] -= 1
        dscores /= num_examples

        # backpropagate gradient to parameters (W, b)
        dW = np.dot(self.X.T, dscores)
        db = np.sum(dscores, axis=0, keepdims=True)

        dW += self.l2*self.W # l2 regularization

        # parameter update
        self.W += -self.lr*dW
        self.b += -self.lr*db

        # compute loss, l2 regularization only, and training accuracy
        self.loss = crossentropy_loss_regularization(self.X, self.y, self.l2)
        training_accuracy = self.eval(self.X, self.y)

        #metrics
        self.loss_history.append(self.loss)
        self.accuracy_history.append(training_accuracy)

        # output tracking
        if i % 10 == 0 and verbose == True :
            print(f"iteration {i}: loss: {self.loss} training_accuracy: {training_accuracy}")

        i += 1

    # exit output
    if verbose == True:
        print(f"iteration {i}: loss: {self.loss} training_accuracy: {training_accuracy}")

# get W
def get_W(self):
    return self.W

# get b
def get_b(self):
    return self.b

```

```

# get params as tuple
def get_parameters(self):
    return self.W, self.b

# evaluate model
def eval(self, X = X, y = y):
    scores = np.dot(X, self.W) + self.b
    pred = np.argmax(scores, axis=1)
    return np.mean(pred == y)

# plot losses
def show_loss(self):
    fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)
    ax.plot(self.loss_history)
    ax.set_title("Linear Classifier Loss")
    ax.set_xlabel('Epoch')
    ax.set_ylabel('Loss')
    ax.grid(True)

# plot training accuracy
def show_accuracy(self):
    fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)
    ax.plot(self.accuracy_history)
    ax.set_title("Linear Classifier Training Accuracy")
    ax.set_xlabel('Epoch')
    ax.set_ylabel('Training Accuracy')
    ax.grid(True)

# post training
def show_classifier(self):
    h = 0.02
    x_min, x_max = self.X[:, 0].min() - 1, self.X[:, 0].max() + 1
    y_min, y_max = self.X[:, 1].min() - 1, self.X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))
    Z = np.dot(np.c_[xx.ravel(), yy.ravel()], self.W) + self.b
    Z = np.argmax(Z, axis=1)
    Z = Z.reshape(xx.shape)
    fig = plt.figure()
    plt.contourf(xx, yy, Z, cmap=plt.cm.Spectral, alpha=0.8)
    plt.scatter(self.X[:, 0], self.X[:, 1], c=self.y, s=40, cmap=plt.cm.
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())

# model summary
def summary(self):
    print(f"Model summary:"
          +f"\n\nW:\n{self.W}"
          +f"\n\nB:\n{self.b}"
          +f"\n\nLearning Rate: {self.lr}"
          +f"\nl2 Regularization: {self.l2}"
          +f"\nEpochs: {self.epochs}"

```



```
+f"\n\nTraining Accuracy: {self.accuracy_history[-1]}"
+f"\nModel Loss: {self.loss_history[-1]}"
```

We next instantiate a `LinearClassifier` model and train it on `X_train`, `y_train`, and initial parameters `W` and `b` using the class method `start()`.

```
In [ ]: # evaluate on training data
Classifier = LinearClassifier(X_train, y_train, W, b, 1e-3, 200, 1e-3)

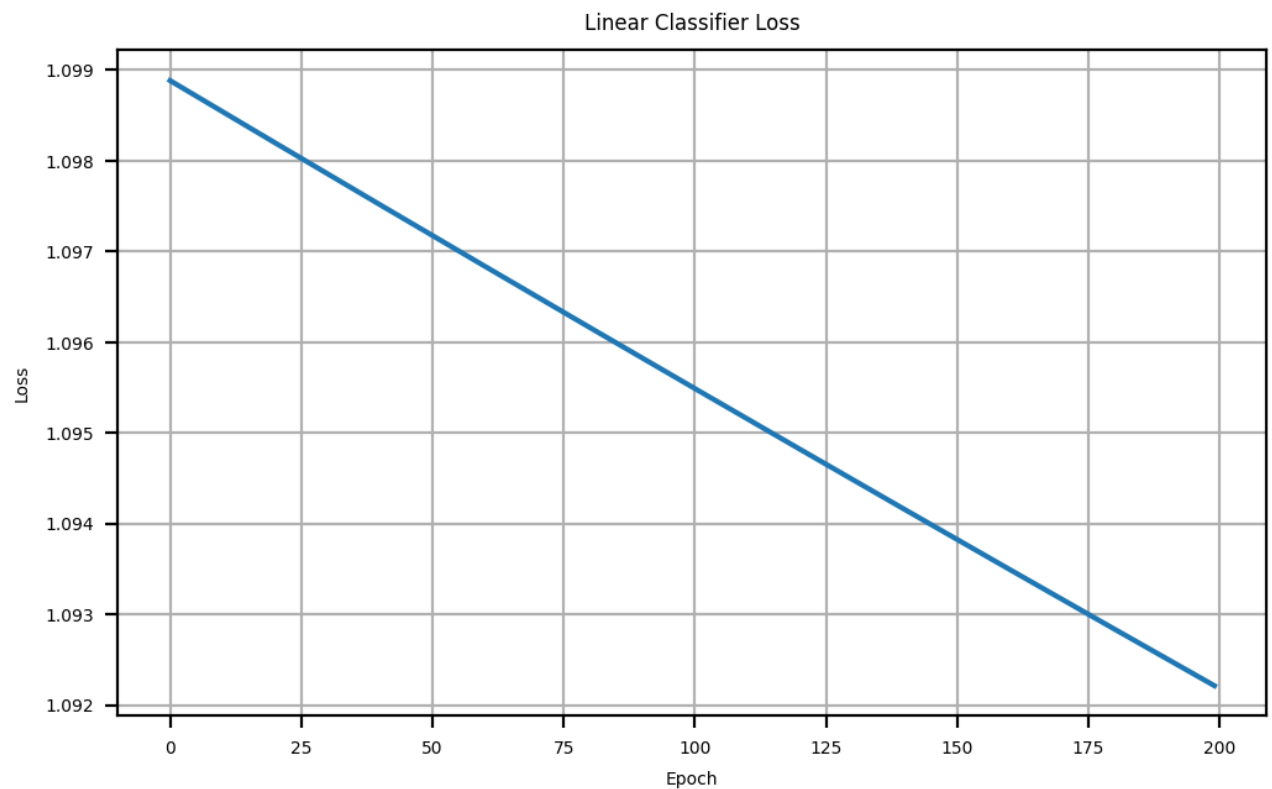

# start training
Classifier.start()
```

```
iteration 0: loss: 1.098884320032409 training_accuracy: 0.3541666666666667
iteration 10: loss: 1.0985408705616455 training_accuracy: 0.4
iteration 20: loss: 1.0981983208685286 training_accuracy: 0.4333333333333333
5
iteration 30: loss: 1.0978566669229148 training_accuracy: 0.4666666666666667
iteration 40: loss: 1.0975159047341516 training_accuracy: 0.525
iteration 50: loss: 1.0971760303504645 training_accuracy: 0.5625
iteration 60: loss: 1.0968370398583502 training_accuracy: 0.5833333333333334
iteration 70: loss: 1.0964989293819818 training_accuracy: 0.6
iteration 80: loss: 1.0961616950826265 training_accuracy: 0.6125
iteration 90: loss: 1.095825333158072 training_accuracy: 0.6166666666666667
iteration 100: loss: 1.095489839842064 training_accuracy: 0.6208333333333333
iteration 110: loss: 1.0951552114037537 training_accuracy: 0.625
iteration 120: loss: 1.0948214441471544 training_accuracy: 0.6208333333333333
3
iteration 130: loss: 1.0944885344106112 training_accuracy: 0.6208333333333333
3
iteration 140: loss: 1.0941564785662745 training_accuracy: 0.6166666666666666
7
iteration 150: loss: 1.09382527301959 training_accuracy: 0.6166666666666667
iteration 160: loss: 1.09349491420879 training_accuracy: 0.6166666666666667
iteration 170: loss: 1.0931653986044014 training_accuracy: 0.6166666666666666
7
iteration 180: loss: 1.0928367227087583 training_accuracy: 0.6208333333333333
3
iteration 190: loss: 1.0925088830555227 training_accuracy: 0.6166666666666666
7
iteration 200: loss: 1.0922145395148164 training_accuracy: 0.6125
```

Calling the methods `show_accuracy()` and `show_loss()` the model's training history can now be shown:

```
In [ ]: # show training accuracy history
Classifier.show_accuracy()

# show loss history
Classifier.show_loss()
```



We now call the `summary` method to show the model summary and a `show_classifier()` method to plot the resulting classifier's decision boundary.

```
In [ ]: Classifier.summary()  
Classifier.show_classifier()
```

Model summary:

W:

```
[[ 0.01463343  0.01229308 -0.01946041]
 [-0.01563398  0.01509063 -0.00507448]]
```

B:

```
[[0.0567085  0.06500469 0.05966792]]
```


Learning Rate: 0.001

l2 Regularization: 0.001

Epochs: 200

Training Accuracy: 0.6125

Model Loss: 1.0922145395148164



Calling the `eval()` method will evaluate the model on the train and test data.

```
In [ ]: # get training accuracy
print(f"training accuracy: {Classifier.eval(X_train, y_train)}")

# get testing accuracy
print(f"test accuracy: {Classifier.eval(X_test, y_test)}")
```

training accuracy: 0.6125

test accuracy: 0.45

A low test accuracy in comparison to the training accuracy typically indicates overfitting, but the radially symmetric distribution of the example data suggests that a linear

classifier is not the appropriate choice for this data.

Cross Validation

For the relatively small number of instances (240) in the training set, we decide to use `sklearn`'s `KFold` class for cross validation. Cross validation consists of 10 folds.

```
In [ ]: from sklearn.model_selection import KFold

# set naive parameters for KFold cross validation
lr = 1e-3
epochs = 300
l2 = 1e-3

# tracking objects
training_history = []
validation_history = []
loss_history = []

# best model tracking
best_accuracy = 0.0

# instantiate KFold object
kf = KFold(n_splits = 10)

# begin cv loop
for i, (train_index, validation_index) in enumerate(kf.split(X_train)):

    # instantiate new LinearClassifier object with cv split data
    cv_classifier = LinearClassifier(X_train[train_index], y_train[train_index])
    cv_classifier.start(verbose = False)

    # obtain tracking metrics
    training_accuracy = cv_classifier.eval(X_train[train_index], y_train[train_index])
    validation_accuracy = cv_classifier.eval(X_train[validation_index], y_train[validation_index])
    loss = cv_classifier.loss

    # update best model
    if training_accuracy > best_accuracy:
        best_accuracy = training_accuracy
        best_classifier = cv_classifier

    # append metrics
    training_history.append(training_accuracy)
    validation_history.append(validation_accuracy)
    loss_history.append(loss)

# return mean values
mean_cv_training = np.mean(training_history)
mean_cv_validation = np.mean(validation_history)
mean_cv_loss = np.mean(loss_history)
```

```
# output
print('Leave One Out Cross-validation results\n'
      f'Mean Training Accuracy: {mean_cv_training}\n'
      f'Mean Validation Accuracy: {mean_cv_validation}\n'
      f'Mean Loss: {mean_cv_loss}')
```

```
Leave One Out Cross-validation results
Mean Training Accuracy: 0.5842592592592593
Mean Validation Accuracy: 0.5666666666666667
Mean Loss: 1.088892209841891
```

```
In [ ]: print("Best Model:")
        best_classifier.summary()

        best_classifier.show_classifier()
```


```
Best Model:
Model summary:
```

```
W:
[[ 0.02072945  0.01775439 -0.02730493]
 [-0.02166191  0.02486712 -0.00782442]]
```

```
B:
[[0.07922916 0.09721065 0.08285023]]
```

```
Learning Rate: 0.001
l2 Regularization: 0.001
Epochs: 300
```

```
Training Accuracy: 0.6203703703703703
Model Loss: 1.0890753164637825
```



Performance Sensitivity: Learning Rate and Gradient Descent

To evaluate model performance sensitivity on various hyperparameters, we modify the above block of code to perform KFold cross validation on models of varying learning rate and epoch size.

Learning Rate Sensitivity

To test the sensitivity of the Linear Classifier, we iteratively test learning rates between `1e-1` to `1e-5` with each instance evaluated through LOO cross-validation.

```
In [ ]: # initialize parameters for the linear classifier iterations
learning_rates = np.geomspace(1e-1, 1e-5, 5)
epochs = 300
l2 = 1e-3

# instantiate KFold object
kf = KFold(n_splits = 10, shuffle = True, random_state=0)

# initialize tracking objects for each linear classifier
model_training_accraccies = []
model_validation_accraccies = []
model_loss = []

# best model tracking
best_accuracy = 0.0
```

```

# run first loop to iterate through learning rates
for lr in learning_rates:
    # tracking objects
    training_history = []
    validation_history = []
    loss_history = []

    # begin CV loop
    for i, (train_index, validation_index) in enumerate(kf.split(X_train)):

        # instantiate new LinearClassifier object with cv split data
        cv_classifier = LinearClassifier(X_train[train_index], y_train[train_index])
        cv_classifier.start(verbose = False)

        # obtain tracking metrics
        training_accuracy = cv_classifier.eval(X_train[train_index], y_train[train_index])
        validation_accuracy = cv_classifier.eval(X_train[validation_index], y_train[validation_index])
        loss = cv_classifier.loss

        # update best model
        if training_accuracy > best_accuracy:
            best_accuracy = training_accuracy
            best_classifier = cv_classifier

        # append metrics
        training_history.append(training_accuracy)
        validation_history.append(validation_accuracy)
        loss_history.append(loss)

    # return mean values
    mean_cv_training = np.mean(training_history)
    mean_cv_validation = np.mean(validation_history)
    mean_cv_loss = np.mean(loss_history)

    # append mean metrics to overall tracking
    model_training_accraccies.append(mean_cv_training)
    model_validation_accraccies.append(mean_cv_validation)
    model_loss.append(mean_cv_loss)

```

Show best model:

```

In [ ]: print("Best Model:")
        best_classifier.summary()

        best_classifier.show_classifier()

```


Best Model:
Model summary:

W:
[[0.02072945 0.01775439 -0.02730493]
[-0.02166191 0.02486712 -0.00782442]]

B:
[[0.07922916 0.09721065 0.08285023]]

Learning Rate: 0.001
l2 Regularization: 0.001
Epochs: 300

Training Accuracy: 0.6203703703703703
Model Loss: 1.0890753164637825



We then plot the pertinent metrics for sensitivity analysis.

```
In [ ]: # plot learning rate sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(learning_rates,
        model_training_accuaccies,
        label = r'Mean training training',
        marker = 'o',
        markersize = 4)

ax.plot(learning_rates,
```



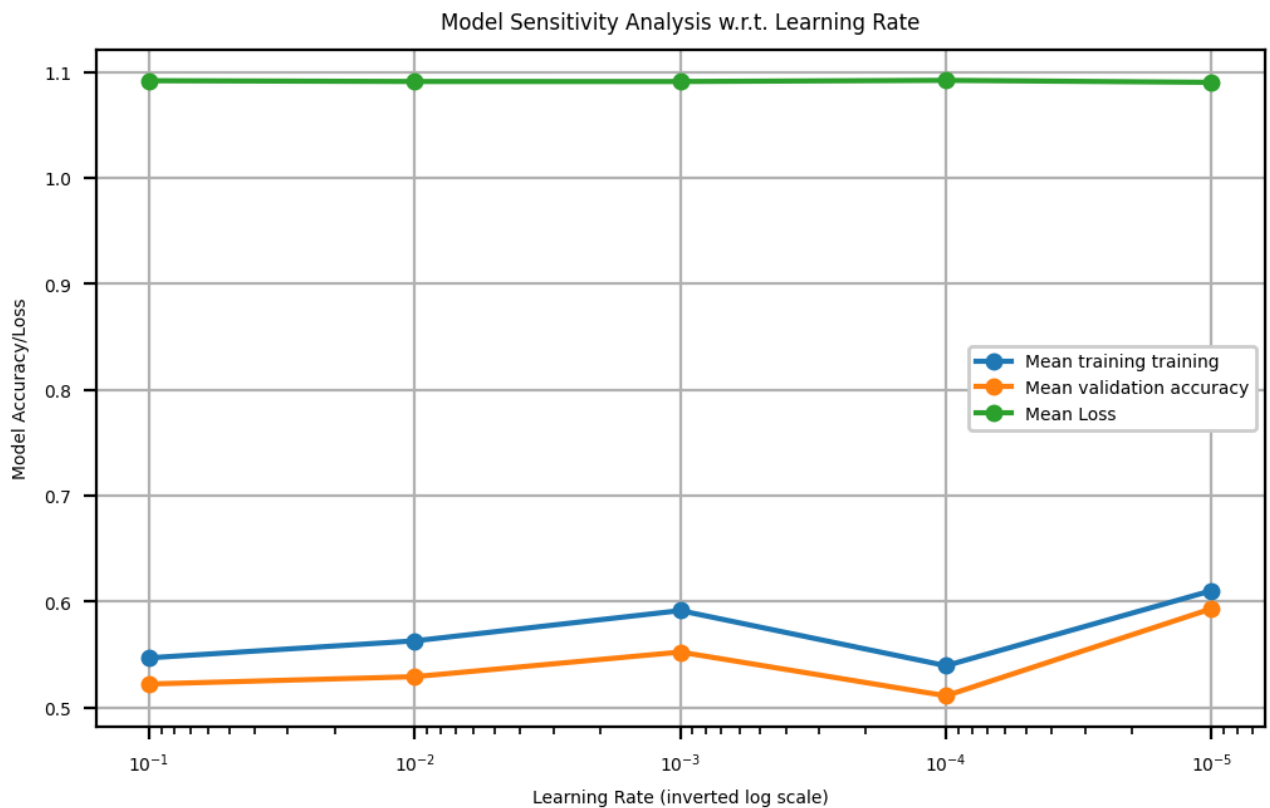
```

model_validation_accuraccies,
label = r'Mean validation accuracy',
marker = 'o',
markersize = 4)

ax.plot(learning_rates,
model_loss,
label = r'Mean Loss',
marker = 'o',
markersize = 4)

ax.set_title("Model Sensitivity Analysis w.r.t. Learning Rate")
ax.set_xlabel('Learning Rate (inverted log scale)')
ax.set_xscale("log")
ax.invert_xaxis()
ax.set_ylabel('Model Accuracy/Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)

```



The above plot illustrates the sensitivity of our Linear Classifier Model w.r.t. different learning rates, evaluated on the `X_train` training data and cross validated with 10 splits. A learning rate of `1e-3` yielded the best training and validation accuracy and subsequent smaller learning rates degraded overall model accuracy.

Gradient Descent Iteration Sensitivity

Similarly, we perform the same task for the number of gradient descent iterations. The

previous best learning rate `1e-3` remains constant in lieu of other hyperparameter tuning schemas such as *grid search* or *random search*.

```
In [ ]: # initialize parameters for the linear classifier iterations
lr = 1e-3
various_epochs = np.linspace(100, 1000, 7, dtype=int)
l2 = 1e-3

# instantiate LeaveOneOut object for CV
kf = KFold(n_splits = 10, shuffle = True, random_state=1)

# initialize tracking objects for each linear classifier
model_training_accuraccies = []
model_validation_accuraccies = []
model_loss = []

# best model tracking
best_accuracy = 0.0

# run first loop to iterate through learning rates
for epochs in various_epochs:
    # tracking objects
    training_history = []
    validation_history = []
    loss_history = []

    # begin CV loop
    for i, (train_index, validation_index) in enumerate(kf.split(X_train)):

        # instatiate new LinearClassifier object with cv split data
        cv_classifier = LinearClassifier(X_train[train_index], y_train[train_index])
        cv_classifier.start(verbose = False)

        # obtain tracking metrics
        training_accuracy = cv_classifier.eval(X_train[train_index], y_train[train_index])
        validation_accuracy = cv_classifier.eval(X_train[validation_index], y_train[validation_index])
        loss = cv_classifier.loss

        # update best model
        if training_accuracy > best_accuracy:
            best_accuracy = training_accuracy
            best_classifier = cv_classifier

        # append metrics
        training_history.append(training_accuracy)
        validation_history.append(validation_accuracy)
        loss_history.append(loss)

    # return mean values
    mean_cv_training = np.mean(training_history)
    mean_cv_validation = np.mean(validation_history)
    mean_cv_loss = np.mean(loss_history)
```

```
# append mean metrics to overall tracking
model_training_accuraccies.append(mean_cv_training)
model_validation_accuraccies.append(mean_cv_validation)
model_loss.append(mean_cv_loss)
```

```
In [ ]: print("Best Model:")
        best_classifier.summary()

        best_classifier.show_classifier()
```


Best Model:
Model summary:

W:
[[0.01054313 0.00635041 -0.00866427]
 [-0.00760066 0.00546519 -0.0042355]]

B:
[[0.02999379 0.03483221 0.03040347]]

Learning Rate: 0.001
l2 Regularization: 0.001
Epochs: 100

Training Accuracy: 0.6435185185185185
Model Loss: 1.0955903096970332

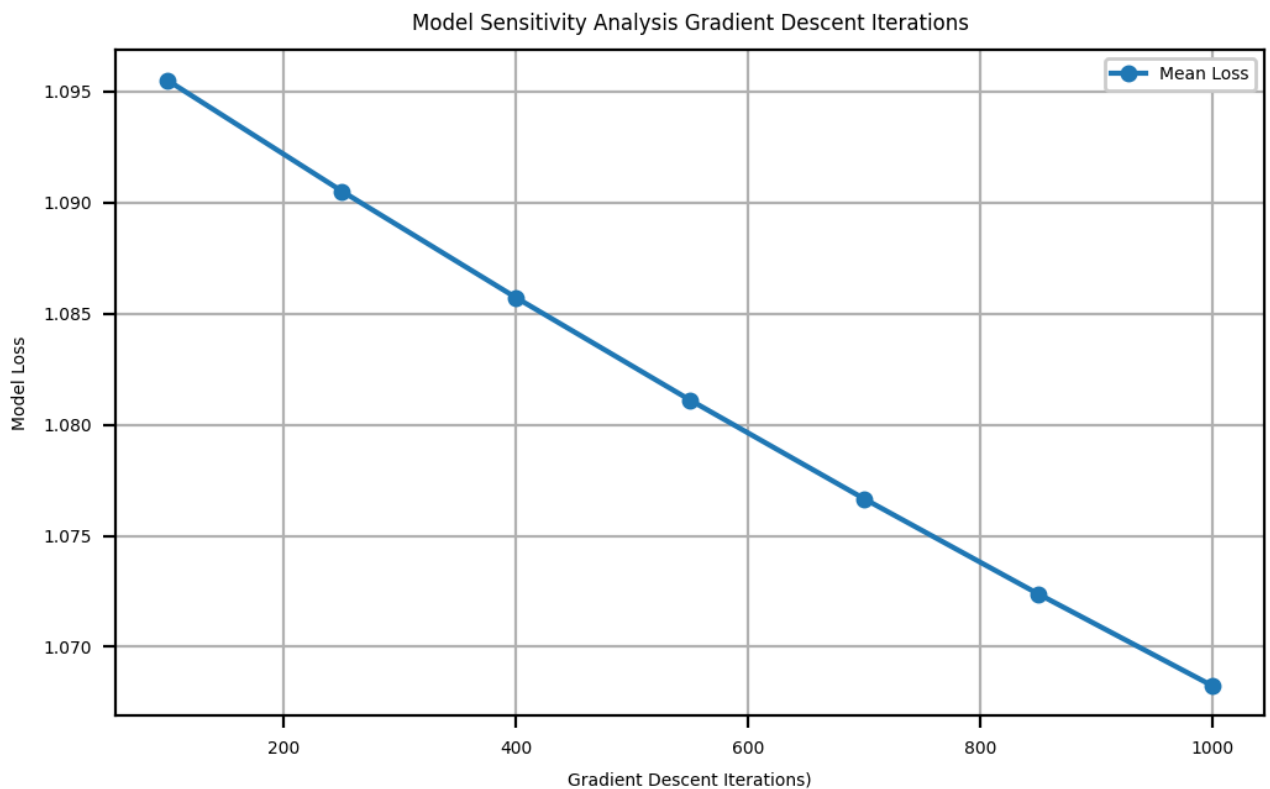


We next plot training/validation accuracy and model loss on seperate plots.

```
In [ ]: # plot learning rate sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(various_epochs,
        model_loss,
        label = r'Mean Loss',
        marker = 'o',
        markersize = 4)

ax.set_title("Model Sensitivity Analysis Gradient Descent Iterations")
ax.set_xlabel('Gradient Descent Iterations')
ax.set_ylabel('Model Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)
```



```
In [ ]: # plot learning rate sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(various_epochs,
        model_training_accuraccies,
        label = r'Mean training training',
        marker = 'o',
        markersize = 4)

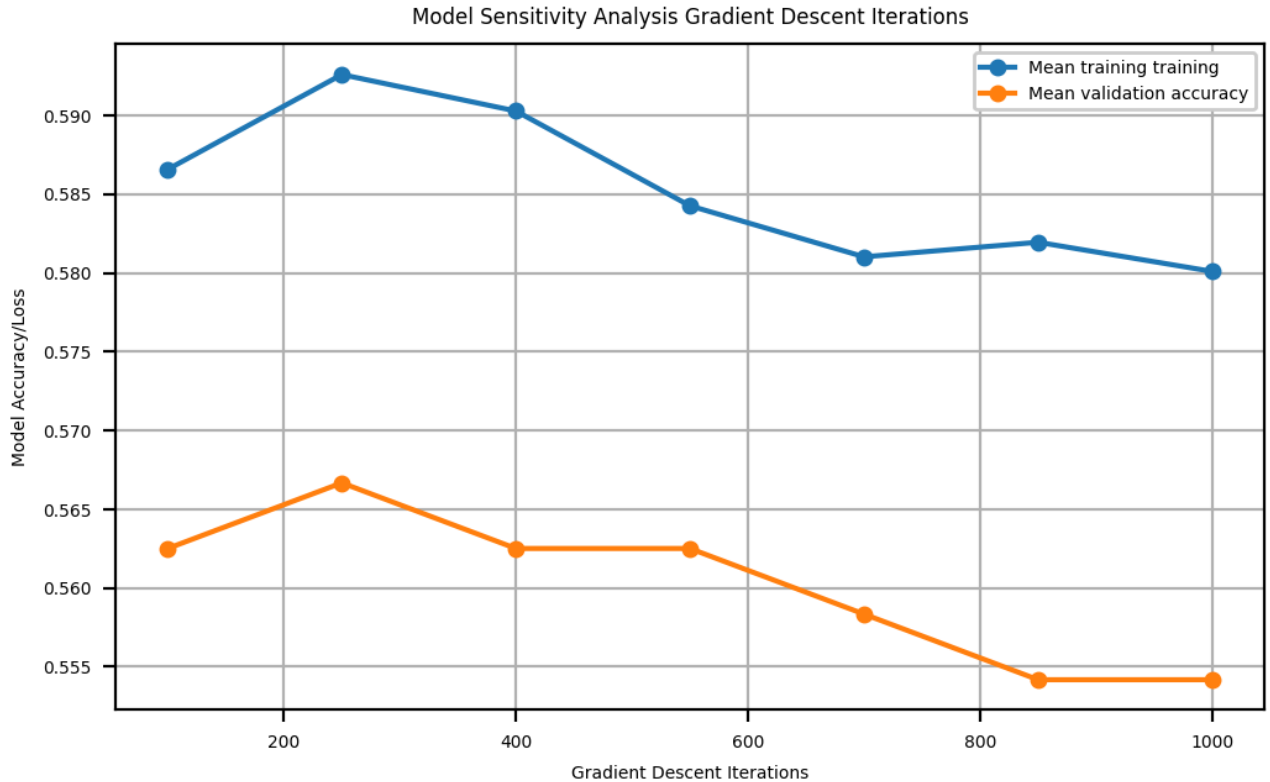
ax.plot(various_epochs,
        model_validation_accuraccies,
        label = r'Mean validation accuracy',
        marker = 'o',
```

```

markersize = 4)

ax.set_title("Model Sensitivity Analysis Gradient Descent Iterations")
ax.set_xlabel('Gradient Descent Iterations')
ax.set_ylabel('Model Accuracy/Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)

```



Change in model sensitivity w.r.t. gradient descent iterations shows an increase in mean validation accuracy up to 250 iterations. Past this point validation and training accuracy appears to degrade. For this reason the number of iterations `250` is the selected to be the best candidate, offering the best training and validation accuracy without unnecessary computational cost and loss of accuracy.

12 Regularization Sensitivity

We will retain the best hyperparameters for the linear model `lr = 1e-3` and `epochs = 250` for this next sensitivity analysis, while testing the effects of different `l2` values including `0.0` (no `l2` regularization).

```

In [ ]: # initialize parameters for the linear classifier iterations
lr = 1e-3
epochs = 250
l2_values = np.append(np.geomspace(1e-1, 1e-5, 5), [0], axis = 0)

# instantiate LeaveOneOut object for CV

```

```

kf = KFold(n_splits = 10, shuffle = True, random_state=2)

# initialize tracking objects for each linear classifier
model_training_accuraccies = []
model_validation_accuraccies = []
model_loss = []

# best model tracking
best_accuracy = 0.0

# run first loop to iterate through learning rates
for l2 in l2_values:
    # tracking objects
    training_history = []
    validation_history = []
    loss_history = []

    # begin CV loop
    for i, (train_index, validation_index) in enumerate(kf.split(X_train)):

        # instatiate new LinearClassifier object with cv split data
        cv_classifier = LinearClassifier(X_train[train_index], y_train[train_index])
        cv_classifier.start(verbose = False)

        # obtain tracking metrics
        training_accuracy = cv_classifier.eval(X_train[train_index], y_train[train_index])
        validation_accuracy = cv_classifier.eval(X_train[validation_index], y_train[validation_index])
        loss = cv_classifier.loss

        # update best model
        if training_accuracy > best_accuracy:
            best_accuracy = training_accuracy
            best_classifier = cv_classifier

        # append metrics
        training_history.append(training_accuracy)
        validation_history.append(validation_accuracy)
        loss_history.append(loss)

    # return mean values
    mean_cv_training = np.mean(training_history)
    mean_cv_validation = np.mean(validation_history)
    mean_cv_loss = np.mean(loss_history)

    # append mean metrics to overall tracking
    model_training_accuraccies.append(mean_cv_training)
    model_validation_accuraccies.append(mean_cv_validation)
    model_loss.append(mean_cv_loss)

```

```

In [ ]: print("Best Model:")
        best_classifier.summary()

        best_classifier.show_classifier()

```

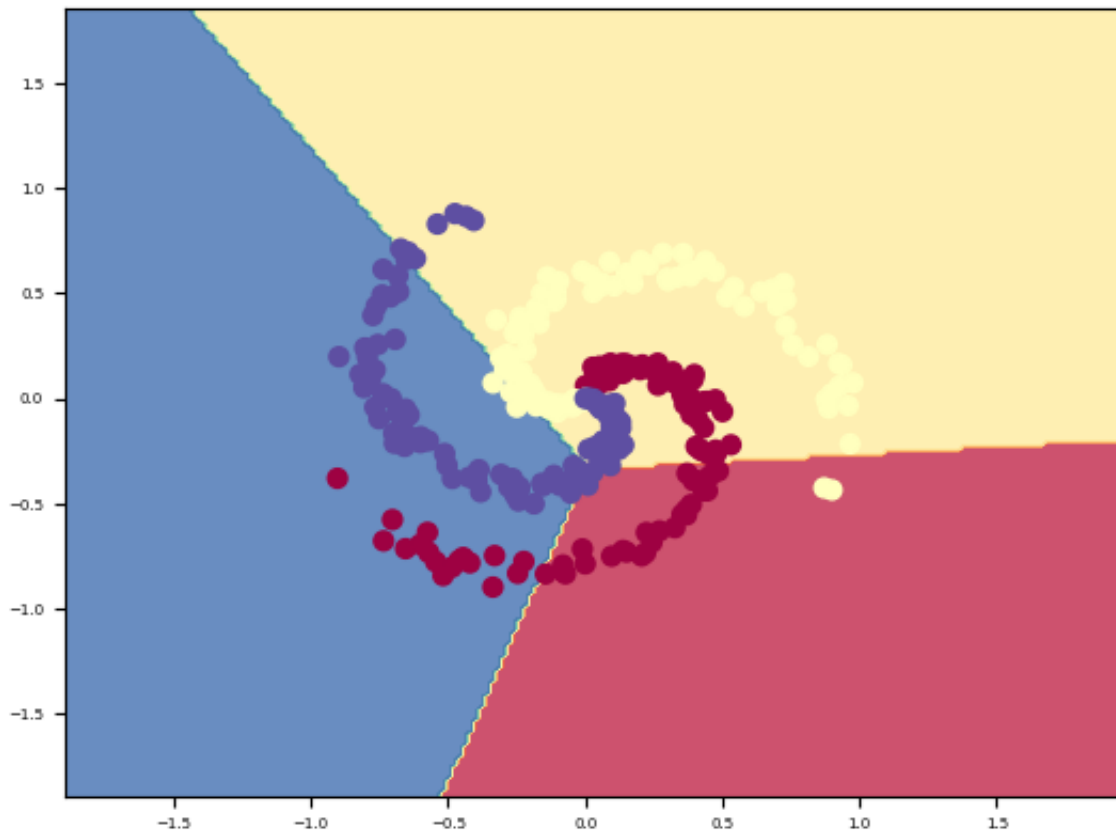
Best Model:
Model summary:

W:
[[0.01741403 0.01463414 -0.02377698]
 [-0.01890681 0.02058674 -0.00483719]]

B:
[[0.0676665 0.08093004 0.07273069]]

Learning Rate: 0.001
l2 Regularization: 0.1
Epochs: 250

Training Accuracy: 0.625
Model Loss: 1.090378951582681



```
In [ ]: # plot learning rate sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

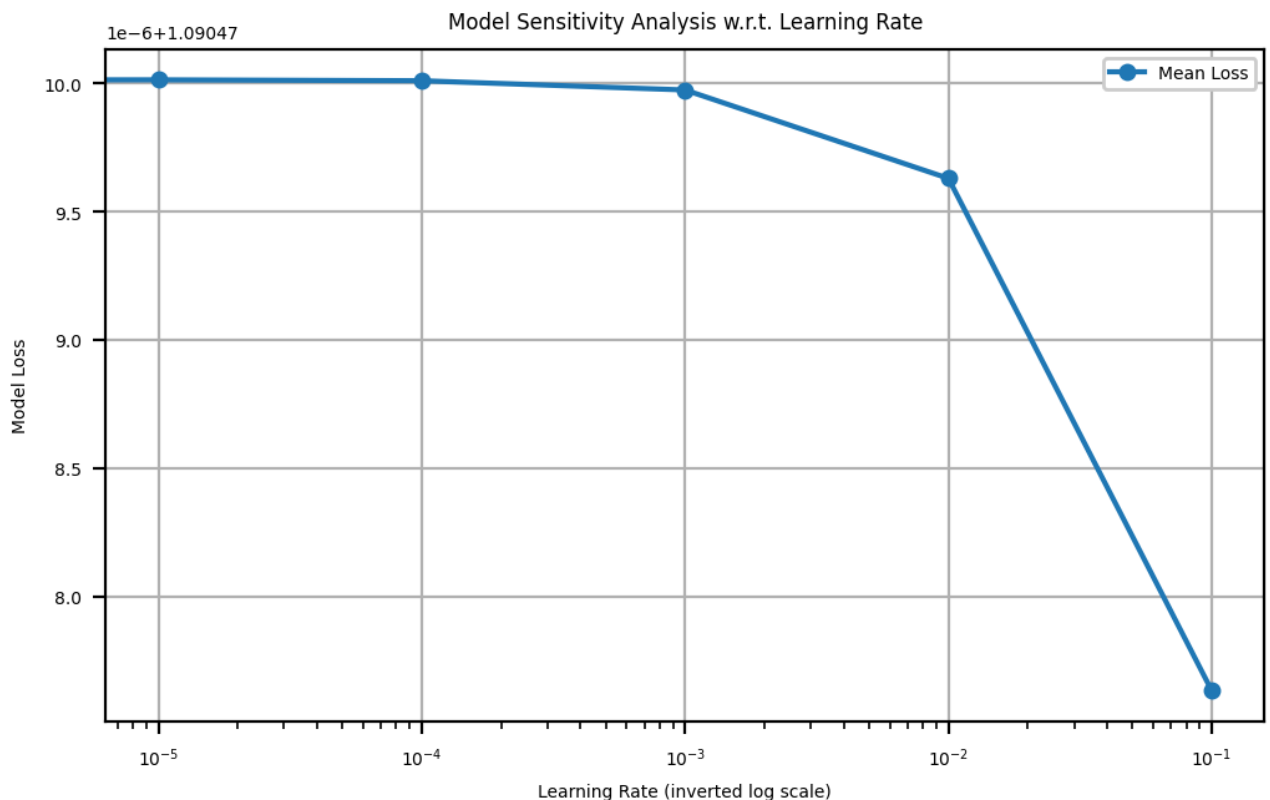
ax.plot(l2_values,
        model_loss,
        label = r'Mean Loss',
        marker = 'o',
        markersize = 4)

ax.set_title("Model Sensitivity Analysis w.r.t. Learning Rate")
ax.set_xlabel('Learning Rate (inverted log scale)')
ax.set_xscale("log")
#ax.invert_xaxis()
```

```

ax.set_ylabel('Model Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)

```



```

In [ ]: # plot learning rate sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(l2_values,
        model_training_accuraccies,
        label = r'Mean training training',
        marker = 'o',
        markersize = 4)

ax.plot(l2_values,
        model_validation_accuraccies,
        label = r'Mean validation accuracy',
        marker = 'o',
        markersize = 4)

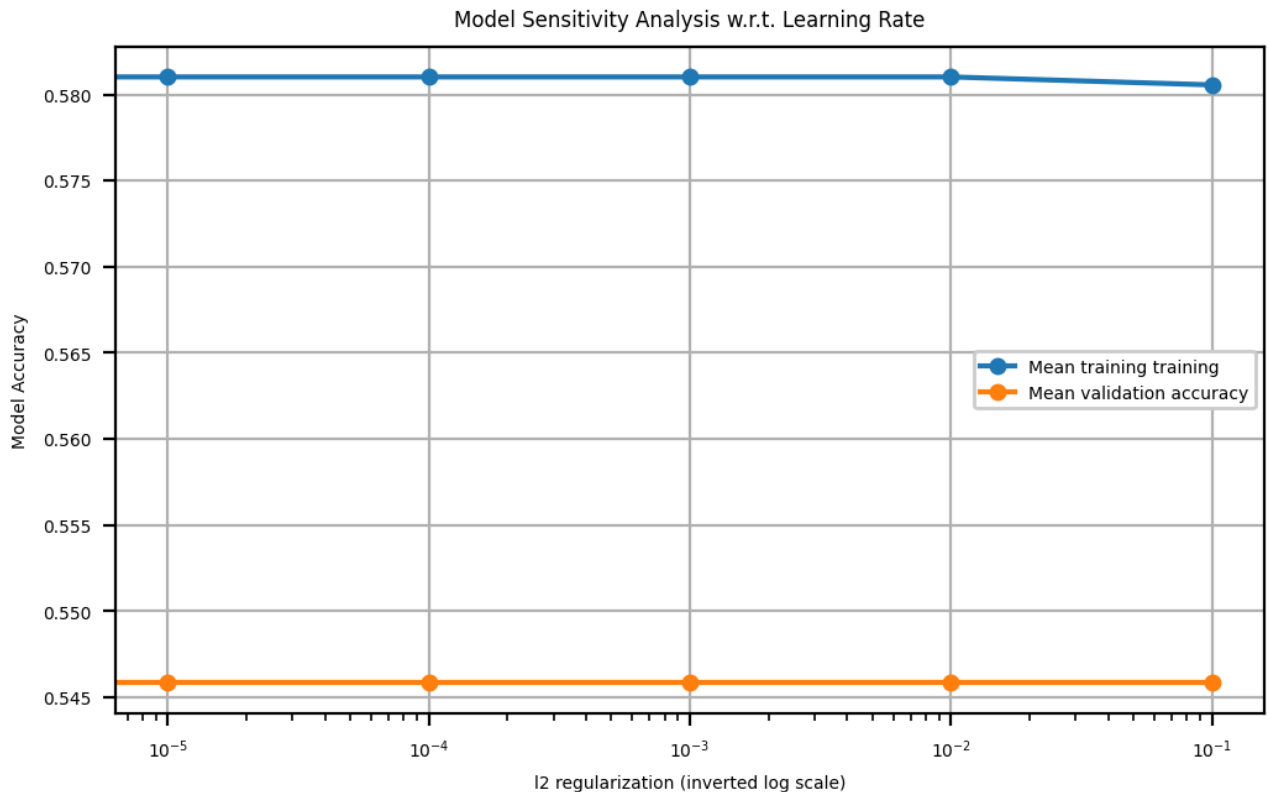
...
ax.plot(l2_values,
        model_loss,
        label = r'Mean Loss')
...

ax.set_title("Model Sensitivity Analysis w.r.t. Learning Rate")
ax.set_xlabel('l2 regularization (inverted log scale)')
ax.set_xscale("log")
#ax.invert_xaxis()

```



```
ax.set_ylabel('Model Accuracy')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)
ax.grid(True)
```



Varying `l2` regularization parameters appears to have minimal effects on the linear classifier model. For values of $l2 \in [0, 10^{-5}]$ there was no change in mean training and validation accuracy. For this reason, `l2` is selected to be `0.0` as an input parameter in order to minimize computational cost.

Train Test Split Sensitivity

The above sensitivity analysis returned the following optimal hyperparameters for the linear classifier model:

- Learning Rate: `lr = 1e-3`
- Gradient Descent Iterations: `epochs = 250`
- Regularization: `l2 = 0.0`

We then use these hyperparameters to analyze the sensitivity of the model w.r.t. train/test splits.

```
In [ ]: # specify train test splits in a list and iterate through splits
        splits = [0.1, 0.2, 0.3, 0.4, 0.5]

        # initialize best hyperparameters for linear classifier
```

```

# l2 regularization = 0.0 default for class constructor
lr = 1e-3
epochs = 250

# reload X and y data from the pickle file
X = pickle.load(open('Misc_files/dataX.pickle', 'rb'))
y = pickle.load(open('Misc_files/dataY.pickle', 'rb'))

# random seed for reproducibility
np.random.seed(576)

# get Array W of shape (2, 3)
W = 0.01*np.random.randn(X.shape[1], max(y)+1)

# get Array K of shape (1, 3)
b = np.zeros((1, max(y)+1))

# initialize tracking objects for each linear classifier
model_training_accuraccies = []
model_validation_accuraccies = []
model_testing_accuracies = []
model_loss = []

# best model tracking
best_accuracy = 0.0

# loop through splits using 10 fold cross-validation
for split in splits:

    # generate new train test split for each iteration, adjust random state
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=spli

    # tracking objects
    training_history = []
    validation_history = []
    test_history = []
    loss_history = []

    # begin CV loop
    for i, (train_index, validation_index) in enumerate(kf.split(X_train)):

        # instantiate new LinearClassifier object with cv split data
        cv_classifier = LinearClassifier(X_train[train_index], y_train[train_index])
        cv_classifier.start(verbose = False)

        # obtain tracking metrics
        training_accuracy = cv_classifier.eval(X_train[train_index], y_train[train_index])
        validation_accuracy = cv_classifier.eval(X_train[validation_index], y_train[validation_index])
        testing_accuracy = cv_classifier.eval(X_test, y_test)
        loss = cv_classifier.loss

        # update best model
        if training_accuracy > best_accuracy:

```

```

        best_accuracy = training_accuracy
        best_classifier = cv_classifier

    # append metrics
    training_history.append(training_accuracy)
    validation_history.append(validation_accuracy)
    test_history.append(testing_accuracy)
    loss_history.append(loss)

    # return mean values
    mean_cv_training = np.mean(training_history)
    mean_cv_validation = np.mean(validation_history)
    mean_cv_testing = np.mean(test_history)
    mean_cv_loss = np.mean(loss_history)

    # append mean metrics to overall tracking
    model_training_accuraccies.append(mean_cv_training)
    model_validation_accuraccies.append(mean_cv_validation)
    model_testing_accuracies.append(mean_cv_testing)
    model_loss.append(mean_cv_loss)

```

```

In [ ]: print("Best Model:")
        best_classifier.summary()

        best_classifier.show_classifier()

```

Best Model:

Model summary:

W:

```

[[ 0.01880898  0.01470664 -0.02396573]
 [-0.02503941  0.01722767 -0.00653695]]

```

B:

```

[[0.081869   0.07055378  0.06879078]]

```


Learning Rate: 0.001

l2 Regularization: 0.001

Epochs: 250

Training Accuracy: 0.6296296296296297

Model Loss: 1.0898555458957988

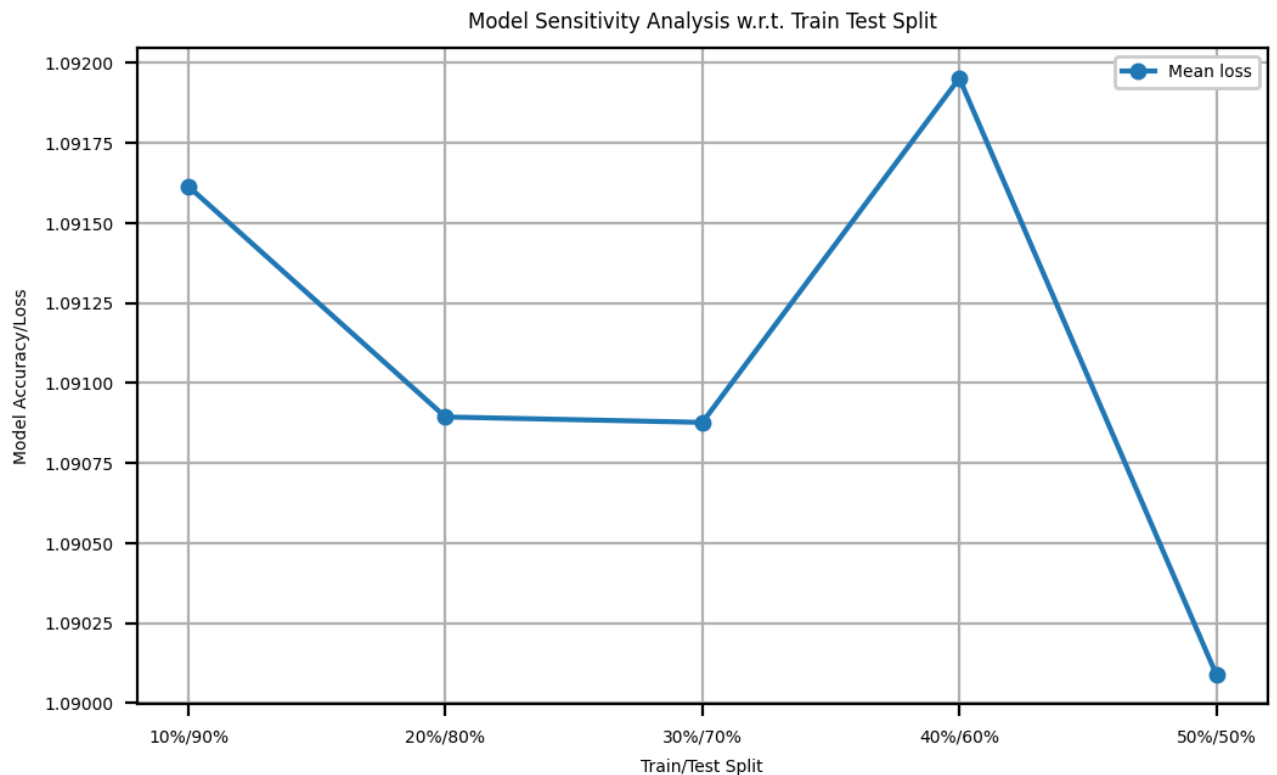


```
In [ ]: # plot train/test split sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(splits,
        model_loss,
        label = r'Mean loss',
        marker = 'o',
        markersize = 4)

ax.set_title("Model Sensitivity Analysis w.r.t. Train Test Split")
ax.set_xlabel('Train/Test Split')
ax.set_ylabel('Model Accuracy/Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)

ax.set_xticks(splits) # Set x-axis tick positions
ax.set_xticklabels([f'{int(split * 100)}%/{100-int(split * 100)}%' for split in splits])
ax.grid(True)
```



```
In [ ]: # plot train/test split sensitivity
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

ax.plot(splits,
        model_training_accuraccies,
        label = r'Mean training training',
        marker = 'o',
        markersize = 4)

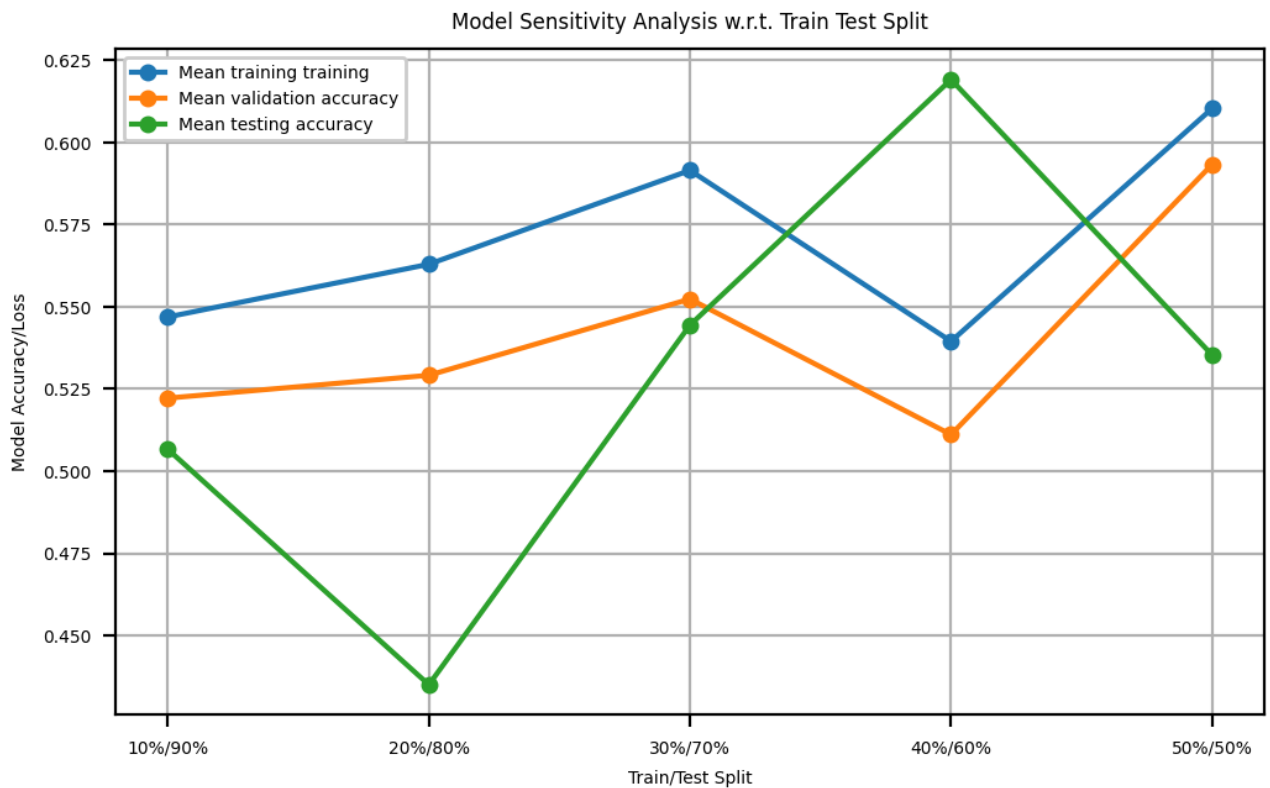
ax.plot(splits,
        model_validation_accuraccies,
        label = r'Mean validation accuracy',
        marker = 'o',
        markersize = 4)

ax.plot(splits,
        model_testing_accuracies,
        label = r'Mean testing accuracy',
        marker = 'o',
        markersize = 4)

ax.set_title("Model Sensitivity Analysis w.r.t. Train Test Split")
ax.set_xlabel('Train/Test Split')
ax.set_ylabel('Model Accuracy/Loss')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)

ax.set_xticks(splits) # Set x-axis tick positions
ax.set_xticklabels([f'{int(split * 100)}%/{100-int(split * 100)}%' for split in splits])
```

```
ax.grid(True)
```



5. Feed Forward Neural Networks

Data Loading and Preprocessing

We first download the [Fashion MNIST dataset](#) through their example code block:

```
In [ ]: # import tensorflow and keras
import tensorflow as tf
import keras

# load data and obtain description for labels
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.fashion_mnist.load_
assert x_train.shape == (60000, 28, 28)
assert x_test.shape == (10000, 28, 28)
assert y_train.shape == (60000,)
assert y_test.shape == (10000,)

# label description
y_description = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat", "Sar
```

Per the keras website:

"This is a dataset of 60,000 28x28 grayscale images of 10 fashion categories, along with a test set of 10,000 images. This dataset can be used as a drop-in replacement for

MNIST."

The classes have the following labels and descriptions:

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

We first plot a selection of the training data for visual analysis. This [website \(geeksforgeeks.org\)](https://www.geeksforgeeks.org/) was used for reference.

```
In [ ]: # obtain data summary
        for i in range(1,10):

            plt.subplot(3,3,i)
            plt.title("Description: "+y_description[y_train[i]])


            plt.imshow(x_train[i], cmap = plt.get_cmap('gray'))

        plt.subplots_adjust(wspace=0, hspace=0.4)

        print(f"Training Data Shape: {x_train.shape}\n"+
              f"Testing Data Shape: {x_test.shape}")
```

Training Data Shape: (60000, 28, 28)

Testing Data Shape: (10000, 28, 28)



We note that the input tensor contains values between 0 and 255, standard for grayscale images, so the data is normalized for model input.

```
In [ ]: # rescale tensor values between 0 and 1
x_train_model = x_train/255.
x_test_model = x_test/255.
```

Model Construction and Architecture

We then construct the 2-layer feedforward neural network using the keras `Sequential()` class. We build the `model_arch()` function in such a way that we can specify the `units` and `activation` function of the resultant `Sequential()` model:

```
In [ ]: from keras import layers

# clear model
keras.backend.clear_session()

# construct model function for implementation
def model_arch(units = 64, activation = "ReLU"):

    # clear and instantiated models
    keras.backend.clear_session()
```



```
# instantiate model
model = keras.Sequential(name = f'units-{units}_activation-{activation}')

# add model input
model.add(keras.Input(shape = (28,28)))

# flatten input of shape (28,28) into a vector
model.add(layers.Flatten())

# add first hidden layer, a naive number of units has been selected as 64
model.add(layers.Dense(name = "hidden_layer", units = units))

if activation == "ReLU":
    model.add(layers.ReLU())
elif activation == "LeakyReLU":
    model.add(layers.LeakyReLU())

# output layer corresponding to the 10 object labels
model.add(layers.Dense(name = "output_layer", units = 10, activation = "softmax"))

# the input shape of our data is 28x28,
return model
```

Next we instantiate a `model_arch()` instance and inspect its architecture with the `summary()` method. Note that we can also see the activation function for the hidden units:

```
In [ ]: model = model_arch(64,"ReLU")
        model.summary()
```

Model: "units-64_activation-ReLU"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
hidden_layer (Dense)	(None, 64)	50240
re_lu (ReLU)	(None, 64)	0
output_layer (Dense)	(None, 10)	650

=====
Total params: 50890 (198.79 KB)
Trainable params: 50890 (198.79 KB)
Non-trainable params: 0 (0.00 Byte)

Model Compilation and Training

Training of the model begins following the assignment of an optimizer, loss function, and

accuracy metric using the `compile()` method.

```
In [ ]: model.compile(optimizer=keras.optimizers.legacy.Adam(learning_rate=1e-3),
                    loss='sparse_categorical_crossentropy',
                    metrics=['sparse_categorical_accuracy'],
                    )
```

Finally, the model is fit to the training data using the `fit()` method where we specify the number of parameters, a validation split, and random shuffling of the training data. Model fitting results in the creation of a `history` object which we can subsequently use to evaluate the model.

```
In [ ]: history = model.fit(
    x_train_model.astype(np.float32), y_train,
    epochs=60,
    validation_split=0.2,
    shuffle = True
)
```

Epoch 1/60

1500/1500 [=====] - 1s 720us/step - loss: 0.5359 - sparse_categorical_accuracy: 0.8142 - val_loss: 0.4989 - val_sparse_categorical_accuracy: 0.8320

Epoch 2/60

1500/1500 [=====] - 1s 676us/step - loss: 0.4027 - sparse_categorical_accuracy: 0.8577 - val_loss: 0.4085 - val_sparse_categorical_accuracy: 0.8503

Epoch 3/60

1500/1500 [=====] - 1s 660us/step - loss: 0.3617 - sparse_categorical_accuracy: 0.8694 - val_loss: 0.3838 - val_sparse_categorical_accuracy: 0.8646

Epoch 4/60

1500/1500 [=====] - 1s 607us/step - loss: 0.3361 - sparse_categorical_accuracy: 0.8770 - val_loss: 0.3652 - val_sparse_categorical_accuracy: 0.8695

Epoch 5/60

1500/1500 [=====] - 1s 671us/step - loss: 0.3165 - sparse_categorical_accuracy: 0.8841 - val_loss: 0.3572 - val_sparse_categorical_accuracy: 0.8717

Epoch 6/60

1500/1500 [=====] - 1s 677us/step - loss: 0.3014 - sparse_categorical_accuracy: 0.8897 - val_loss: 0.3875 - val_sparse_categorical_accuracy: 0.8575

Epoch 7/60

1500/1500 [=====] - 1s 656us/step - loss: 0.2894 - sparse_categorical_accuracy: 0.8925 - val_loss: 0.3370 - val_sparse_categorical_accuracy: 0.8832

Epoch 8/60

1500/1500 [=====] - 1s 656us/step - loss: 0.2774 - sparse_categorical_accuracy: 0.8968 - val_loss: 0.3674 - val_sparse_categorical_accuracy: 0.8718

Epoch 9/60

1500/1500 [=====] - 1s 654us/step - loss: 0.2673 -
sparse_categorical_accuracy: 0.9013 - val_loss: 0.3419 - val_sparse_categorical_accuracy: 0.8783
Epoch 10/60
1500/1500 [=====] - 1s 663us/step - loss: 0.2586 -
sparse_categorical_accuracy: 0.9041 - val_loss: 0.3378 - val_sparse_categorical_accuracy: 0.8814
Epoch 11/60
1500/1500 [=====] - 1s 649us/step - loss: 0.2519 -
sparse_categorical_accuracy: 0.9068 - val_loss: 0.3338 - val_sparse_categorical_accuracy: 0.8838
Epoch 12/60
1500/1500 [=====] - 1s 665us/step - loss: 0.2415 -
sparse_categorical_accuracy: 0.9099 - val_loss: 0.3357 - val_sparse_categorical_accuracy: 0.8838
Epoch 13/60
1500/1500 [=====] - 1s 651us/step - loss: 0.2360 -
sparse_categorical_accuracy: 0.9121 - val_loss: 0.3290 - val_sparse_categorical_accuracy: 0.8857
Epoch 14/60
1500/1500 [=====] - 1s 688us/step - loss: 0.2318 -
sparse_categorical_accuracy: 0.9143 - val_loss: 0.3333 - val_sparse_categorical_accuracy: 0.8861
Epoch 15/60
1500/1500 [=====] - 1s 653us/step - loss: 0.2228 -
sparse_categorical_accuracy: 0.9165 - val_loss: 0.3383 - val_sparse_categorical_accuracy: 0.8842
Epoch 16/60
1500/1500 [=====] - 1s 671us/step - loss: 0.2202 -
sparse_categorical_accuracy: 0.9175 - val_loss: 0.3246 - val_sparse_categorical_accuracy: 0.8913
Epoch 17/60
1500/1500 [=====] - 1s 646us/step - loss: 0.2135 -
sparse_categorical_accuracy: 0.9218 - val_loss: 0.3377 - val_sparse_categorical_accuracy: 0.8843
Epoch 18/60
1500/1500 [=====] - 1s 661us/step - loss: 0.2097 -
sparse_categorical_accuracy: 0.9224 - val_loss: 0.3401 - val_sparse_categorical_accuracy: 0.8852
Epoch 19/60
1500/1500 [=====] - 1s 660us/step - loss: 0.2044 -
sparse_categorical_accuracy: 0.9229 - val_loss: 0.3423 - val_sparse_categorical_accuracy: 0.8852
Epoch 20/60
1500/1500 [=====] - 1s 651us/step - loss: 0.2023 -
sparse_categorical_accuracy: 0.9246 - val_loss: 0.3356 - val_sparse_categorical_accuracy: 0.8905
Epoch 21/60
1500/1500 [=====] - 1s 623us/step - loss: 0.1942 -
sparse_categorical_accuracy: 0.9279 - val_loss: 0.3537 - val_sparse_categorical_accuracy: 0.8826
Epoch 22/60
1500/1500 [=====] - 1s 619us/step - loss: 0.1918 -

sparse_categorical_accuracy: 0.9288 - val_loss: 0.3624 - val_sparse_categorical_accuracy: 0.8841
Epoch 23/60
1500/1500 [=====] - 1s 625us/step - loss: 0.1860 - sparse_categorical_accuracy: 0.9317 - val_loss: 0.3466 - val_sparse_categorical_accuracy: 0.8873
Epoch 24/60
1500/1500 [=====] - 1s 653us/step - loss: 0.1838 - sparse_categorical_accuracy: 0.9316 - val_loss: 0.3558 - val_sparse_categorical_accuracy: 0.8854
Epoch 25/60
1500/1500 [=====] - 1s 638us/step - loss: 0.1791 - sparse_categorical_accuracy: 0.9336 - val_loss: 0.3614 - val_sparse_categorical_accuracy: 0.8884
Epoch 26/60
1500/1500 [=====] - 1s 647us/step - loss: 0.1787 - sparse_categorical_accuracy: 0.9338 - val_loss: 0.3639 - val_sparse_categorical_accuracy: 0.8871
Epoch 27/60
1500/1500 [=====] - 1s 664us/step - loss: 0.1726 - sparse_categorical_accuracy: 0.9369 - val_loss: 0.3548 - val_sparse_categorical_accuracy: 0.8886
Epoch 28/60
1500/1500 [=====] - 1s 683us/step - loss: 0.1713 - sparse_categorical_accuracy: 0.9366 - val_loss: 0.3544 - val_sparse_categorical_accuracy: 0.8908
Epoch 29/60
1500/1500 [=====] - 1s 669us/step - loss: 0.1656 - sparse_categorical_accuracy: 0.9383 - val_loss: 0.3780 - val_sparse_categorical_accuracy: 0.8816
Epoch 30/60
1500/1500 [=====] - 1s 644us/step - loss: 0.1648 - sparse_categorical_accuracy: 0.9387 - val_loss: 0.3671 - val_sparse_categorical_accuracy: 0.8867
Epoch 31/60
1500/1500 [=====] - 1s 666us/step - loss: 0.1612 - sparse_categorical_accuracy: 0.9399 - val_loss: 0.3833 - val_sparse_categorical_accuracy: 0.8843
Epoch 32/60
1500/1500 [=====] - 1s 677us/step - loss: 0.1562 - sparse_categorical_accuracy: 0.9421 - val_loss: 0.3773 - val_sparse_categorical_accuracy: 0.8873
Epoch 33/60
1500/1500 [=====] - 1s 638us/step - loss: 0.1574 - sparse_categorical_accuracy: 0.9416 - val_loss: 0.3801 - val_sparse_categorical_accuracy: 0.8881
Epoch 34/60
1500/1500 [=====] - 1s 672us/step - loss: 0.1553 - sparse_categorical_accuracy: 0.9415 - val_loss: 0.3801 - val_sparse_categorical_accuracy: 0.8908
Epoch 35/60
1500/1500 [=====] - 1s 669us/step - loss: 0.1501 - sparse_categorical_accuracy: 0.9454 - val_loss: 0.3833 - val_sparse_categorical_accuracy: 0.8908

cal_accuracy: 0.8860
Epoch 36/60
1500/1500 [=====] - 1s 653us/step - loss: 0.1491 -
sparse_categorical_accuracy: 0.9443 - val_loss: 0.3892 - val_sparse_categori
cal_accuracy: 0.8887
Epoch 37/60
1500/1500 [=====] - 1s 641us/step - loss: 0.1426 -
sparse_categorical_accuracy: 0.9477 - val_loss: 0.4026 - val_sparse_categori
cal_accuracy: 0.8878
Epoch 38/60
1500/1500 [=====] - 1s 636us/step - loss: 0.1437 -
sparse_categorical_accuracy: 0.9461 - val_loss: 0.4190 - val_sparse_categori
cal_accuracy: 0.8808
Epoch 39/60
1500/1500 [=====] - 1s 682us/step - loss: 0.1400 -
sparse_categorical_accuracy: 0.9477 - val_loss: 0.4055 - val_sparse_categori
cal_accuracy: 0.8852
Epoch 40/60
1500/1500 [=====] - 1s 638us/step - loss: 0.1387 -
sparse_categorical_accuracy: 0.9486 - val_loss: 0.4092 - val_sparse_categori
cal_accuracy: 0.8863
Epoch 41/60
1500/1500 [=====] - 1s 599us/step - loss: 0.1369 -
sparse_categorical_accuracy: 0.9506 - val_loss: 0.4293 - val_sparse_categori
cal_accuracy: 0.8867
Epoch 42/60
1500/1500 [=====] - 1s 633us/step - loss: 0.1348 -
sparse_categorical_accuracy: 0.9500 - val_loss: 0.4059 - val_sparse_categori
cal_accuracy: 0.8898
Epoch 43/60
1500/1500 [=====] - 1s 591us/step - loss: 0.1321 -
sparse_categorical_accuracy: 0.9517 - val_loss: 0.4242 - val_sparse_categori
cal_accuracy: 0.8874
Epoch 44/60
1500/1500 [=====] - 1s 615us/step - loss: 0.1315 -
sparse_categorical_accuracy: 0.9516 - val_loss: 0.4282 - val_sparse_categori
cal_accuracy: 0.8858
Epoch 45/60
1500/1500 [=====] - 1s 670us/step - loss: 0.1309 -
sparse_categorical_accuracy: 0.9512 - val_loss: 0.4246 - val_sparse_categori
cal_accuracy: 0.8876
Epoch 46/60
1500/1500 [=====] - 1s 678us/step - loss: 0.1268 -
sparse_categorical_accuracy: 0.9525 - val_loss: 0.4389 - val_sparse_categori
cal_accuracy: 0.8877
Epoch 47/60
1500/1500 [=====] - 1s 703us/step - loss: 0.1241 -
sparse_categorical_accuracy: 0.9538 - val_loss: 0.4766 - val_sparse_categori
cal_accuracy: 0.8818
Epoch 48/60
1500/1500 [=====] - 1s 707us/step - loss: 0.1234 -
sparse_categorical_accuracy: 0.9539 - val_loss: 0.4511 - val_sparse_categori
cal_accuracy: 0.8826

```

Epoch 49/60
1500/1500 [=====] - 1s 683us/step - loss: 0.1224 -
sparse_categorical_accuracy: 0.9549 - val_loss: 0.4675 - val_sparse_categori
cal_accuracy: 0.8844
Epoch 50/60
1500/1500 [=====] - 1s 655us/step - loss: 0.1205 -
sparse_categorical_accuracy: 0.9557 - val_loss: 0.4449 - val_sparse_categori
cal_accuracy: 0.8860
Epoch 51/60
1500/1500 [=====] - 1s 667us/step - loss: 0.1189 -
sparse_categorical_accuracy: 0.9568 - val_loss: 0.4577 - val_sparse_categori
cal_accuracy: 0.8872
Epoch 52/60
1500/1500 [=====] - 1s 656us/step - loss: 0.1169 -
sparse_categorical_accuracy: 0.9565 - val_loss: 0.4771 - val_sparse_categori
cal_accuracy: 0.8848
Epoch 53/60
1500/1500 [=====] - 1s 685us/step - loss: 0.1145 -
sparse_categorical_accuracy: 0.9580 - val_loss: 0.4645 - val_sparse_categori
cal_accuracy: 0.8842
Epoch 54/60
1500/1500 [=====] - 1s 670us/step - loss: 0.1107 -
sparse_categorical_accuracy: 0.9596 - val_loss: 0.4865 - val_sparse_categori
cal_accuracy: 0.8853
Epoch 55/60
1500/1500 [=====] - 1s 656us/step - loss: 0.1133 -
sparse_categorical_accuracy: 0.9584 - val_loss: 0.4817 - val_sparse_categori
cal_accuracy: 0.8832
Epoch 56/60
1500/1500 [=====] - 1s 694us/step - loss: 0.1111 -
sparse_categorical_accuracy: 0.9588 - val_loss: 0.5076 - val_sparse_categori
cal_accuracy: 0.8779
Epoch 57/60
1500/1500 [=====] - 1s 662us/step - loss: 0.1074 -
sparse_categorical_accuracy: 0.9592 - val_loss: 0.5116 - val_sparse_categori
cal_accuracy: 0.8828
Epoch 58/60
1500/1500 [=====] - 1s 686us/step - loss: 0.1080 -
sparse_categorical_accuracy: 0.9611 - val_loss: 0.4881 - val_sparse_categori
cal_accuracy: 0.8843
Epoch 59/60
1500/1500 [=====] - 1s 593us/step - loss: 0.1069 -
sparse_categorical_accuracy: 0.9603 - val_loss: 0.5196 - val_sparse_categori
cal_accuracy: 0.8823
Epoch 60/60
1500/1500 [=====] - 1s 598us/step - loss: 0.1050 -
sparse_categorical_accuracy: 0.9617 - val_loss: 0.5085 - val_sparse_categori
cal_accuracy: 0.8840

```

We can now observe the models training history w.r.t. validation and training accuracies and losses.

```
In [ ]: # create function to show model training accuracy and loss history
```

```

def model_history(history):
    # plot training history
    fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

    ax.plot(history.history['sparse_categorical_accuracy'],
            label = r'Training Accuracy',
            #marker = 'o',
            markersize = 4)

    ax.plot(history.history['val_sparse_categorical_accuracy'],
            label = r'Validation Accuracy',
            #marker = 'o',
            markersize = 4)

    ax.set_title("Model Sensitivity Analysis w.r.t. Train Test Split")
    ax.set_xlabel('Train/Test Split')
    ax.set_ylabel('Model Accuracy/Loss')
    ax.legend()
    ax.legend().get_frame().set_alpha(1.0)

    ax.grid(True)

    # plot loss history
    fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

    ax.plot(history.history['loss'],
            label = r'Training Loss',
            #marker = 'o',
            markersize = 4)



    ax.plot(history.history['val_loss'],
            label = r'Validation Loss',
            #marker = 'o',
            markersize = 4)

    ax.set_title("Model Sensitivity Analysis w.r.t. Train Test Split")
    ax.set_xlabel('Train/Test Split')
    ax.set_ylabel('Model Accuracy/Loss')
    ax.set_ylim(0,0.6)
    ax.legend()
    ax.legend().get_frame().set_alpha(1.0)

    ax.grid(True)

    # plot training history using the function
    model_history(history)

```



The above loss and accuracy plots are a clear indicator of overfitting from the model. While training accuracy and loss improves with subsequent epochs, validation accuracy plateaus and validation loss appears to increase with subsequent epochs.

We can finally evaluate the model on the test data using sci-kit learn's

`accuracy_score()` method.

```
In [ ]: from sklearn.metrics import accuracy_score

predicted_softmax = model.predict(x_test.astype(np.float32))
y_pred = np.argmax(predicted_softmax, axis = 1)

accuracy_score(y_test, y_pred)
```

313/313 [=====] - 1s 2ms/step

Out[]: 0.862

The test accuracy of 0.862 is in line with the lower validation accuracy as described.

Model Evaluation: leaky ReLU

We next evaluate the Fashion MNIST dataset using a similar model with the keras `LeakyReLU` activation. We can call the previously created `model_arch()` function with the appropriate parameters to create this model.

```
In [ ]: # build identical model with LeakyReLU activation
model = model_arch(64, "LeakyReLU")

# show model architecture
model.summary()
```

Model: "units-64_activation-LeakyReLU"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
hidden_layer (Dense)	(None, 64)	50240
leaky_re_lu (LeakyReLU)	(None, 64)	0
output_layer (Dense)	(None, 10)	650

=====
Total params: 50890 (198.79 KB)
Trainable params: 50890 (198.79 KB)
Non-trainable params: 0 (0.00 Byte)

We then similarly compile and fit the new model.

```
In [ ]: # compile
model.compile(optimizer=keras.optimizers.legacy.Adam(learning_rate=1e-3),
              loss='sparse_categorical_crossentropy',
              metrics=['sparse_categorical_accuracy'],
```

```

    )

# fit
history = model.fit(
    x_train_model.astype(np.float32), y_train,
    epochs=60,
    validation_split=0.2,
    shuffle = True
)

```

Epoch 1/60

1500/1500 [=====] - 1s 711us/step - loss: 0.5424 - sparse_categorical_accuracy: 0.8124 - val_loss: 0.4755 - val_sparse_categorical_accuracy: 0.8313

Epoch 2/60

1500/1500 [=====] - 1s 720us/step - loss: 0.4253 - sparse_categorical_accuracy: 0.8503 - val_loss: 0.4256 - val_sparse_categorical_accuracy: 0.8519

Epoch 3/60

1500/1500 [=====] - 1s 710us/step - loss: 0.3905 - sparse_categorical_accuracy: 0.8597 - val_loss: 0.4038 - val_sparse_categorical_accuracy: 0.8548

Epoch 4/60

1500/1500 [=====] - 1s 673us/step - loss: 0.3676 - sparse_categorical_accuracy: 0.8678 - val_loss: 0.3748 - val_sparse_categorical_accuracy: 0.8670

Epoch 5/60

1500/1500 [=====] - 1s 649us/step - loss: 0.3510 - sparse_categorical_accuracy: 0.8753 - val_loss: 0.3755 - val_sparse_categorical_accuracy: 0.8641

Epoch 6/60

1500/1500 [=====] - 1s 647us/step - loss: 0.3366 - sparse_categorical_accuracy: 0.8791 - val_loss: 0.4112 - val_sparse_categorical_accuracy: 0.8506

Epoch 7/60

1500/1500 [=====] - 1s 623us/step - loss: 0.3253 - sparse_categorical_accuracy: 0.8811 - val_loss: 0.3694 - val_sparse_categorical_accuracy: 0.8656

Epoch 8/60

1500/1500 [=====] - 1s 635us/step - loss: 0.3156 - sparse_categorical_accuracy: 0.8847 - val_loss: 0.3691 - val_sparse_categorical_accuracy: 0.8661

Epoch 9/60

1500/1500 [=====] - 1s 652us/step - loss: 0.3084 - sparse_categorical_accuracy: 0.8881 - val_loss: 0.3475 - val_sparse_categorical_accuracy: 0.8761

Epoch 10/60

1500/1500 [=====] - 1s 701us/step - loss: 0.3005 - sparse_categorical_accuracy: 0.8900 - val_loss: 0.3507 - val_sparse_categorical_accuracy: 0.8741

Epoch 11/60

1500/1500 [=====] - 1s 661us/step - loss: 0.2931 - sparse_categorical_accuracy: 0.8922 - val_loss: 0.3419 - val_sparse_categorical_accuracy: 0.8777

Epoch 12/60
1500/1500 [=====] - 1s 704us/step - loss: 0.2855 -
sparse_categorical_accuracy: 0.8942 - val_loss: 0.3711 - val_sparse_categori
cal_accuracy: 0.8665
Epoch 13/60
1500/1500 [=====] - 1s 666us/step - loss: 0.2808 -
sparse_categorical_accuracy: 0.8972 - val_loss: 0.4133 - val_sparse_categori
cal_accuracy: 0.8512
Epoch 14/60
1500/1500 [=====] - 1s 684us/step - loss: 0.2790 -
sparse_categorical_accuracy: 0.8981 - val_loss: 0.3562 - val_sparse_categori
cal_accuracy: 0.8749
Epoch 15/60
1500/1500 [=====] - 1s 680us/step - loss: 0.2710 -
sparse_categorical_accuracy: 0.8991 - val_loss: 0.3536 - val_sparse_categori
cal_accuracy: 0.8749
Epoch 16/60
1500/1500 [=====] - 1s 665us/step - loss: 0.2660 -
sparse_categorical_accuracy: 0.9011 - val_loss: 0.3673 - val_sparse_categori
cal_accuracy: 0.8745
Epoch 17/60
1500/1500 [=====] - 1s 656us/step - loss: 0.2630 -
sparse_categorical_accuracy: 0.9021 - val_loss: 0.3655 - val_sparse_categori
cal_accuracy: 0.8757
Epoch 18/60
1500/1500 [=====] - 1s 674us/step - loss: 0.2585 -
sparse_categorical_accuracy: 0.9038 - val_loss: 0.3483 - val_sparse_categori
cal_accuracy: 0.8821
Epoch 19/60
1500/1500 [=====] - 1s 733us/step - loss: 0.2540 -
sparse_categorical_accuracy: 0.9050 - val_loss: 0.3555 - val_sparse_categori
cal_accuracy: 0.8790
Epoch 20/60
1500/1500 [=====] - 1s 692us/step - loss: 0.2505 -
sparse_categorical_accuracy: 0.9072 - val_loss: 0.3479 - val_sparse_categori
cal_accuracy: 0.8806
Epoch 21/60
1500/1500 [=====] - 1s 741us/step - loss: 0.2481 -
sparse_categorical_accuracy: 0.9089 - val_loss: 0.3603 - val_sparse_categori
cal_accuracy: 0.8794
Epoch 22/60
1500/1500 [=====] - 1s 712us/step - loss: 0.2421 -
sparse_categorical_accuracy: 0.9103 - val_loss: 0.3560 - val_sparse_categori
cal_accuracy: 0.8824
Epoch 23/60
1500/1500 [=====] - 1s 735us/step - loss: 0.2422 -
sparse_categorical_accuracy: 0.9106 - val_loss: 0.3480 - val_sparse_categori
cal_accuracy: 0.8825
Epoch 24/60
1500/1500 [=====] - 1s 692us/step - loss: 0.2374 -
sparse_categorical_accuracy: 0.9106 - val_loss: 0.3493 - val_sparse_categori
cal_accuracy: 0.8849
Epoch 25/60

1500/1500 [=====] - 1s 785us/step - loss: 0.2346 -
sparse_categorical_accuracy: 0.9131 - val_loss: 0.4125 - val_sparse_categorical_accuracy: 0.8616
Epoch 26/60
1500/1500 [=====] - 1s 735us/step - loss: 0.2304 -
sparse_categorical_accuracy: 0.9147 - val_loss: 0.3629 - val_sparse_categorical_accuracy: 0.8777
Epoch 27/60
1500/1500 [=====] - 1s 663us/step - loss: 0.2292 -
sparse_categorical_accuracy: 0.9151 - val_loss: 0.3411 - val_sparse_categorical_accuracy: 0.8852
Epoch 28/60
1500/1500 [=====] - 1s 717us/step - loss: 0.2281 -
sparse_categorical_accuracy: 0.9153 - val_loss: 0.3466 - val_sparse_categorical_accuracy: 0.8842
Epoch 29/60
1500/1500 [=====] - 1s 687us/step - loss: 0.2242 -
sparse_categorical_accuracy: 0.9180 - val_loss: 0.3478 - val_sparse_categorical_accuracy: 0.8843
Epoch 30/60
1500/1500 [=====] - 1s 788us/step - loss: 0.2199 -
sparse_categorical_accuracy: 0.9174 - val_loss: 0.3614 - val_sparse_categorical_accuracy: 0.8832
Epoch 31/60
1500/1500 [=====] - 1s 716us/step - loss: 0.2180 -
sparse_categorical_accuracy: 0.9183 - val_loss: 0.3494 - val_sparse_categorical_accuracy: 0.8833
Epoch 32/60
1500/1500 [=====] - 1s 670us/step - loss: 0.2155 -
sparse_categorical_accuracy: 0.9208 - val_loss: 0.3588 - val_sparse_categorical_accuracy: 0.8852
Epoch 33/60
1500/1500 [=====] - 1s 729us/step - loss: 0.2168 -
sparse_categorical_accuracy: 0.9199 - val_loss: 0.3550 - val_sparse_categorical_accuracy: 0.8882
Epoch 34/60
1500/1500 [=====] - 1s 706us/step - loss: 0.2115 -
sparse_categorical_accuracy: 0.9208 - val_loss: 0.3649 - val_sparse_categorical_accuracy: 0.8842
Epoch 35/60
1500/1500 [=====] - 1s 770us/step - loss: 0.2081 -
sparse_categorical_accuracy: 0.9216 - val_loss: 0.3685 - val_sparse_categorical_accuracy: 0.8838
Epoch 36/60
1500/1500 [=====] - 1s 721us/step - loss: 0.2066 -
sparse_categorical_accuracy: 0.9231 - val_loss: 0.3707 - val_sparse_categorical_accuracy: 0.8842
Epoch 37/60
1500/1500 [=====] - 1s 699us/step - loss: 0.2063 -
sparse_categorical_accuracy: 0.9242 - val_loss: 0.3702 - val_sparse_categorical_accuracy: 0.8825
Epoch 38/60
1500/1500 [=====] - 1s 681us/step - loss: 0.2059 -

sparse_categorical_accuracy: 0.9244 - val_loss: 0.4002 - val_sparse_categorical_accuracy: 0.8791
Epoch 39/60
1500/1500 [=====] - 1s 698us/step - loss: 0.2007 - sparse_categorical_accuracy: 0.9250 - val_loss: 0.3693 - val_sparse_categorical_accuracy: 0.8856
Epoch 40/60
1500/1500 [=====] - 1s 670us/step - loss: 0.2020 - sparse_categorical_accuracy: 0.9253 - val_loss: 0.3669 - val_sparse_categorical_accuracy: 0.8855
Epoch 41/60
1500/1500 [=====] - 1s 697us/step - loss: 0.1976 - sparse_categorical_accuracy: 0.9267 - val_loss: 0.4182 - val_sparse_categorical_accuracy: 0.8699
Epoch 42/60
1500/1500 [=====] - 1s 685us/step - loss: 0.1993 - sparse_categorical_accuracy: 0.9260 - val_loss: 0.4068 - val_sparse_categorical_accuracy: 0.8734
Epoch 43/60
1500/1500 [=====] - 1s 674us/step - loss: 0.1929 - sparse_categorical_accuracy: 0.9289 - val_loss: 0.3699 - val_sparse_categorical_accuracy: 0.8883
Epoch 44/60
1500/1500 [=====] - 1s 698us/step - loss: 0.1952 - sparse_categorical_accuracy: 0.9273 - val_loss: 0.3901 - val_sparse_categorical_accuracy: 0.8787
Epoch 45/60
1500/1500 [=====] - 1s 732us/step - loss: 0.1889 - sparse_categorical_accuracy: 0.9306 - val_loss: 0.3821 - val_sparse_categorical_accuracy: 0.8867
Epoch 46/60
1500/1500 [=====] - 1s 723us/step - loss: 0.1889 - sparse_categorical_accuracy: 0.9292 - val_loss: 0.3679 - val_sparse_categorical_accuracy: 0.8860
Epoch 47/60
1500/1500 [=====] - 1s 654us/step - loss: 0.1877 - sparse_categorical_accuracy: 0.9306 - val_loss: 0.3839 - val_sparse_categorical_accuracy: 0.8840
Epoch 48/60
1500/1500 [=====] - 1s 657us/step - loss: 0.1914 - sparse_categorical_accuracy: 0.9292 - val_loss: 0.3881 - val_sparse_categorical_accuracy: 0.8850
Epoch 49/60
1500/1500 [=====] - 1s 803us/step - loss: 0.1824 - sparse_categorical_accuracy: 0.9324 - val_loss: 0.3805 - val_sparse_categorical_accuracy: 0.8866
Epoch 50/60
1500/1500 [=====] - 1s 684us/step - loss: 0.1832 - sparse_categorical_accuracy: 0.9317 - val_loss: 0.3936 - val_sparse_categorical_accuracy: 0.8834
Epoch 51/60
1500/1500 [=====] - 1s 690us/step - loss: 0.1829 - sparse_categorical_accuracy: 0.9321 - val_loss: 0.4595 - val_sparse_categorical_accuracy: 0.8834

```

cal_accuracy: 0.8718
Epoch 52/60
1500/1500 [=====] - 1s 694us/step - loss: 0.1814 -
sparse_categorical_accuracy: 0.9332 - val_loss: 0.4155 - val_sparse_categori
cal_accuracy: 0.8797
Epoch 53/60
1500/1500 [=====] - 1s 682us/step - loss: 0.1808 -
sparse_categorical_accuracy: 0.9329 - val_loss: 0.3798 - val_sparse_categori
cal_accuracy: 0.8874
Epoch 54/60
1500/1500 [=====] - 1s 699us/step - loss: 0.1781 -
sparse_categorical_accuracy: 0.9344 - val_loss: 0.4059 - val_sparse_categori
cal_accuracy: 0.8836
Epoch 55/60
1500/1500 [=====] - 1s 737us/step - loss: 0.1790 -
sparse_categorical_accuracy: 0.9334 - val_loss: 0.4286 - val_sparse_categori
cal_accuracy: 0.8736
Epoch 56/60
1500/1500 [=====] - 1s 684us/step - loss: 0.1752 -
sparse_categorical_accuracy: 0.9366 - val_loss: 0.4094 - val_sparse_categori
cal_accuracy: 0.8803
Epoch 57/60
1500/1500 [=====] - 1s 687us/step - loss: 0.1720 -
sparse_categorical_accuracy: 0.9368 - val_loss: 0.4100 - val_sparse_categori
cal_accuracy: 0.8808
Epoch 58/60
1500/1500 [=====] - 1s 714us/step - loss: 0.1746 -
sparse_categorical_accuracy: 0.9361 - val_loss: 0.3923 - val_sparse_categori
cal_accuracy: 0.8858
Epoch 59/60
1500/1500 [=====] - 1s 698us/step - loss: 0.1698 -
sparse_categorical_accuracy: 0.9371 - val_loss: 0.4017 - val_sparse_categori
cal_accuracy: 0.8836
Epoch 60/60
1500/1500 [=====] - 1s 723us/step - loss: 0.1682 -
sparse_categorical_accuracy: 0.9381 - val_loss: 0.4382 - val_sparse_categori
cal_accuracy: 0.8784

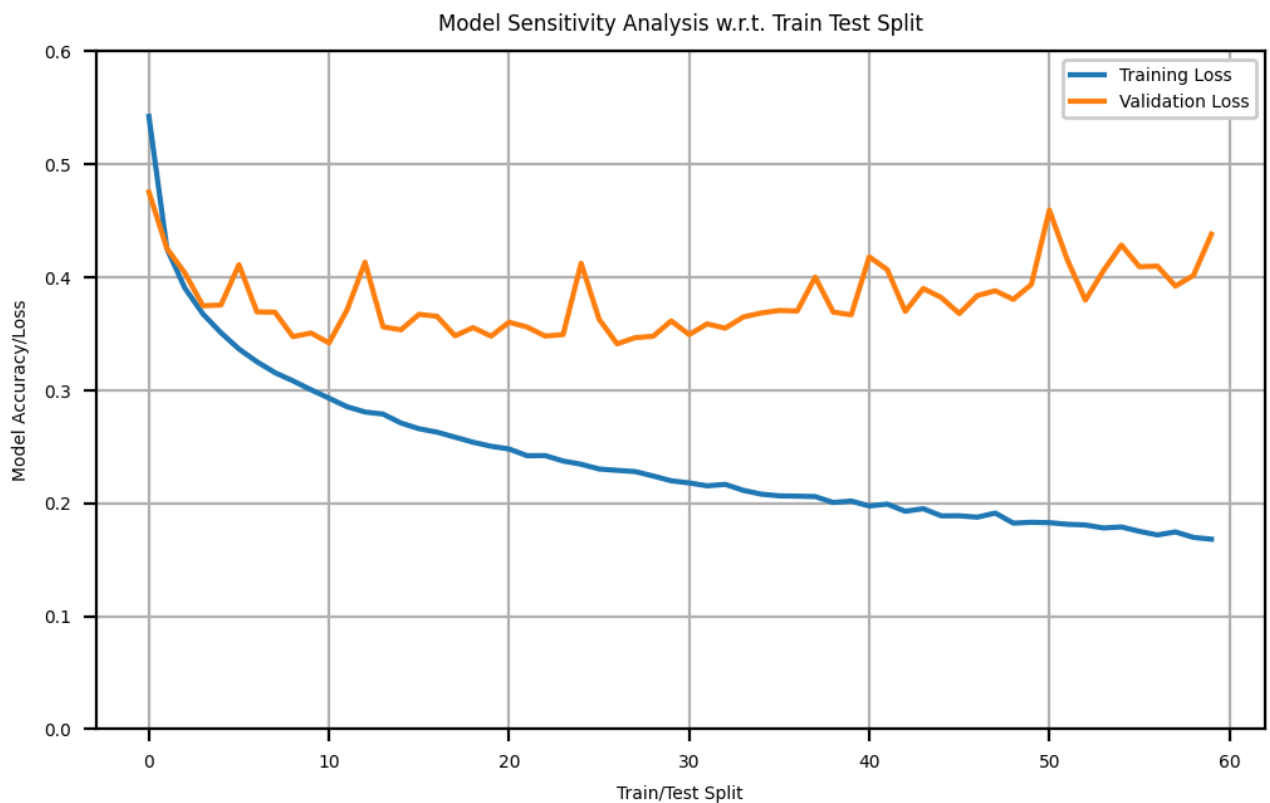

```

The LeakyReLU model training history is as follows:

```

In [ ]: # plot training history
        model_history(history)

```



```
In [ ]: predicted_softmax = model.predict(x_test.astype(np.float32))
        y_pred = np.argmax(predicted_softmax, axis = 1)

        accuracy_score(y_test, y_pred)
```

313/313 [=====] - 1s 2ms/step

Out[]: 0.8536

The model utilizing `LeakyReLU` had a test accuracy of `0.8536`, comparable to the former `ReLU` model. Should a model be selected for further refinement or deployment, on the basis of activation function alone, the `LeakyReLU` model would be the preferred choice due to better loss history.

Optimizer Choice Influence

From the [keras](#) website there are a number of optimizers to choose from for model architecture. The following optimizers were selected for analysis:

- [SGD](#): Stochastic Gradient Descent with momentum optimizer.
- [RMSprop](#): Optimizer that implements the RMSprop algorithm. Maintain a moving (discounted) average of the square of gradients, and divide the gradient by the root of this average.
- [ADAM](#): Adam optimization is a stochastic gradient descent method that is based on adaptive estimation of first-order and second-order moments.
- [Adadelata](#): Adadelata optimization is a stochastic gradient descent method that is based on adaptive learning rate per dimension.
- [Ftrl](#): "Follow The Regularized Leader" (FTRL) is an optimization algorithm developed at Google for click-through rate prediction in the early 2010s. It is most suitable for shallow models with large and sparse feature spaces.

Similar to the methodology presented in **Problem 4**, we will maintain the best hyperparameters from the previous **ffn** analysis and vary the subject optimizer hyperparameter for evaluation.

```
In [ ]: # create list of optimizers and names for loop implementation
optimizers = [
    # instantiate optimizer objects with the same learning rates
    keras.optimizers.legacy.SGD(learning_rate=1e-3),
    keras.optimizers.legacy.RMSprop(learning_rate=1e-3),
    keras.optimizers.legacy.Adam(learning_rate=1e-3),
    keras.optimizers.legacy.Adadelata(learning_rate=1e-3),
    keras.optimizers.legacy.Ftrl(learning_rate=1e-3)
]

names = ["SGD", "RMSprop", "ADAM", "Adadelata", "Ftrl"]

# instantiate tracking objects for analysis
histories = [] # empty list to save history objects for each model
train_loss = [] # save train loss for each mode
train_accuracies = [] # save train accuracy for each model
test_accuracies = [] # save test accuracy for each model
```



```

# begin training loop
print("Begin evaluation:")

for name, optimizer in zip(names, optimizers):

    # clear environment of previous model
    keras.backend.clear_session()

    # create ffn model using previously evaluated best hyperparameters
    model = model_arch(64, "LeakyReLU")

    # compile model with each optimizer
    model.compile(optimizer=optimizer,
                  loss='sparse_categorical_crossentropy',
                  metrics=['sparse_categorical_accuracy'],
                  )

    # begin training, for the sake of brevity, number of epochs has been halved
    print(f"\nTraining, {name} optimizer\n")
    history = model.fit(
        x_train_model.astype(np.float32), y_train,
        epochs=30,
        validation_split=0.2,
        shuffle = True
    )

    # evaluate on test data
    print(f"\n {name} Test Accuracy:")
    predicted_softmax = model.predict(x_test.astype(np.float32))
    y_pred = np.argmax(predicted_softmax, axis = 1)

    test_accuracy = accuracy_score(y_test, y_pred)

    print(test_accuracy)

    # append tracking objects
    histories.append(history)
    train_loss.append(history.history["loss"][-1])
    train_accuracies.append(history.history["sparse_categorical_accuracy"][-1])
    test_accuracies.append(test_accuracy)

```

Begin evaluation:

Training, SGD optimizer

Epoch 1/30

1500/1500 [=====] - 1s 618us/step - loss: 1.5208 - sparse_categorical_accuracy: 0.5461 - val_loss: 1.1038 - val_sparse_categorical_accuracy: 0.6734

Epoch 2/30

1500/1500 [=====] - 1s 590us/step - loss: 0.9747 - sparse_categorical_accuracy: 0.6941 - val_loss: 0.8665 - val_sparse_categorical_accuracy: 0.7182

Epoch 3/30

1500/1500 [=====] - 1s 571us/step - loss: 0.8250 -
sparse_categorical_accuracy: 0.7289 - val_loss: 0.7710 - val_sparse_categorical_accuracy: 0.7442
Epoch 4/30
1500/1500 [=====] - 1s 567us/step - loss: 0.7509 -
sparse_categorical_accuracy: 0.7512 - val_loss: 0.7153 - val_sparse_categorical_accuracy: 0.7619
Epoch 5/30
1500/1500 [=====] - 1s 585us/step - loss: 0.7028 -
sparse_categorical_accuracy: 0.7685 - val_loss: 0.6767 - val_sparse_categorical_accuracy: 0.7747
Epoch 6/30
1500/1500 [=====] - 1s 577us/step - loss: 0.6677 -
sparse_categorical_accuracy: 0.7801 - val_loss: 0.6475 - val_sparse_categorical_accuracy: 0.7793
Epoch 7/30
1500/1500 [=====] - 1s 559us/step - loss: 0.6410 -
sparse_categorical_accuracy: 0.7893 - val_loss: 0.6259 - val_sparse_categorical_accuracy: 0.7887
Epoch 8/30
1500/1500 [=====] - 1s 553us/step - loss: 0.6193 -
sparse_categorical_accuracy: 0.7953 - val_loss: 0.6065 - val_sparse_categorical_accuracy: 0.7928
Epoch 9/30
1500/1500 [=====] - 1s 493us/step - loss: 0.6015 -
sparse_categorical_accuracy: 0.8014 - val_loss: 0.5922 - val_sparse_categorical_accuracy: 0.7986
Epoch 10/30
1500/1500 [=====] - 1s 515us/step - loss: 0.5862 -
sparse_categorical_accuracy: 0.8055 - val_loss: 0.5778 - val_sparse_categorical_accuracy: 0.8018
Epoch 11/30
1500/1500 [=====] - 1s 490us/step - loss: 0.5731 -
sparse_categorical_accuracy: 0.8095 - val_loss: 0.5671 - val_sparse_categorical_accuracy: 0.8075
Epoch 12/30
1500/1500 [=====] - 1s 533us/step - loss: 0.5618 -
sparse_categorical_accuracy: 0.8136 - val_loss: 0.5565 - val_sparse_categorical_accuracy: 0.8091
Epoch 13/30
1500/1500 [=====] - 1s 562us/step - loss: 0.5518 -
sparse_categorical_accuracy: 0.8171 - val_loss: 0.5480 - val_sparse_categorical_accuracy: 0.8125
Epoch 14/30
1500/1500 [=====] - 1s 481us/step - loss: 0.5430 -
sparse_categorical_accuracy: 0.8186 - val_loss: 0.5405 - val_sparse_categorical_accuracy: 0.8138
Epoch 15/30
1500/1500 [=====] - 1s 499us/step - loss: 0.5349 -
sparse_categorical_accuracy: 0.8213 - val_loss: 0.5349 - val_sparse_categorical_accuracy: 0.8168
Epoch 16/30
1500/1500 [=====] - 1s 533us/step - loss: 0.5278 -

sparse_categorical_accuracy: 0.8233 - val_loss: 0.5280 - val_sparse_categorical_accuracy: 0.8174
Epoch 17/30
1500/1500 [=====] - 1s 478us/step - loss: 0.5211 - sparse_categorical_accuracy: 0.8244 - val_loss: 0.5215 - val_sparse_categorical_accuracy: 0.8210
Epoch 18/30
1500/1500 [=====] - 1s 496us/step - loss: 0.5154 - sparse_categorical_accuracy: 0.8274 - val_loss: 0.5171 - val_sparse_categorical_accuracy: 0.8201
Epoch 19/30
1500/1500 [=====] - 1s 514us/step - loss: 0.5100 - sparse_categorical_accuracy: 0.8280 - val_loss: 0.5127 - val_sparse_categorical_accuracy: 0.8204
Epoch 20/30
1500/1500 [=====] - 1s 472us/step - loss: 0.5050 - sparse_categorical_accuracy: 0.8301 - val_loss: 0.5090 - val_sparse_categorical_accuracy: 0.8230
Epoch 21/30
1500/1500 [=====] - 1s 485us/step - loss: 0.5004 - sparse_categorical_accuracy: 0.8313 - val_loss: 0.5034 - val_sparse_categorical_accuracy: 0.8275
Epoch 22/30
1500/1500 [=====] - 1s 503us/step - loss: 0.4962 - sparse_categorical_accuracy: 0.8325 - val_loss: 0.4996 - val_sparse_categorical_accuracy: 0.8264
Epoch 23/30
1500/1500 [=====] - 1s 469us/step - loss: 0.4922 - sparse_categorical_accuracy: 0.8335 - val_loss: 0.4956 - val_sparse_categorical_accuracy: 0.8285
Epoch 24/30
1500/1500 [=====] - 1s 494us/step - loss: 0.4884 - sparse_categorical_accuracy: 0.8346 - val_loss: 0.4932 - val_sparse_categorical_accuracy: 0.8277
Epoch 25/30
1500/1500 [=====] - 1s 484us/step - loss: 0.4847 - sparse_categorical_accuracy: 0.8356 - val_loss: 0.4893 - val_sparse_categorical_accuracy: 0.8298
Epoch 26/30
1500/1500 [=====] - 1s 465us/step - loss: 0.4815 - sparse_categorical_accuracy: 0.8365 - val_loss: 0.4868 - val_sparse_categorical_accuracy: 0.8305
Epoch 27/30
1500/1500 [=====] - 1s 487us/step - loss: 0.4783 - sparse_categorical_accuracy: 0.8369 - val_loss: 0.4868 - val_sparse_categorical_accuracy: 0.8288
Epoch 28/30
1500/1500 [=====] - 1s 468us/step - loss: 0.4755 - sparse_categorical_accuracy: 0.8382 - val_loss: 0.4817 - val_sparse_categorical_accuracy: 0.8317
Epoch 29/30
1500/1500 [=====] - 1s 483us/step - loss: 0.4727 - sparse_categorical_accuracy: 0.8388 - val_loss: 0.4784 - val_sparse_categorical_accuracy: 0.8317

cal_accuracy: 0.8339
Epoch 30/30
1500/1500 [=====] - 1s 486us/step - loss: 0.4699 -
sparse_categorical_accuracy: 0.8390 - val_loss: 0.4764 - val_sparse_categori
cal_accuracy: 0.8353

SGD Test Accuracy:
313/313 [=====] - 0s 293us/step
0.7912

Training, RMSprop optimizer

Epoch 1/30
1500/1500 [=====] - 1s 580us/step - loss: 0.5607 -
sparse_categorical_accuracy: 0.8040 - val_loss: 0.4686 - val_sparse_categori
cal_accuracy: 0.8333
Epoch 2/30
1500/1500 [=====] - 1s 526us/step - loss: 0.4326 -
sparse_categorical_accuracy: 0.8485 - val_loss: 0.4216 - val_sparse_categori
cal_accuracy: 0.8480
Epoch 3/30
1500/1500 [=====] - 1s 547us/step - loss: 0.3992 -
sparse_categorical_accuracy: 0.8594 - val_loss: 0.3869 - val_sparse_categori
cal_accuracy: 0.8640
Epoch 4/30
1500/1500 [=====] - 1s 542us/step - loss: 0.3770 -
sparse_categorical_accuracy: 0.8656 - val_loss: 0.3778 - val_sparse_categori
cal_accuracy: 0.8650
Epoch 5/30
1500/1500 [=====] - 1s 526us/step - loss: 0.3609 -
sparse_categorical_accuracy: 0.8699 - val_loss: 0.3854 - val_sparse_categori
cal_accuracy: 0.8612
Epoch 6/30
1500/1500 [=====] - 1s 540us/step - loss: 0.3482 -
sparse_categorical_accuracy: 0.8759 - val_loss: 0.3852 - val_sparse_categori
cal_accuracy: 0.8656
Epoch 7/30
1500/1500 [=====] - 1s 540us/step - loss: 0.3382 -
sparse_categorical_accuracy: 0.8790 - val_loss: 0.3748 - val_sparse_categori
cal_accuracy: 0.8668
Epoch 8/30
1500/1500 [=====] - 1s 543us/step - loss: 0.3290 -
sparse_categorical_accuracy: 0.8815 - val_loss: 0.3513 - val_sparse_categori
cal_accuracy: 0.8761
Epoch 9/30
1500/1500 [=====] - 1s 525us/step - loss: 0.3199 -
sparse_categorical_accuracy: 0.8855 - val_loss: 0.3871 - val_sparse_categori
cal_accuracy: 0.8671
Epoch 10/30
1500/1500 [=====] - 1s 543us/step - loss: 0.3145 -
sparse_categorical_accuracy: 0.8873 - val_loss: 0.3934 - val_sparse_categori
cal_accuracy: 0.8642
Epoch 11/30

1500/1500 [=====] - 1s 554us/step - loss: 0.3087 -
sparse_categorical_accuracy: 0.8882 - val_loss: 0.3642 - val_sparse_categorical_accuracy: 0.8723
Epoch 12/30
1500/1500 [=====] - 1s 540us/step - loss: 0.3027 -
sparse_categorical_accuracy: 0.8918 - val_loss: 0.3548 - val_sparse_categorical_accuracy: 0.8771
Epoch 13/30
1500/1500 [=====] - 1s 570us/step - loss: 0.2994 -
sparse_categorical_accuracy: 0.8920 - val_loss: 0.3762 - val_sparse_categorical_accuracy: 0.8723
Epoch 14/30
1500/1500 [=====] - 1s 629us/step - loss: 0.2944 -
sparse_categorical_accuracy: 0.8955 - val_loss: 0.3468 - val_sparse_categorical_accuracy: 0.8799
Epoch 15/30
1500/1500 [=====] - 1s 565us/step - loss: 0.2899 -
sparse_categorical_accuracy: 0.8956 - val_loss: 0.3688 - val_sparse_categorical_accuracy: 0.8755
Epoch 16/30
1500/1500 [=====] - 1s 562us/step - loss: 0.2863 -
sparse_categorical_accuracy: 0.8963 - val_loss: 0.3492 - val_sparse_categorical_accuracy: 0.8849
Epoch 17/30
1500/1500 [=====] - 1s 560us/step - loss: 0.2837 -
sparse_categorical_accuracy: 0.8987 - val_loss: 0.3660 - val_sparse_categorical_accuracy: 0.8763
Epoch 18/30
1500/1500 [=====] - 1s 630us/step - loss: 0.2805 -
sparse_categorical_accuracy: 0.9000 - val_loss: 0.3613 - val_sparse_categorical_accuracy: 0.8817
Epoch 19/30
1500/1500 [=====] - 1s 601us/step - loss: 0.2761 -
sparse_categorical_accuracy: 0.9002 - val_loss: 0.3621 - val_sparse_categorical_accuracy: 0.8798
Epoch 20/30
1500/1500 [=====] - 1s 634us/step - loss: 0.2739 -
sparse_categorical_accuracy: 0.9018 - val_loss: 0.3469 - val_sparse_categorical_accuracy: 0.8827
Epoch 21/30
1500/1500 [=====] - 1s 633us/step - loss: 0.2716 -
sparse_categorical_accuracy: 0.9026 - val_loss: 0.3698 - val_sparse_categorical_accuracy: 0.8759
Epoch 22/30
1500/1500 [=====] - 1s 592us/step - loss: 0.2694 -
sparse_categorical_accuracy: 0.9031 - val_loss: 0.3730 - val_sparse_categorical_accuracy: 0.8795
Epoch 23/30
1500/1500 [=====] - 1s 578us/step - loss: 0.2661 -
sparse_categorical_accuracy: 0.9041 - val_loss: 0.4004 - val_sparse_categorical_accuracy: 0.8683
Epoch 24/30
1500/1500 [=====] - 1s 587us/step - loss: 0.2638 -

sparse_categorical_accuracy: 0.9046 - val_loss: 0.3804 - val_sparse_categorical_accuracy: 0.8793
Epoch 25/30
1500/1500 [=====] - 1s 581us/step - loss: 0.2629 - sparse_categorical_accuracy: 0.9063 - val_loss: 0.3738 - val_sparse_categorical_accuracy: 0.8783
Epoch 26/30
1500/1500 [=====] - 1s 581us/step - loss: 0.2616 - sparse_categorical_accuracy: 0.9059 - val_loss: 0.3957 - val_sparse_categorical_accuracy: 0.8757
Epoch 27/30
1500/1500 [=====] - 1s 573us/step - loss: 0.2572 - sparse_categorical_accuracy: 0.9064 - val_loss: 0.3666 - val_sparse_categorical_accuracy: 0.8815
Epoch 28/30
1500/1500 [=====] - 1s 596us/step - loss: 0.2569 - sparse_categorical_accuracy: 0.9075 - val_loss: 0.3555 - val_sparse_categorical_accuracy: 0.8837
Epoch 29/30
1500/1500 [=====] - 1s 590us/step - loss: 0.2543 - sparse_categorical_accuracy: 0.9097 - val_loss: 0.3651 - val_sparse_categorical_accuracy: 0.8817
Epoch 30/30
1500/1500 [=====] - 1s 584us/step - loss: 0.2518 - sparse_categorical_accuracy: 0.9093 - val_loss: 0.3615 - val_sparse_categorical_accuracy: 0.8861

RMSprop Test Accuracy:
313/313 [=====] - 0s 313us/step
0.8504

Training, ADAM optimizer

Epoch 1/30
1500/1500 [=====] - 1s 659us/step - loss: 0.5376 - sparse_categorical_accuracy: 0.8130 - val_loss: 0.4813 - val_sparse_categorical_accuracy: 0.8274
Epoch 2/30
1500/1500 [=====] - 1s 660us/step - loss: 0.4219 - sparse_categorical_accuracy: 0.8515 - val_loss: 0.4405 - val_sparse_categorical_accuracy: 0.8424
Epoch 3/30
1500/1500 [=====] - 1s 624us/step - loss: 0.3856 - sparse_categorical_accuracy: 0.8621 - val_loss: 0.4302 - val_sparse_categorical_accuracy: 0.8398
Epoch 4/30
1500/1500 [=====] - 1s 626us/step - loss: 0.3659 - sparse_categorical_accuracy: 0.8687 - val_loss: 0.4037 - val_sparse_categorical_accuracy: 0.8610
Epoch 5/30
1500/1500 [=====] - 1s 635us/step - loss: 0.3486 - sparse_categorical_accuracy: 0.8753 - val_loss: 0.3726 - val_sparse_categorical_accuracy: 0.8661

Epoch 6/30
1500/1500 [=====] - 1s 672us/step - loss: 0.3350 - sparse_categorical_accuracy: 0.8771 - val_loss: 0.3649 - val_sparse_categorical_accuracy: 0.8717

Epoch 7/30
1500/1500 [=====] - 1s 619us/step - loss: 0.3227 - sparse_categorical_accuracy: 0.8826 - val_loss: 0.3489 - val_sparse_categorical_accuracy: 0.8763

Epoch 8/30
1500/1500 [=====] - 1s 619us/step - loss: 0.3160 - sparse_categorical_accuracy: 0.8841 - val_loss: 0.3629 - val_sparse_categorical_accuracy: 0.8697

Epoch 9/30
1500/1500 [=====] - 1s 621us/step - loss: 0.3074 - sparse_categorical_accuracy: 0.8882 - val_loss: 0.3698 - val_sparse_categorical_accuracy: 0.8673

Epoch 10/30
1500/1500 [=====] - 1s 618us/step - loss: 0.2988 - sparse_categorical_accuracy: 0.8910 - val_loss: 0.3424 - val_sparse_categorical_accuracy: 0.8796

Epoch 11/30
1500/1500 [=====] - 1s 611us/step - loss: 0.2932 - sparse_categorical_accuracy: 0.8935 - val_loss: 0.3488 - val_sparse_categorical_accuracy: 0.8757

Epoch 12/30
1500/1500 [=====] - 1s 674us/step - loss: 0.2871 - sparse_categorical_accuracy: 0.8941 - val_loss: 0.3414 - val_sparse_categorical_accuracy: 0.8808

Epoch 13/30
1500/1500 [=====] - 1s 659us/step - loss: 0.2797 - sparse_categorical_accuracy: 0.8975 - val_loss: 0.3459 - val_sparse_categorical_accuracy: 0.8769

Epoch 14/30
1500/1500 [=====] - 1s 706us/step - loss: 0.2756 - sparse_categorical_accuracy: 0.8995 - val_loss: 0.3410 - val_sparse_categorical_accuracy: 0.8802

Epoch 15/30
1500/1500 [=====] - 1s 675us/step - loss: 0.2674 - sparse_categorical_accuracy: 0.9020 - val_loss: 0.3501 - val_sparse_categorical_accuracy: 0.8760

Epoch 16/30
1500/1500 [=====] - 1s 658us/step - loss: 0.2672 - sparse_categorical_accuracy: 0.9021 - val_loss: 0.4051 - val_sparse_categorical_accuracy: 0.8577

Epoch 17/30
1500/1500 [=====] - 1s 659us/step - loss: 0.2627 - sparse_categorical_accuracy: 0.9036 - val_loss: 0.3680 - val_sparse_categorical_accuracy: 0.8733

Epoch 18/30
1500/1500 [=====] - 1s 655us/step - loss: 0.2593 - sparse_categorical_accuracy: 0.9046 - val_loss: 0.3499 - val_sparse_categorical_accuracy: 0.8787

Epoch 19/30

1500/1500 [=====] - 1s 638us/step - loss: 0.2534 -
sparse_categorical_accuracy: 0.9066 - val_loss: 0.3582 - val_sparse_categorical_accuracy: 0.8752
Epoch 20/30
1500/1500 [=====] - 1s 607us/step - loss: 0.2521 -
sparse_categorical_accuracy: 0.9074 - val_loss: 0.3459 - val_sparse_categorical_accuracy: 0.8827
Epoch 21/30
1500/1500 [=====] - 1s 673us/step - loss: 0.2448 -
sparse_categorical_accuracy: 0.9101 - val_loss: 0.3464 - val_sparse_categorical_accuracy: 0.8787
Epoch 22/30
1500/1500 [=====] - 1s 659us/step - loss: 0.2429 -
sparse_categorical_accuracy: 0.9105 - val_loss: 0.3792 - val_sparse_categorical_accuracy: 0.8697
Epoch 23/30
1500/1500 [=====] - 1s 672us/step - loss: 0.2390 -
sparse_categorical_accuracy: 0.9120 - val_loss: 0.3391 - val_sparse_categorical_accuracy: 0.8823
Epoch 24/30
1500/1500 [=====] - 1s 615us/step - loss: 0.2363 -
sparse_categorical_accuracy: 0.9118 - val_loss: 0.3570 - val_sparse_categorical_accuracy: 0.8802
Epoch 25/30
1500/1500 [=====] - 1s 618us/step - loss: 0.2314 -
sparse_categorical_accuracy: 0.9146 - val_loss: 0.3469 - val_sparse_categorical_accuracy: 0.8849
Epoch 26/30
1500/1500 [=====] - 1s 629us/step - loss: 0.2280 -
sparse_categorical_accuracy: 0.9160 - val_loss: 0.3405 - val_sparse_categorical_accuracy: 0.8843
Epoch 27/30
1500/1500 [=====] - 1s 622us/step - loss: 0.2272 -
sparse_categorical_accuracy: 0.9156 - val_loss: 0.3461 - val_sparse_categorical_accuracy: 0.8835
Epoch 28/30
1500/1500 [=====] - 1s 627us/step - loss: 0.2265 -
sparse_categorical_accuracy: 0.9169 - val_loss: 0.3821 - val_sparse_categorical_accuracy: 0.8720
Epoch 29/30
1500/1500 [=====] - 1s 617us/step - loss: 0.2220 -
sparse_categorical_accuracy: 0.9171 - val_loss: 0.3581 - val_sparse_categorical_accuracy: 0.8816
Epoch 30/30
1500/1500 [=====] - 1s 629us/step - loss: 0.2238 -
sparse_categorical_accuracy: 0.9162 - val_loss: 0.3690 - val_sparse_categorical_accuracy: 0.8791

ADAM Test Accuracy:

313/313 [=====] - 0s 321us/step
0.863

Training, Adadelata optimizer

Epoch 1/30
1500/1500 [=====] - 1s 701us/step - loss: 2.2446 - sparse_categorical_accuracy: 0.1600 - val_loss: 2.0975 - val_sparse_categorical_accuracy: 0.2637

Epoch 2/30
1500/1500 [=====] - 1s 677us/step - loss: 1.9932 - sparse_categorical_accuracy: 0.3221 - val_loss: 1.8843 - val_sparse_categorical_accuracy: 0.3969

Epoch 3/30
1500/1500 [=====] - 1s 675us/step - loss: 1.8054 - sparse_categorical_accuracy: 0.4533 - val_loss: 1.7145 - val_sparse_categorical_accuracy: 0.5110

Epoch 4/30
1500/1500 [=====] - 1s 673us/step - loss: 1.6520 - sparse_categorical_accuracy: 0.5379 - val_loss: 1.5739 - val_sparse_categorical_accuracy: 0.5702

Epoch 5/30
1500/1500 [=====] - 1s 679us/step - loss: 1.5238 - sparse_categorical_accuracy: 0.5798 - val_loss: 1.4561 - val_sparse_categorical_accuracy: 0.6007

Epoch 6/30
1500/1500 [=====] - 1s 674us/step - loss: 1.4166 - sparse_categorical_accuracy: 0.6068 - val_loss: 1.3576 - val_sparse_categorical_accuracy: 0.6241

Epoch 7/30
1500/1500 [=====] - 1s 675us/step - loss: 1.3270 - sparse_categorical_accuracy: 0.6279 - val_loss: 1.2760 - val_sparse_categorical_accuracy: 0.6434

Epoch 8/30
1500/1500 [=====] - 1s 692us/step - loss: 1.2530 - sparse_categorical_accuracy: 0.6436 - val_loss: 1.2081 - val_sparse_categorical_accuracy: 0.6559

Epoch 9/30
1500/1500 [=====] - 1s 678us/step - loss: 1.1909 - sparse_categorical_accuracy: 0.6546 - val_loss: 1.1513 - val_sparse_categorical_accuracy: 0.6689

Epoch 10/30
1500/1500 [=====] - 1s 675us/step - loss: 1.1383 - sparse_categorical_accuracy: 0.6656 - val_loss: 1.1026 - val_sparse_categorical_accuracy: 0.6790

Epoch 11/30
1500/1500 [=====] - 1s 630us/step - loss: 1.0934 - sparse_categorical_accuracy: 0.6734 - val_loss: 1.0608 - val_sparse_categorical_accuracy: 0.6862

Epoch 12/30
1500/1500 [=====] - 1s 641us/step - loss: 1.0544 - sparse_categorical_accuracy: 0.6799 - val_loss: 1.0243 - val_sparse_categorical_accuracy: 0.6933

Epoch 13/30
1500/1500 [=====] - 1s 641us/step - loss: 1.0203 - sparse_categorical_accuracy: 0.6863 - val_loss: 0.9923 - val_sparse_categorical_accuracy: 0.6978

Epoch 14/30
1500/1500 [=====] - 1s 640us/step - loss: 0.9904 -
sparse_categorical_accuracy: 0.6918 - val_loss: 0.9644 - val_sparse_categori
cal_accuracy: 0.7034
Epoch 15/30
1500/1500 [=====] - 1s 684us/step - loss: 0.9640 -
sparse_categorical_accuracy: 0.6968 - val_loss: 0.9395 - val_sparse_categori
cal_accuracy: 0.7089
Epoch 16/30
1500/1500 [=====] - 1s 648us/step - loss: 0.9404 -
sparse_categorical_accuracy: 0.7024 - val_loss: 0.9172 - val_sparse_categori
cal_accuracy: 0.7143
Epoch 17/30
1500/1500 [=====] - 1s 627us/step - loss: 0.9191 -
sparse_categorical_accuracy: 0.7083 - val_loss: 0.8971 - val_sparse_categori
cal_accuracy: 0.7171
Epoch 18/30
1500/1500 [=====] - 1s 630us/step - loss: 0.9000 -
sparse_categorical_accuracy: 0.7129 - val_loss: 0.8791 - val_sparse_categori
cal_accuracy: 0.7220
Epoch 19/30
1500/1500 [=====] - 1s 622us/step - loss: 0.8827 -
sparse_categorical_accuracy: 0.7178 - val_loss: 0.8626 - val_sparse_categori
cal_accuracy: 0.7266
Epoch 20/30
1500/1500 [=====] - 1s 629us/step - loss: 0.8666 -
sparse_categorical_accuracy: 0.7218 - val_loss: 0.8474 - val_sparse_categori
cal_accuracy: 0.7309
Epoch 21/30
1500/1500 [=====] - 1s 635us/step - loss: 0.8520 -
sparse_categorical_accuracy: 0.7260 - val_loss: 0.8338 - val_sparse_categori
cal_accuracy: 0.7344
Epoch 22/30
1500/1500 [=====] - 1s 620us/step - loss: 0.8387 -
sparse_categorical_accuracy: 0.7298 - val_loss: 0.8212 - val_sparse_categori
cal_accuracy: 0.7375
Epoch 23/30
1500/1500 [=====] - 1s 628us/step - loss: 0.8264 -
sparse_categorical_accuracy: 0.7332 - val_loss: 0.8094 - val_sparse_categori
cal_accuracy: 0.7409
Epoch 24/30
1500/1500 [=====] - 1s 625us/step - loss: 0.8148 -
sparse_categorical_accuracy: 0.7364 - val_loss: 0.7986 - val_sparse_categori
cal_accuracy: 0.7437
Epoch 25/30
1500/1500 [=====] - 1s 615us/step - loss: 0.8041 -
sparse_categorical_accuracy: 0.7400 - val_loss: 0.7882 - val_sparse_categori
cal_accuracy: 0.7456
Epoch 26/30
1500/1500 [=====] - 1s 625us/step - loss: 0.7938 -
sparse_categorical_accuracy: 0.7430 - val_loss: 0.7785 - val_sparse_categori
cal_accuracy: 0.7498
Epoch 27/30

1500/1500 [=====] - 1s 623us/step - loss: 0.7843 -
sparse_categorical_accuracy: 0.7463 - val_loss: 0.7696 - val_sparse_categorical_accuracy: 0.7530
Epoch 28/30
1500/1500 [=====] - 1s 620us/step - loss: 0.7754 -
sparse_categorical_accuracy: 0.7486 - val_loss: 0.7612 - val_sparse_categorical_accuracy: 0.7553
Epoch 29/30
1500/1500 [=====] - 1s 647us/step - loss: 0.7669 -
sparse_categorical_accuracy: 0.7513 - val_loss: 0.7528 - val_sparse_categorical_accuracy: 0.7566
Epoch 30/30
1500/1500 [=====] - 1s 669us/step - loss: 0.7586 -
sparse_categorical_accuracy: 0.7544 - val_loss: 0.7451 - val_sparse_categorical_accuracy: 0.7584

Adadelata Test Accuracy:

313/313 [=====] - 0s 327us/step
0.738

Training, Ftrl optimizer

Epoch 1/30
1500/1500 [=====] - 1s 762us/step - loss: 2.2342 -
sparse_categorical_accuracy: 0.2295 - val_loss: 2.0115 - val_sparse_categorical_accuracy: 0.2884
Epoch 2/30
1500/1500 [=====] - 1s 712us/step - loss: 1.7400 -
sparse_categorical_accuracy: 0.3511 - val_loss: 1.5382 - val_sparse_categorical_accuracy: 0.4240
Epoch 3/30
1500/1500 [=====] - 1s 711us/step - loss: 1.4203 -
sparse_categorical_accuracy: 0.5265 - val_loss: 1.3115 - val_sparse_categorical_accuracy: 0.5709
Epoch 4/30
1500/1500 [=====] - 1s 667us/step - loss: 1.2377 -
sparse_categorical_accuracy: 0.5799 - val_loss: 1.1580 - val_sparse_categorical_accuracy: 0.6132
Epoch 5/30
1500/1500 [=====] - 1s 663us/step - loss: 1.1078 -
sparse_categorical_accuracy: 0.6179 - val_loss: 1.0460 - val_sparse_categorical_accuracy: 0.6366
Epoch 6/30
1500/1500 [=====] - 1s 686us/step - loss: 1.0130 -
sparse_categorical_accuracy: 0.6402 - val_loss: 0.9642 - val_sparse_categorical_accuracy: 0.6565
Epoch 7/30
1500/1500 [=====] - 1s 665us/step - loss: 0.9443 -
sparse_categorical_accuracy: 0.6561 - val_loss: 0.9059 - val_sparse_categorical_accuracy: 0.6684
Epoch 8/30
1500/1500 [=====] - 1s 668us/step - loss: 0.8945 -
sparse_categorical_accuracy: 0.6672 - val_loss: 0.8633 - val_sparse_categorical_accuracy: 0.6784

cal_accuracy: 0.6837
Epoch 9/30
1500/1500 [=====] - 1s 664us/step - loss: 0.8573 -
sparse_categorical_accuracy: 0.6800 - val_loss: 0.8311 - val_sparse_categori
cal_accuracy: 0.6911
Epoch 10/30
1500/1500 [=====] - 1s 674us/step - loss: 0.8282 -
sparse_categorical_accuracy: 0.6900 - val_loss: 0.8058 - val_sparse_categori
cal_accuracy: 0.7019
Epoch 11/30
1500/1500 [=====] - 1s 674us/step - loss: 0.8047 -
sparse_categorical_accuracy: 0.6994 - val_loss: 0.7848 - val_sparse_categori
cal_accuracy: 0.7059
Epoch 12/30
1500/1500 [=====] - 1s 677us/step - loss: 0.7850 -
sparse_categorical_accuracy: 0.7059 - val_loss: 0.7679 - val_sparse_categori
cal_accuracy: 0.7166
Epoch 13/30
1500/1500 [=====] - 1s 675us/step - loss: 0.7682 -
sparse_categorical_accuracy: 0.7139 - val_loss: 0.7525 - val_sparse_categori
cal_accuracy: 0.7206
Epoch 14/30
1500/1500 [=====] - 1s 666us/step - loss: 0.7535 -
sparse_categorical_accuracy: 0.7197 - val_loss: 0.7391 - val_sparse_categori
cal_accuracy: 0.7278
Epoch 15/30
1500/1500 [=====] - 1s 673us/step - loss: 0.7404 -
sparse_categorical_accuracy: 0.7254 - val_loss: 0.7272 - val_sparse_categori
cal_accuracy: 0.7337
Epoch 16/30
1500/1500 [=====] - 1s 665us/step - loss: 0.7286 -
sparse_categorical_accuracy: 0.7317 - val_loss: 0.7165 - val_sparse_categori
cal_accuracy: 0.7380
Epoch 17/30
1500/1500 [=====] - 1s 678us/step - loss: 0.7180 -
sparse_categorical_accuracy: 0.7356 - val_loss: 0.7068 - val_sparse_categori
cal_accuracy: 0.7427
Epoch 18/30
1500/1500 [=====] - 1s 673us/step - loss: 0.7082 -
sparse_categorical_accuracy: 0.7399 - val_loss: 0.6983 - val_sparse_categori
cal_accuracy: 0.7462
Epoch 19/30
1500/1500 [=====] - 1s 672us/step - loss: 0.6992 -
sparse_categorical_accuracy: 0.7437 - val_loss: 0.6898 - val_sparse_categori
cal_accuracy: 0.7500
Epoch 20/30
1500/1500 [=====] - 1s 675us/step - loss: 0.6909 -
sparse_categorical_accuracy: 0.7474 - val_loss: 0.6821 - val_sparse_categori
cal_accuracy: 0.7536
Epoch 21/30
1500/1500 [=====] - 1s 669us/step - loss: 0.6833 -
sparse_categorical_accuracy: 0.7511 - val_loss: 0.6751 - val_sparse_categori
cal_accuracy: 0.7550

```

Epoch 22/30
1500/1500 [=====] - 1s 675us/step - loss: 0.6761 -
sparse_categorical_accuracy: 0.7543 - val_loss: 0.6687 - val_sparse_categori
cal_accuracy: 0.7571
Epoch 23/30
1500/1500 [=====] - 1s 695us/step - loss: 0.6694 -
sparse_categorical_accuracy: 0.7574 - val_loss: 0.6627 - val_sparse_categori
cal_accuracy: 0.7600
Epoch 24/30
1500/1500 [=====] - 1s 694us/step - loss: 0.6632 -
sparse_categorical_accuracy: 0.7598 - val_loss: 0.6568 - val_sparse_categori
cal_accuracy: 0.7619
Epoch 25/30
1500/1500 [=====] - 1s 692us/step - loss: 0.6574 -
sparse_categorical_accuracy: 0.7624 - val_loss: 0.6514 - val_sparse_categori
cal_accuracy: 0.7645
Epoch 26/30
1500/1500 [=====] - 1s 690us/step - loss: 0.6519 -
sparse_categorical_accuracy: 0.7647 - val_loss: 0.6465 - val_sparse_categori
cal_accuracy: 0.7673
Epoch 27/30
1500/1500 [=====] - 1s 689us/step - loss: 0.6468 -
sparse_categorical_accuracy: 0.7666 - val_loss: 0.6416 - val_sparse_categori
cal_accuracy: 0.7688
Epoch 28/30
1500/1500 [=====] - 1s 676us/step - loss: 0.6418 -
sparse_categorical_accuracy: 0.7681 - val_loss: 0.6372 - val_sparse_categori
cal_accuracy: 0.7697
Epoch 29/30
1500/1500 [=====] - 1s 652us/step - loss: 0.6373 -
sparse_categorical_accuracy: 0.7705 - val_loss: 0.6330 - val_sparse_categori
cal_accuracy: 0.7722
Epoch 30/30
1500/1500 [=====] - 1s 681us/step - loss: 0.6329 -
sparse_categorical_accuracy: 0.7727 - val_loss: 0.6288 - val_sparse_categori
cal_accuracy: 0.7728

Ftrl Test Accuracy:
313/313 [=====] - 0s 305us/step
0.7527

```

```

In [ ]: # plot training histories for each model on same plot
import matplotlib.cm as cm

fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

colors = cm.cool(np.linspace(0, 1, len(names)))

# add training
for name, history, color in zip(names, histories, colors):
    # training accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['sparse_categorical_accuracy'],
            label = f"{name} Training",

```

```

        color = color,
        linewidth=1)

    # validation accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['val_sparse_categorical_accuracy'],
            label = f"{name} Validation",
            color = color,
            linewidth=0.5,
            linestyle="--")

    ax.set_title("Accuracy Comparison of FFN Optimizers")
    ax.set_xlabel('Epoch')
    ax.set_xlim(1, len(histories[0].history["loss"]))
    ax.set_ylabel('Accuracy')
    ax.legend()
    ax.legend().get_frame().set_alpha(1.0)

    ax.grid(True)

fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

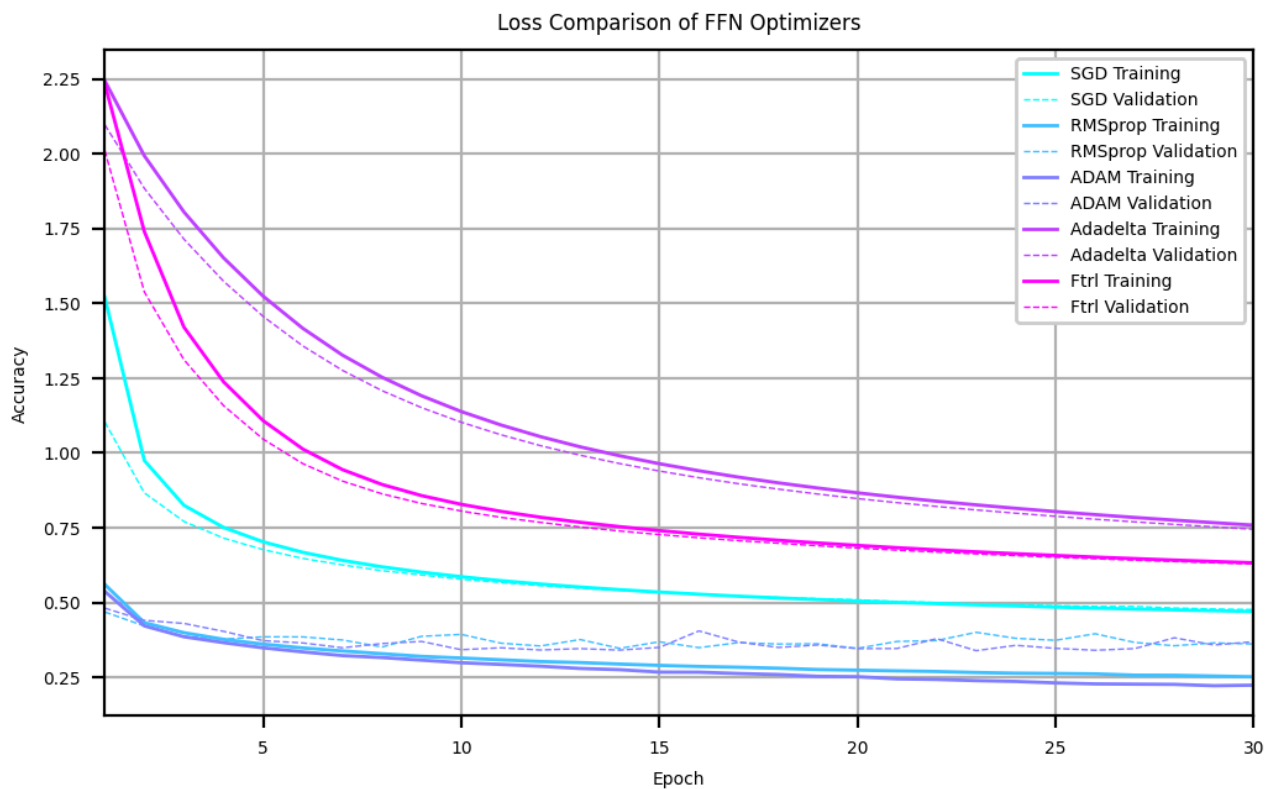

# add training
for name, history, color in zip(names, histories, colors):
    # training accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['loss'],
            label = f"{name} Training",
            color = color,
            linewidth=1)

    # validation accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['val_loss'],
            label = f"{name} Validation",
            color = color,
            linewidth=0.5,
            linestyle="--")

    ax.set_title("Loss Comparison of FFN Optimizers")
    ax.set_xlabel('Epoch')
    ax.set_xlim(1, len(histories[0].history["loss"]))
    ax.set_ylabel('Accuracy')
    ax.legend()
    ax.legend().get_frame().set_alpha(1.0)

    ax.grid(True)

```



```
In [ ]: import pandas as pd

# create summary table
optimizer_summary = pd.DataFrame(
    list(zip(names, train_loss, train_accuracies, test_accuracies)),
    columns=["Optimizer", "Training Loss", "Training Accuracy", "Test Accuracy"])
```

```
).set_index("Optimizer")
```

```
optimizer_summary
```

Out []:

	Training Loss	Training Accuracy	Test Accuracy
Optimizer			
SGD	0.469906	0.839021	0.7912
RMSprop	0.251783	0.909292	0.8504
ADAM	0.223793	0.916229	0.8630
Adadelta	0.758649	0.754375	0.7380
Ftrl	0.632887	0.772667	0.7527

Of the five optimizers analyzed, the Adam and RMSProp optimizers performed the best. The similar performance is most likely due to the Adam optimizer implementing the same technique as RMSProp, using the second momentum method.

Hidden Units Influence

```
In [ ]: # create list of hidden units to evaluate
units = [8, 64, 1024]

# instantiate tracking objects for analysis
histories = [] # empty list to save history objects for each model
train_loss = [] # save train loss for each mode
train_accuracies = [] # save train accuracy for each model
test_accuracies = [] # save test accuracy for each model

# begin training loop
print("Begin evaluation:")

for unit in units:

    # clear enviornment of previous model
    keras.backend.clear_session()

    # create ffn models using varaible hidden units
    model = model_arch(unit, "LeakyReLU")

    # compile models with Adam optimizer
    model.compile(optimizer=keras.optimizers.legacy.Adam(learning_rate=1e-3),
                  loss='sparse_categorical_crossentropy',
                  metrics=['sparse_categorical_accuracy'],
                  )

    # begin training, for the sake of brevity, number of epochs has been hal
    print(f"\nTraining, {unit} Hidden Units\n")
```



```
# show model summary and number of trainable parameters
model.summary()

history = model.fit(
    x_train_model.astype(np.float32), y_train,
    epochs=30,
    validation_split=0.2,
    shuffle = True
)

# evaluate on test data
print(f"\nTest Accuracy:")
predicted_softmax = model.predict(x_test.astype(np.float32))
y_pred = np.argmax(predicted_softmax, axis = 1)

test_accuracy = accuracy_score(y_test, y_pred)

print(test_accuracy)

# append tracking objects
histories.append(history)
train_loss.append(history.history["loss"][-1])
train_accuracies.append(history.history["sparse_categorical_accuracy"][-1])
test_accuracies.append(test_accuracy)
```

Begin evaluation:

Training, 8 Hidden Units

Model: "units-8_activation-LeakyReLU"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
hidden_layer (Dense)	(None, 8)	6280
leaky_re_lu (LeakyReLU)	(None, 8)	0
output_layer (Dense)	(None, 10)	90

=====
Total params: 6370 (24.88 KB)
Trainable params: 6370 (24.88 KB)
Non-trainable params: 0 (0.00 Byte)

Epoch 1/30
1500/1500 [=====] - 1s 524us/step - loss: 0.7383 - sparse_categorical_accuracy: 0.7495 - val_loss: 0.5326 - val_sparse_categorical_accuracy: 0.8173
Epoch 2/30
1500/1500 [=====] - 1s 459us/step - loss: 0.4994 - sparse_categorical_accuracy: 0.8288 - val_loss: 0.4845 - val_sparse_categorical_accuracy: 0.8375

cal_accuracy: 0.8315
Epoch 3/30
1500/1500 [=====] - 1s 425us/step - loss: 0.4591 -
sparse_categorical_accuracy: 0.8390 - val_loss: 0.4739 - val_sparse_categori
cal_accuracy: 0.8332
Epoch 4/30
1500/1500 [=====] - 1s 457us/step - loss: 0.4393 -
sparse_categorical_accuracy: 0.8468 - val_loss: 0.4524 - val_sparse_categori
cal_accuracy: 0.8381
Epoch 5/30
1500/1500 [=====] - 1s 447us/step - loss: 0.4270 -
sparse_categorical_accuracy: 0.8507 - val_loss: 0.4417 - val_sparse_categori
cal_accuracy: 0.8473
Epoch 6/30
1500/1500 [=====] - 1s 430us/step - loss: 0.4183 -
sparse_categorical_accuracy: 0.8535 - val_loss: 0.4375 - val_sparse_categori
cal_accuracy: 0.8483
Epoch 7/30
1500/1500 [=====] - 1s 446us/step - loss: 0.4111 -
sparse_categorical_accuracy: 0.8561 - val_loss: 0.4359 - val_sparse_categori
cal_accuracy: 0.8472
Epoch 8/30
1500/1500 [=====] - 1s 446us/step - loss: 0.4065 -
sparse_categorical_accuracy: 0.8588 - val_loss: 0.4509 - val_sparse_categori
cal_accuracy: 0.8428
Epoch 9/30
1500/1500 [=====] - 1s 429us/step - loss: 0.4016 -
sparse_categorical_accuracy: 0.8589 - val_loss: 0.4207 - val_sparse_categori
cal_accuracy: 0.8535
Epoch 10/30
1500/1500 [=====] - 1s 449us/step - loss: 0.3972 -
sparse_categorical_accuracy: 0.8607 - val_loss: 0.4252 - val_sparse_categori
cal_accuracy: 0.8528
Epoch 11/30
1500/1500 [=====] - 1s 445us/step - loss: 0.3938 -
sparse_categorical_accuracy: 0.8614 - val_loss: 0.4257 - val_sparse_categori
cal_accuracy: 0.8509
Epoch 12/30
1500/1500 [=====] - 1s 436us/step - loss: 0.3912 -
sparse_categorical_accuracy: 0.8615 - val_loss: 0.4184 - val_sparse_categori
cal_accuracy: 0.8523
Epoch 13/30
1500/1500 [=====] - 1s 447us/step - loss: 0.3864 -
sparse_categorical_accuracy: 0.8638 - val_loss: 0.4160 - val_sparse_categori
cal_accuracy: 0.8546
Epoch 14/30
1500/1500 [=====] - 1s 426us/step - loss: 0.3853 -
sparse_categorical_accuracy: 0.8633 - val_loss: 0.4212 - val_sparse_categori
cal_accuracy: 0.8530
Epoch 15/30
1500/1500 [=====] - 1s 442us/step - loss: 0.3834 -
sparse_categorical_accuracy: 0.8642 - val_loss: 0.4293 - val_sparse_categori
cal_accuracy: 0.8480

Epoch 16/30
1500/1500 [=====] - 1s 459us/step - loss: 0.3812 - sparse_categorical_accuracy: 0.8653 - val_loss: 0.4185 - val_sparse_categorical_accuracy: 0.8503

Epoch 17/30
1500/1500 [=====] - 1s 426us/step - loss: 0.3786 - sparse_categorical_accuracy: 0.8673 - val_loss: 0.4206 - val_sparse_categorical_accuracy: 0.8512

Epoch 18/30
1500/1500 [=====] - 1s 444us/step - loss: 0.3766 - sparse_categorical_accuracy: 0.8669 - val_loss: 0.4266 - val_sparse_categorical_accuracy: 0.8509

Epoch 19/30
1500/1500 [=====] - 1s 447us/step - loss: 0.3746 - sparse_categorical_accuracy: 0.8669 - val_loss: 0.4120 - val_sparse_categorical_accuracy: 0.8562

Epoch 20/30
1500/1500 [=====] - 1s 428us/step - loss: 0.3745 - sparse_categorical_accuracy: 0.8672 - val_loss: 0.4181 - val_sparse_categorical_accuracy: 0.8530

Epoch 21/30
1500/1500 [=====] - 1s 447us/step - loss: 0.3724 - sparse_categorical_accuracy: 0.8691 - val_loss: 0.4138 - val_sparse_categorical_accuracy: 0.8526

Epoch 22/30
1500/1500 [=====] - 1s 430us/step - loss: 0.3701 - sparse_categorical_accuracy: 0.8685 - val_loss: 0.4169 - val_sparse_categorical_accuracy: 0.8556

Epoch 23/30
1500/1500 [=====] - 1s 449us/step - loss: 0.3689 - sparse_categorical_accuracy: 0.8693 - val_loss: 0.4168 - val_sparse_categorical_accuracy: 0.8545

Epoch 24/30
1500/1500 [=====] - 1s 442us/step - loss: 0.3680 - sparse_categorical_accuracy: 0.8686 - val_loss: 0.4138 - val_sparse_categorical_accuracy: 0.8538

Epoch 25/30
1500/1500 [=====] - 1s 427us/step - loss: 0.3665 - sparse_categorical_accuracy: 0.8702 - val_loss: 0.4111 - val_sparse_categorical_accuracy: 0.8581

Epoch 26/30
1500/1500 [=====] - 1s 442us/step - loss: 0.3654 - sparse_categorical_accuracy: 0.8706 - val_loss: 0.4169 - val_sparse_categorical_accuracy: 0.8541

Epoch 27/30
1500/1500 [=====] - 1s 442us/step - loss: 0.3635 - sparse_categorical_accuracy: 0.8705 - val_loss: 0.4220 - val_sparse_categorical_accuracy: 0.8547

Epoch 28/30
1500/1500 [=====] - 1s 429us/step - loss: 0.3640 - sparse_categorical_accuracy: 0.8705 - val_loss: 0.4118 - val_sparse_categorical_accuracy: 0.8558

Epoch 29/30

1500/1500 [=====] - 1s 453us/step - loss: 0.3633 - sparse_categorical_accuracy: 0.8715 - val_loss: 0.4236 - val_sparse_categorical_accuracy: 0.8521
Epoch 30/30
1500/1500 [=====] - 1s 447us/step - loss: 0.3601 - sparse_categorical_accuracy: 0.8724 - val_loss: 0.4139 - val_sparse_categorical_accuracy: 0.8551

Test Accuracy:

313/313 [=====] - 0s 235us/step
0.7652

Training, 64 Hidden Units

Model: "units-64_activation-LeakyReLU"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
hidden_layer (Dense)	(None, 64)	50240
leaky_re_lu (LeakyReLU)	(None, 64)	0
output_layer (Dense)	(None, 10)	650

=====
Total params: 50890 (198.79 KB)
Trainable params: 50890 (198.79 KB)
Non-trainable params: 0 (0.00 Byte)

Epoch 1/30

1500/1500 [=====] - 1s 646us/step - loss: 0.5420 - sparse_categorical_accuracy: 0.8101 - val_loss: 0.4382 - val_sparse_categorical_accuracy: 0.8465

Epoch 2/30

1500/1500 [=====] - 1s 625us/step - loss: 0.4217 - sparse_categorical_accuracy: 0.8521 - val_loss: 0.4008 - val_sparse_categorical_accuracy: 0.8589

Epoch 3/30

1500/1500 [=====] - 1s 635us/step - loss: 0.3894 - sparse_categorical_accuracy: 0.8617 - val_loss: 0.3732 - val_sparse_categorical_accuracy: 0.8679

Epoch 4/30

1500/1500 [=====] - 1s 625us/step - loss: 0.3648 - sparse_categorical_accuracy: 0.8694 - val_loss: 0.3754 - val_sparse_categorical_accuracy: 0.8670

Epoch 5/30

1500/1500 [=====] - 1s 632us/step - loss: 0.3502 - sparse_categorical_accuracy: 0.8737 - val_loss: 0.3807 - val_sparse_categorical_accuracy: 0.8615

Epoch 6/30

1500/1500 [=====] - 1s 629us/step - loss: 0.3386 -

sparse_categorical_accuracy: 0.8765 - val_loss: 0.3710 - val_sparse_categorical_accuracy: 0.8637
Epoch 7/30
1500/1500 [=====] - 1s 631us/step - loss: 0.3225 - sparse_categorical_accuracy: 0.8821 - val_loss: 0.3481 - val_sparse_categorical_accuracy: 0.8734
Epoch 8/30
1500/1500 [=====] - 1s 633us/step - loss: 0.3128 - sparse_categorical_accuracy: 0.8847 - val_loss: 0.3562 - val_sparse_categorical_accuracy: 0.8737
Epoch 9/30
1500/1500 [=====] - 1s 632us/step - loss: 0.3070 - sparse_categorical_accuracy: 0.8880 - val_loss: 0.3754 - val_sparse_categorical_accuracy: 0.8652
Epoch 10/30
1500/1500 [=====] - 1s 633us/step - loss: 0.3002 - sparse_categorical_accuracy: 0.8897 - val_loss: 0.3448 - val_sparse_categorical_accuracy: 0.8769
Epoch 11/30
1500/1500 [=====] - 1s 627us/step - loss: 0.2909 - sparse_categorical_accuracy: 0.8935 - val_loss: 0.3589 - val_sparse_categorical_accuracy: 0.8725
Epoch 12/30
1500/1500 [=====] - 1s 632us/step - loss: 0.2842 - sparse_categorical_accuracy: 0.8952 - val_loss: 0.3515 - val_sparse_categorical_accuracy: 0.8738
Epoch 13/30
1500/1500 [=====] - 1s 630us/step - loss: 0.2799 - sparse_categorical_accuracy: 0.8967 - val_loss: 0.3382 - val_sparse_categorical_accuracy: 0.8811
Epoch 14/30
1500/1500 [=====] - 1s 632us/step - loss: 0.2751 - sparse_categorical_accuracy: 0.8982 - val_loss: 0.3491 - val_sparse_categorical_accuracy: 0.8729
Epoch 15/30
1500/1500 [=====] - 1s 624us/step - loss: 0.2711 - sparse_categorical_accuracy: 0.8999 - val_loss: 0.3476 - val_sparse_categorical_accuracy: 0.8772
Epoch 16/30
1500/1500 [=====] - 1s 631us/step - loss: 0.2632 - sparse_categorical_accuracy: 0.9023 - val_loss: 0.3466 - val_sparse_categorical_accuracy: 0.8808
Epoch 17/30
1500/1500 [=====] - 1s 627us/step - loss: 0.2596 - sparse_categorical_accuracy: 0.9040 - val_loss: 0.3373 - val_sparse_categorical_accuracy: 0.8838
Epoch 18/30
1500/1500 [=====] - 1s 655us/step - loss: 0.2571 - sparse_categorical_accuracy: 0.9045 - val_loss: 0.3425 - val_sparse_categorical_accuracy: 0.8788
Epoch 19/30
1500/1500 [=====] - 1s 627us/step - loss: 0.2540 - sparse_categorical_accuracy: 0.9057 - val_loss: 0.3332 - val_sparse_categorical_accuracy: 0.8788

```
cal_accuracy: 0.8824
Epoch 20/30
1500/1500 [=====] - 1s 637us/step - loss: 0.2498 -
sparse_categorical_accuracy: 0.9078 - val_loss: 0.3735 - val_sparse_categori
cal_accuracy: 0.8699
Epoch 21/30
1500/1500 [=====] - 1s 636us/step - loss: 0.2449 -
sparse_categorical_accuracy: 0.9088 - val_loss: 0.3396 - val_sparse_categori
cal_accuracy: 0.8834
Epoch 22/30
1500/1500 [=====] - 1s 629us/step - loss: 0.2405 -
sparse_categorical_accuracy: 0.9112 - val_loss: 0.3499 - val_sparse_categori
cal_accuracy: 0.8776
Epoch 23/30
1500/1500 [=====] - 1s 635us/step - loss: 0.2367 -
sparse_categorical_accuracy: 0.9131 - val_loss: 0.3708 - val_sparse_categori
cal_accuracy: 0.8765
Epoch 24/30
1500/1500 [=====] - 1s 640us/step - loss: 0.2334 -
sparse_categorical_accuracy: 0.9134 - val_loss: 0.3501 - val_sparse_categori
cal_accuracy: 0.8824
Epoch 25/30
1500/1500 [=====] - 1s 642us/step - loss: 0.2337 -
sparse_categorical_accuracy: 0.9134 - val_loss: 0.3629 - val_sparse_categori
cal_accuracy: 0.8777
Epoch 26/30
1500/1500 [=====] - 1s 640us/step - loss: 0.2297 -
sparse_categorical_accuracy: 0.9147 - val_loss: 0.3801 - val_sparse_categori
cal_accuracy: 0.8723
Epoch 27/30
1500/1500 [=====] - 1s 641us/step - loss: 0.2266 -
sparse_categorical_accuracy: 0.9162 - val_loss: 0.3602 - val_sparse_categori
cal_accuracy: 0.8778
Epoch 28/30
1500/1500 [=====] - 1s 618us/step - loss: 0.2248 -
sparse_categorical_accuracy: 0.9180 - val_loss: 0.3919 - val_sparse_categori
cal_accuracy: 0.8651
Epoch 29/30
1500/1500 [=====] - 1s 646us/step - loss: 0.2243 -
sparse_categorical_accuracy: 0.9160 - val_loss: 0.3542 - val_sparse_categori
cal_accuracy: 0.8811
Epoch 30/30
1500/1500 [=====] - 1s 669us/step - loss: 0.2184 -
sparse_categorical_accuracy: 0.9192 - val_loss: 0.3671 - val_sparse_categori
cal_accuracy: 0.8783

Test Accuracy:
313/313 [=====] - 0s 384us/step
0.8526
```

Training, 1024 Hidden Units

Model: "units-1024_activation-LeakyReLU"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
hidden_layer (Dense)	(None, 1024)	803840
leaky_re_lu (LeakyReLU)	(None, 1024)	0
output_layer (Dense)	(None, 10)	10250

```

=====
Total params: 814090 (3.11 MB)
Trainable params: 814090 (3.11 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```

```

Epoch 1/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.5212 - sparse_categorical_accuracy: 0.8164 - val_loss: 0.4235 - val_sparse_categorical_accuracy: 0.8499
Epoch 2/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.4192 - sparse_categorical_accuracy: 0.8493 - val_loss: 0.4915 - val_sparse_categorical_accuracy: 0.8162
Epoch 3/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3893 - sparse_categorical_accuracy: 0.8608 - val_loss: 0.3805 - val_sparse_categorical_accuracy: 0.8652
Epoch 4/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3693 - sparse_categorical_accuracy: 0.8676 - val_loss: 0.3787 - val_sparse_categorical_accuracy: 0.8681
Epoch 5/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3530 - sparse_categorical_accuracy: 0.8721 - val_loss: 0.3649 - val_sparse_categorical_accuracy: 0.8718
Epoch 6/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3373 - sparse_categorical_accuracy: 0.8774 - val_loss: 0.3542 - val_sparse_categorical_accuracy: 0.8722
Epoch 7/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3277 - sparse_categorical_accuracy: 0.8806 - val_loss: 0.3822 - val_sparse_categorical_accuracy: 0.8714
Epoch 8/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3124 - sparse_categorical_accuracy: 0.8865 - val_loss: 0.3824 - val_sparse_categorical_accuracy: 0.8705
Epoch 9/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.3102 - sparse_categorical_accuracy: 0.8855 - val_loss: 0.3504 - val_sparse_categorical_accuracy: 0.8769
Epoch 10/30

```

1500/1500 [=====] - 3s 2ms/step - loss: 0.3001 - sparse_categorical_accuracy: 0.8905 - val_loss: 0.3786 - val_sparse_categorical_accuracy: 0.8761
Epoch 11/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2913 - sparse_categorical_accuracy: 0.8931 - val_loss: 0.3994 - val_sparse_categorical_accuracy: 0.8670
Epoch 12/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2886 - sparse_categorical_accuracy: 0.8949 - val_loss: 0.3640 - val_sparse_categorical_accuracy: 0.8748
Epoch 13/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2832 - sparse_categorical_accuracy: 0.8964 - val_loss: 0.3867 - val_sparse_categorical_accuracy: 0.8714
Epoch 14/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2765 - sparse_categorical_accuracy: 0.8985 - val_loss: 0.3561 - val_sparse_categorical_accuracy: 0.8806
Epoch 15/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2747 - sparse_categorical_accuracy: 0.9005 - val_loss: 0.3644 - val_sparse_categorical_accuracy: 0.8780
Epoch 16/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2628 - sparse_categorical_accuracy: 0.9026 - val_loss: 0.3844 - val_sparse_categorical_accuracy: 0.8753
Epoch 17/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2664 - sparse_categorical_accuracy: 0.9030 - val_loss: 0.3542 - val_sparse_categorical_accuracy: 0.8824
Epoch 18/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2589 - sparse_categorical_accuracy: 0.9055 - val_loss: 0.3453 - val_sparse_categorical_accuracy: 0.8857
Epoch 19/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2531 - sparse_categorical_accuracy: 0.9054 - val_loss: 0.4316 - val_sparse_categorical_accuracy: 0.8590
Epoch 20/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2482 - sparse_categorical_accuracy: 0.9089 - val_loss: 0.3606 - val_sparse_categorical_accuracy: 0.8853
Epoch 21/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2441 - sparse_categorical_accuracy: 0.9105 - val_loss: 0.3888 - val_sparse_categorical_accuracy: 0.8798
Epoch 22/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2401 - sparse_categorical_accuracy: 0.9116 - val_loss: 0.3819 - val_sparse_categorical_accuracy: 0.8793
Epoch 23/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2401 - sparse_categorical_accuracy: 0.9116 - val_loss: 0.3819 - val_sparse_categorical_accuracy: 0.8793


```

arse_categorical_accuracy: 0.9121 - val_loss: 0.3585 - val_sparse_categorical_accuracy: 0.8836
Epoch 24/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2341 - sparse_categorical_accuracy: 0.9141 - val_loss: 0.4133 - val_sparse_categorical_accuracy: 0.8726
Epoch 25/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2300 - sparse_categorical_accuracy: 0.9146 - val_loss: 0.3711 - val_sparse_categorical_accuracy: 0.8846
Epoch 26/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2334 - sparse_categorical_accuracy: 0.9152 - val_loss: 0.3789 - val_sparse_categorical_accuracy: 0.8860
Epoch 27/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2198 - sparse_categorical_accuracy: 0.9201 - val_loss: 0.4031 - val_sparse_categorical_accuracy: 0.8759
Epoch 28/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2273 - sparse_categorical_accuracy: 0.9164 - val_loss: 0.3730 - val_sparse_categorical_accuracy: 0.8844
Epoch 29/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2192 - sparse_categorical_accuracy: 0.9200 - val_loss: 0.3823 - val_sparse_categorical_accuracy: 0.8838
Epoch 30/30
1500/1500 [=====] - 3s 2ms/step - loss: 0.2141 - sparse_categorical_accuracy: 0.9218 - val_loss: 0.4112 - val_sparse_categorical_accuracy: 0.8802

Test Accuracy:
313/313 [=====] - 0s 670us/step
0.857

```

```

In [ ]: # plot training histories for each model on same plot
fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)

colors = cm.Set1(np.linspace(0, 1, len(names)))

# add training
for unit, history, color in zip(units, histories, colors):
    # training accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['sparse_categorical_accuracy'],
            label = f'{unit} Hidden Units Training',
            color = color,
            linewidth=1)

    # validation accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['val_sparse_categorical_accuracy'],
            label = f'{unit} Hidden Units Validation',
            color = color,

```

```

        linewidth=0.5,
        linestyle="--")

ax.set_title("Accuracy Comparison of FFN Number of Hidden Units")
ax.set_xlabel('Epoch')
ax.set_xlim(1, len(histories[0].history["loss"]))
ax.set_ylabel('Accuracy')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)

ax.grid(True)

fig, ax = plt.subplots(figsize=(6,3.5), dpi=200)



# add training
for unit, history, color in zip(units, histories, colors):
    # training accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['loss'],
            label = f"{unit} Hidden Units Training",
            color = color,
            linewidth=1)

    # validation accuracy
    ax.plot(range(1, len(histories[0].history["loss"])+1),
            history.history['val_loss'],
            label = f"{unit} Hidden Units Validation",
            color = color,
            linewidth=0.5,
            linestyle="--")

ax.set_title("Loss Comparison of FFN Number of Hidden Units")
ax.set_xlabel('Epoch')
ax.set_xlim(1, len(histories[0].history["loss"]))
ax.set_ylabel('Accuracy')
ax.legend()
ax.legend().get_frame().set_alpha(1.0)

ax.grid(True)

```



```
In [ ]: # create summary table
optimizer_summary = pd.DataFrame(
    list(zip(units, train_loss, train_accuracies, test_accuracies)),
    columns=["Hidden Units", "Training Loss", "Training Accuracy", "Test Accuracy"],
    index=[0, 1, 2]).set_index("Hidden Units")

optimizer_summary
```

Out[]:

	Training Loss	Training Accuracy	Test Accuracy
Hidden Units			
8	0.360062	0.872417	0.7652
64	0.218399	0.919229	0.8526
1024	0.214138	0.921792	0.8570

In general, increasing the number of hidden units increases test accuracy up to a point. But from the summary table and loss plot it is apparent that these increases are diminishing in value and prone to overfitting.

6. Team Information

Team Github: [Team4](#)

Team Members:

- Arheum Kim: [ahreum239](#)
- Isaac Salvador: [isalva2](#)
- Sadjad Bazarnovi: [sadjad33](#)